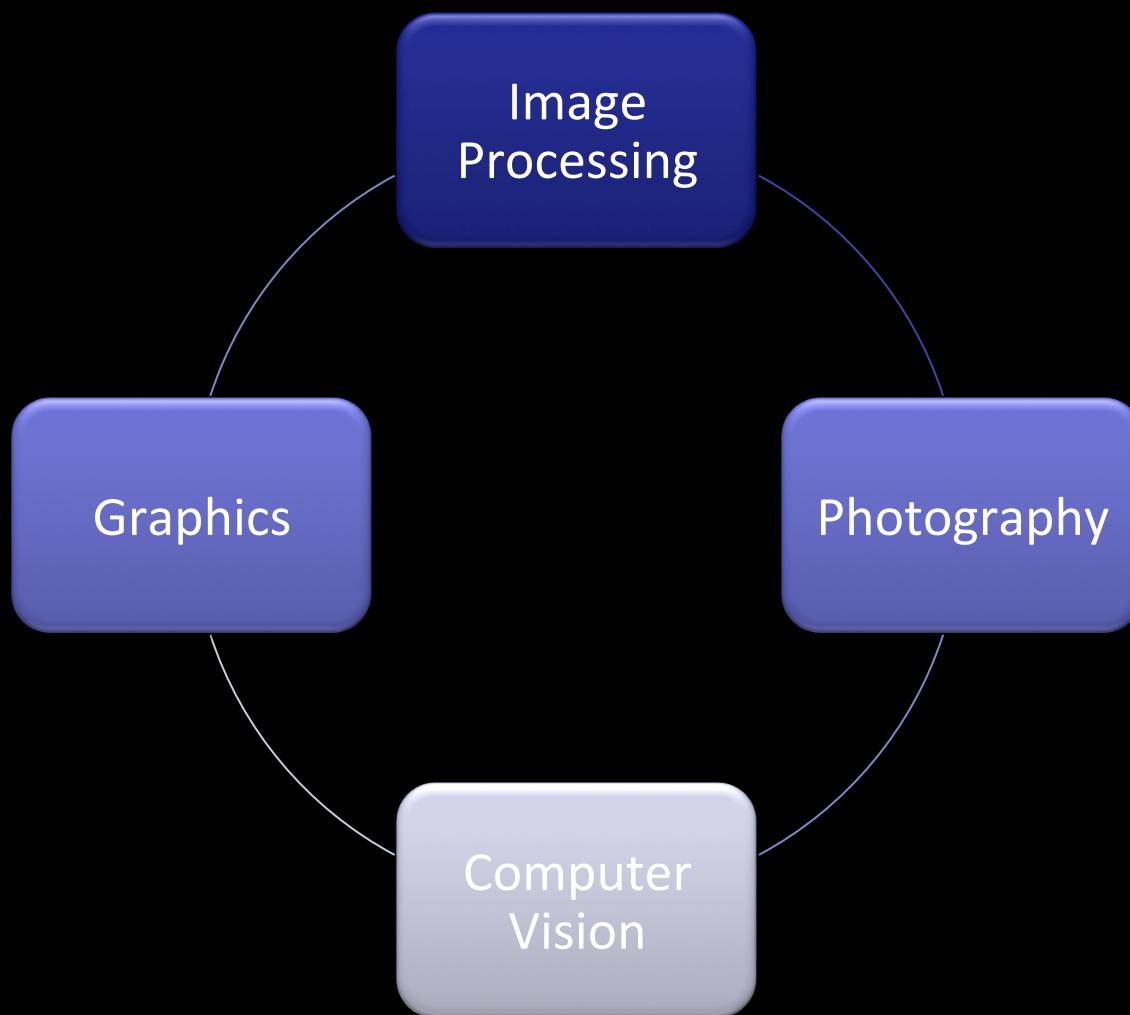


# Computational Imaging:

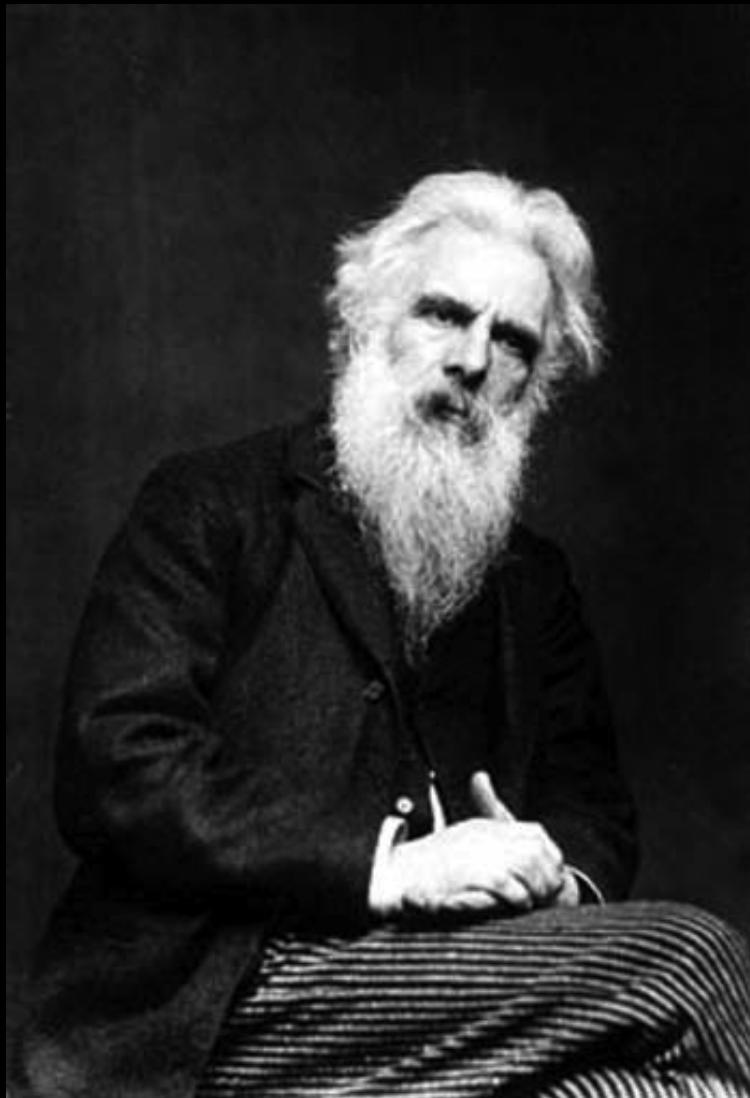
From Photons to Photos

Peyman Milanfar

# What is Computational Imaging?

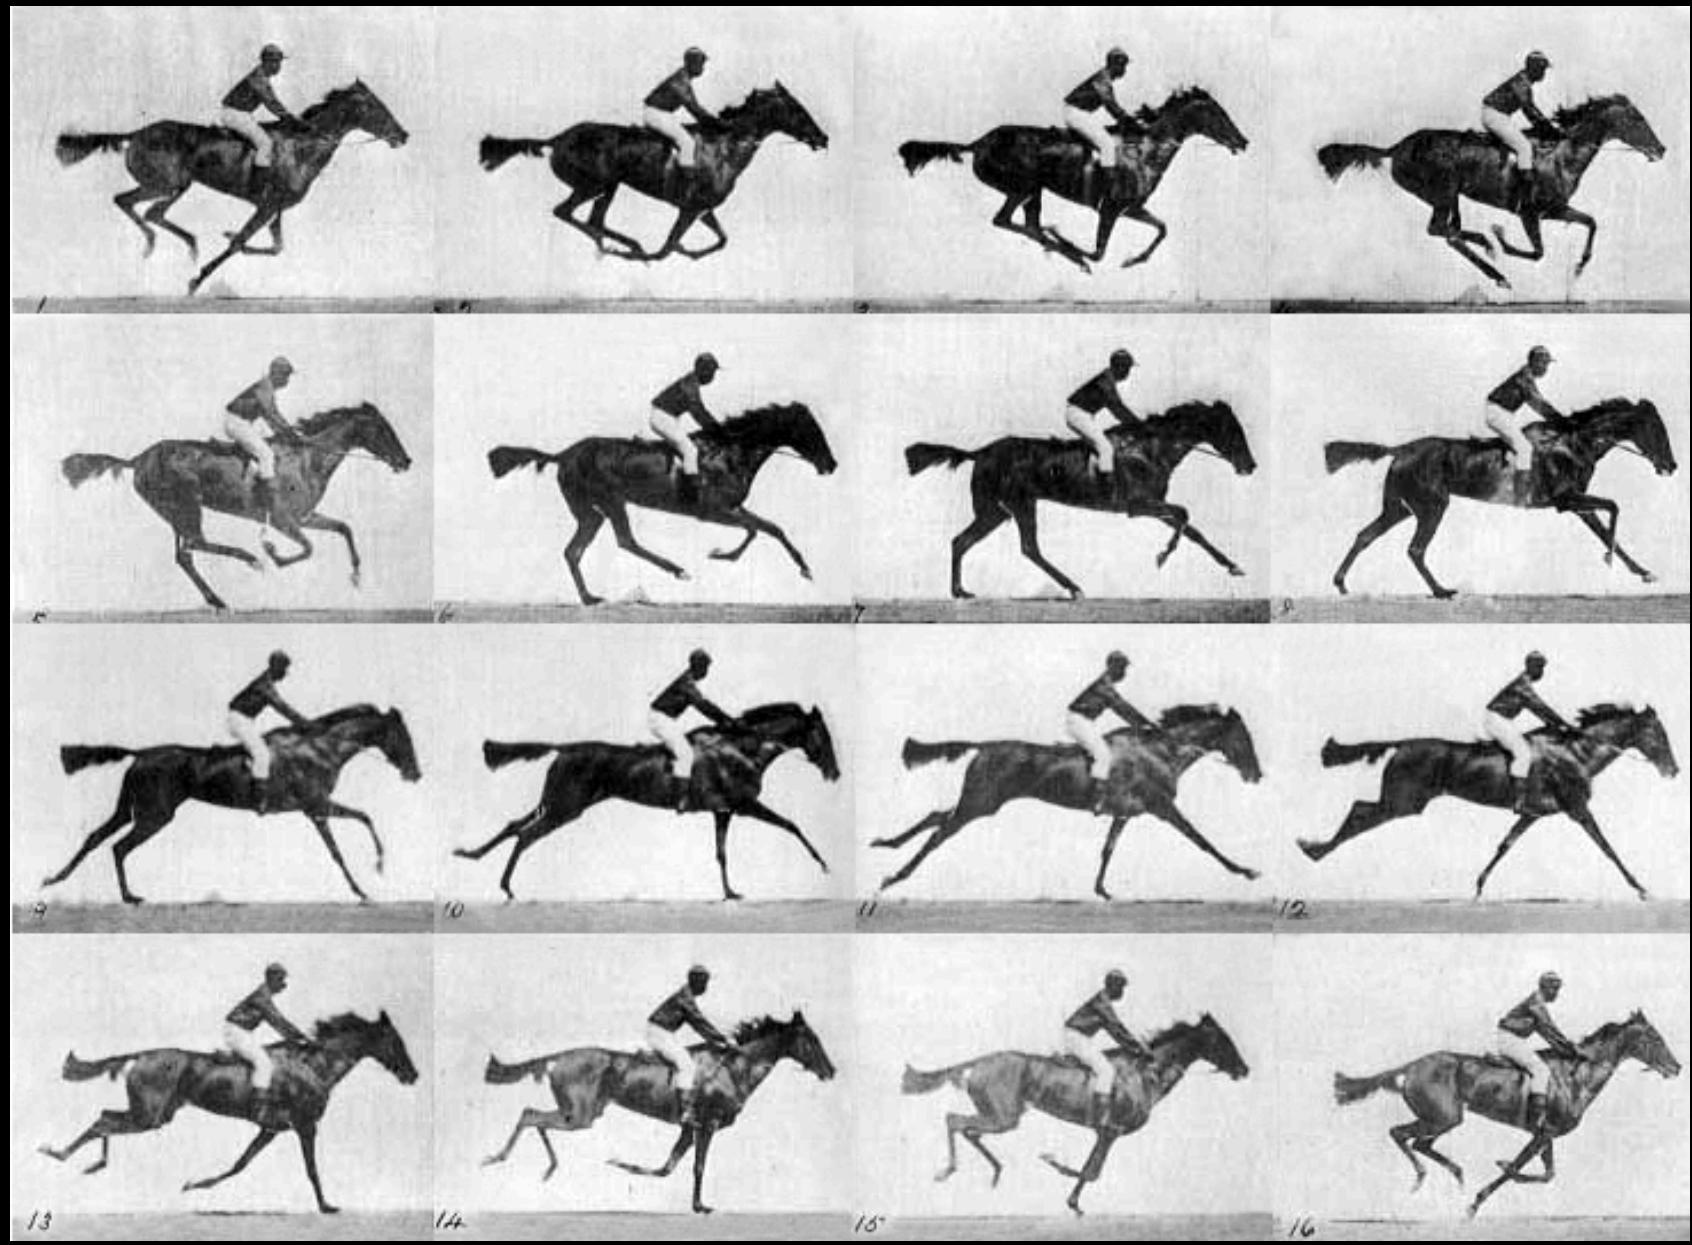


# Long and fascinating history

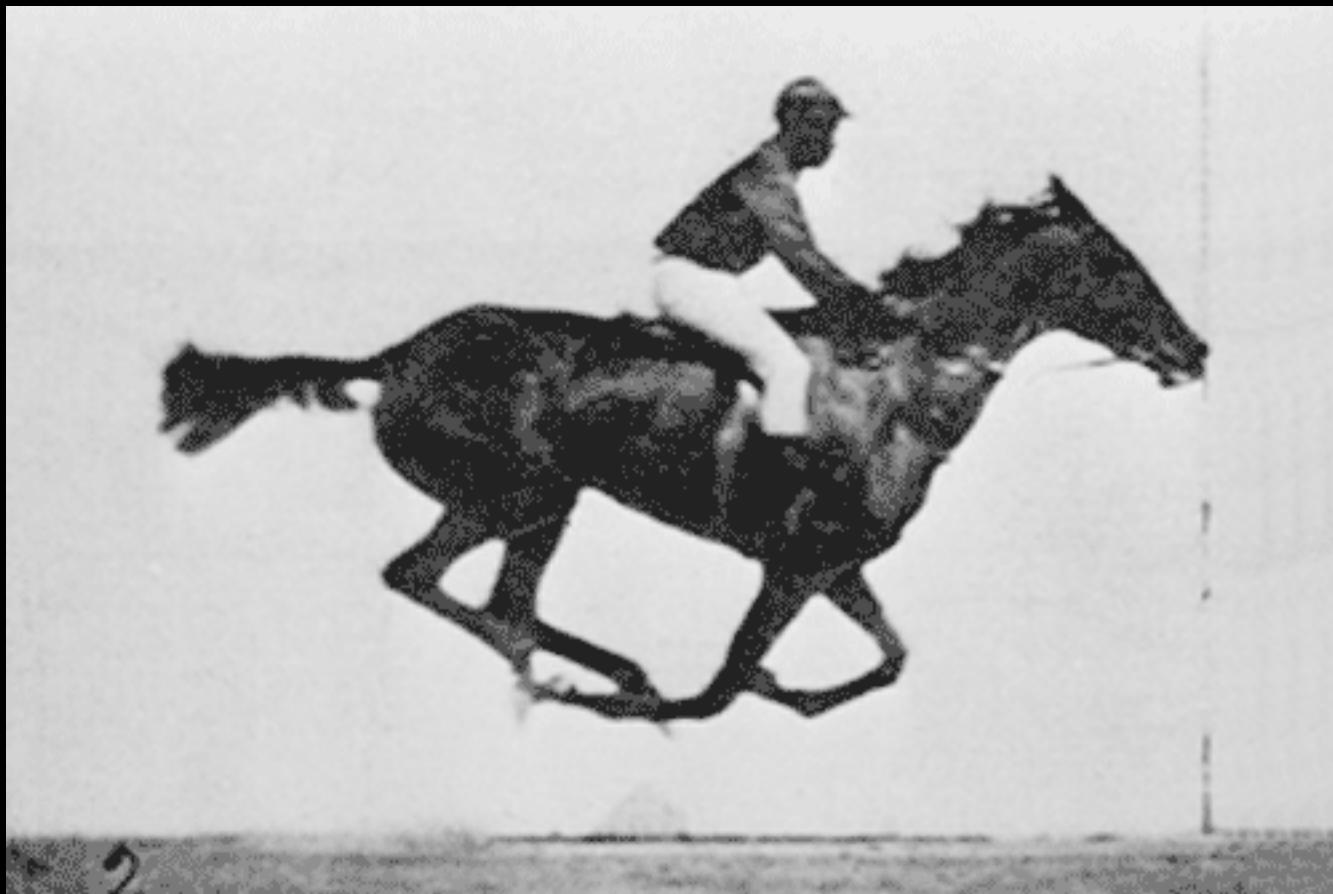


Eadweard Muybridge 1830 – 1904

# Galloping Horse 1878

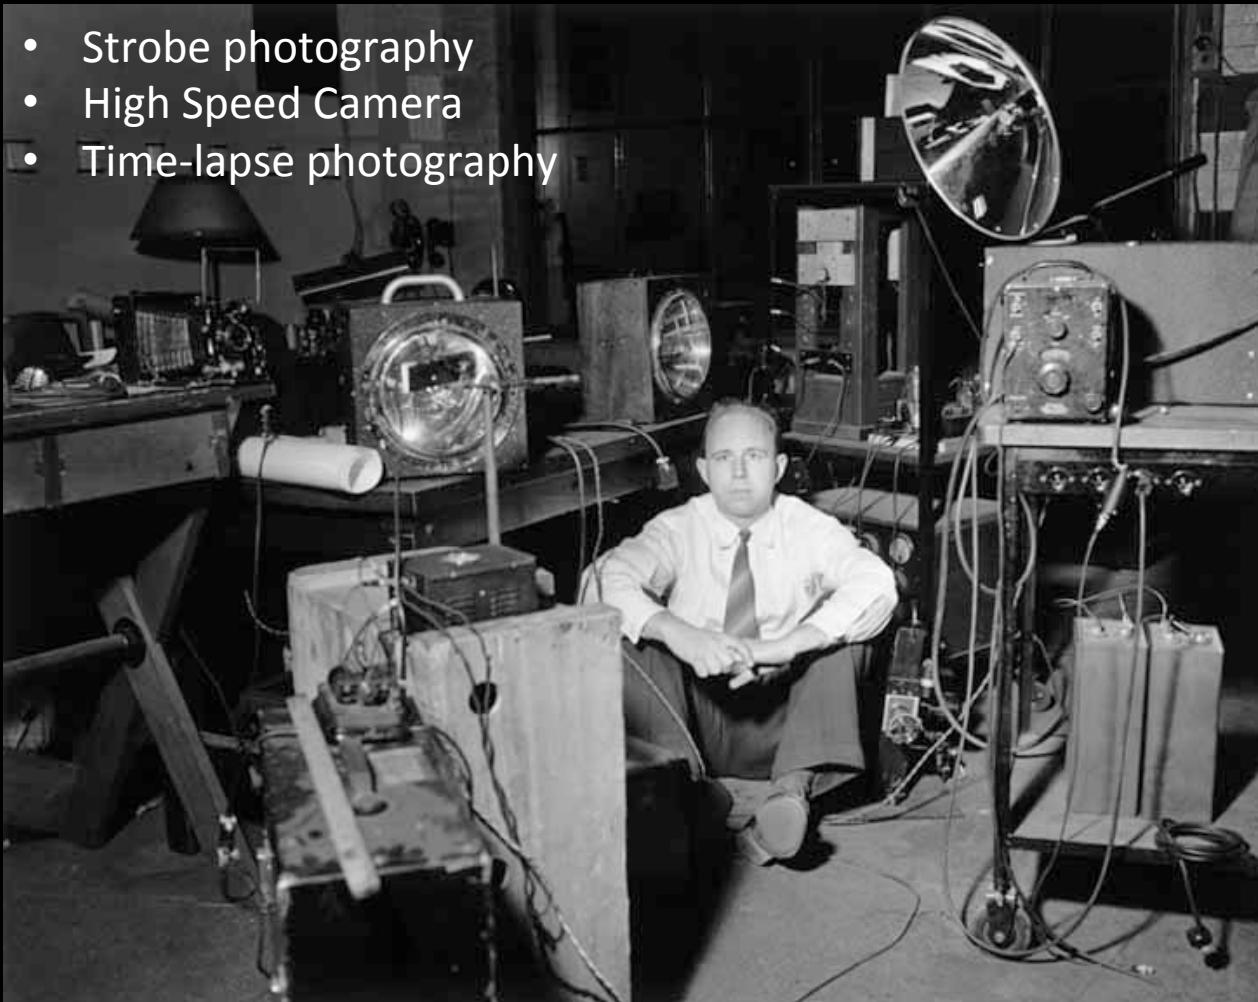


# GIF made in 2006



# Another Pioneer

- Strobe photography
- High Speed Camera
- Time-lapse photography



Harold "Doc" Edgerton (1903 – 1990)

# “Stopping Time”



Microphone picks up acoustic shock wave, triggers the flash.

Shot in Darkness, with sync-strobe, exposure time  $\sim 10^{-7}$  s

# 1980s and 90s



- HDR
- wearable computers
- super-resolution
- panorama stitching

Not Computational Imaging

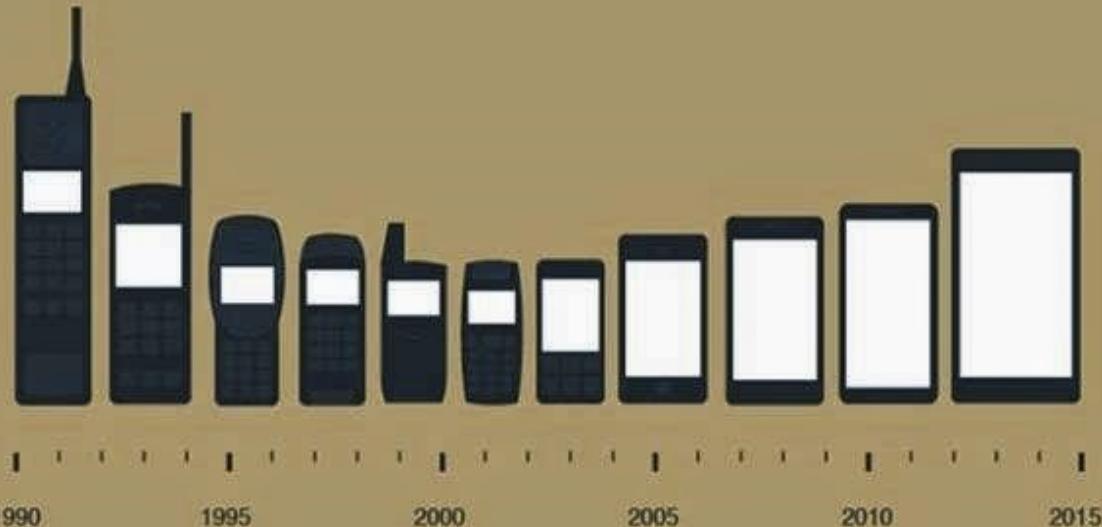


# The Modern Era of Easy Photography

Take it  
See it  
Share it

# Evolution of (cheap) Photography

# Evolution of the (Smart) phone





# 2015 -- Paris



AP

2015 -- Moscow



AP

# 2015 -- Sydney



Reuters

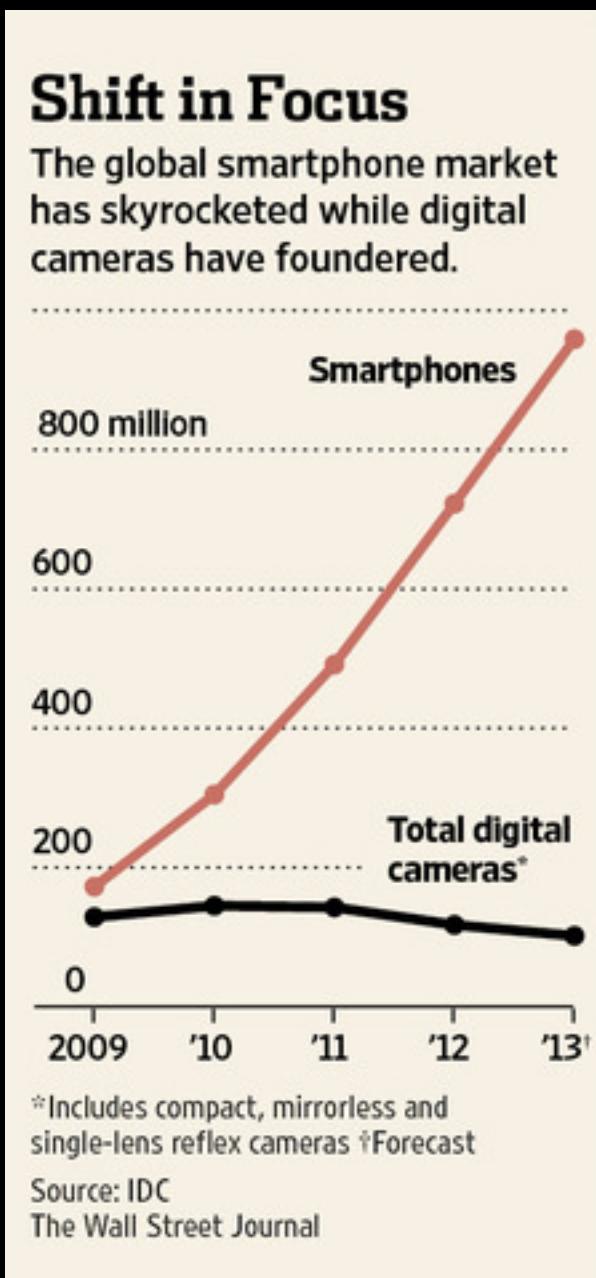
1,000,000,000,000

Roughly a **trillion** photos *shared* on  
social media in 2014

7,000,000,000

As many as **seven billion** were “selfies”

# Mobile Imaging is Totally Dominant



# Dual Aims of (Computational) Imaging

# Capture what I see (Photography)

- Take a nice picture. (1870s - )
  - Make me a better photographer. (1970s - )
  - Do it with my simple camera. (2000s - )

# Let's have a look

# Sensors: Form vs. Function

\$500

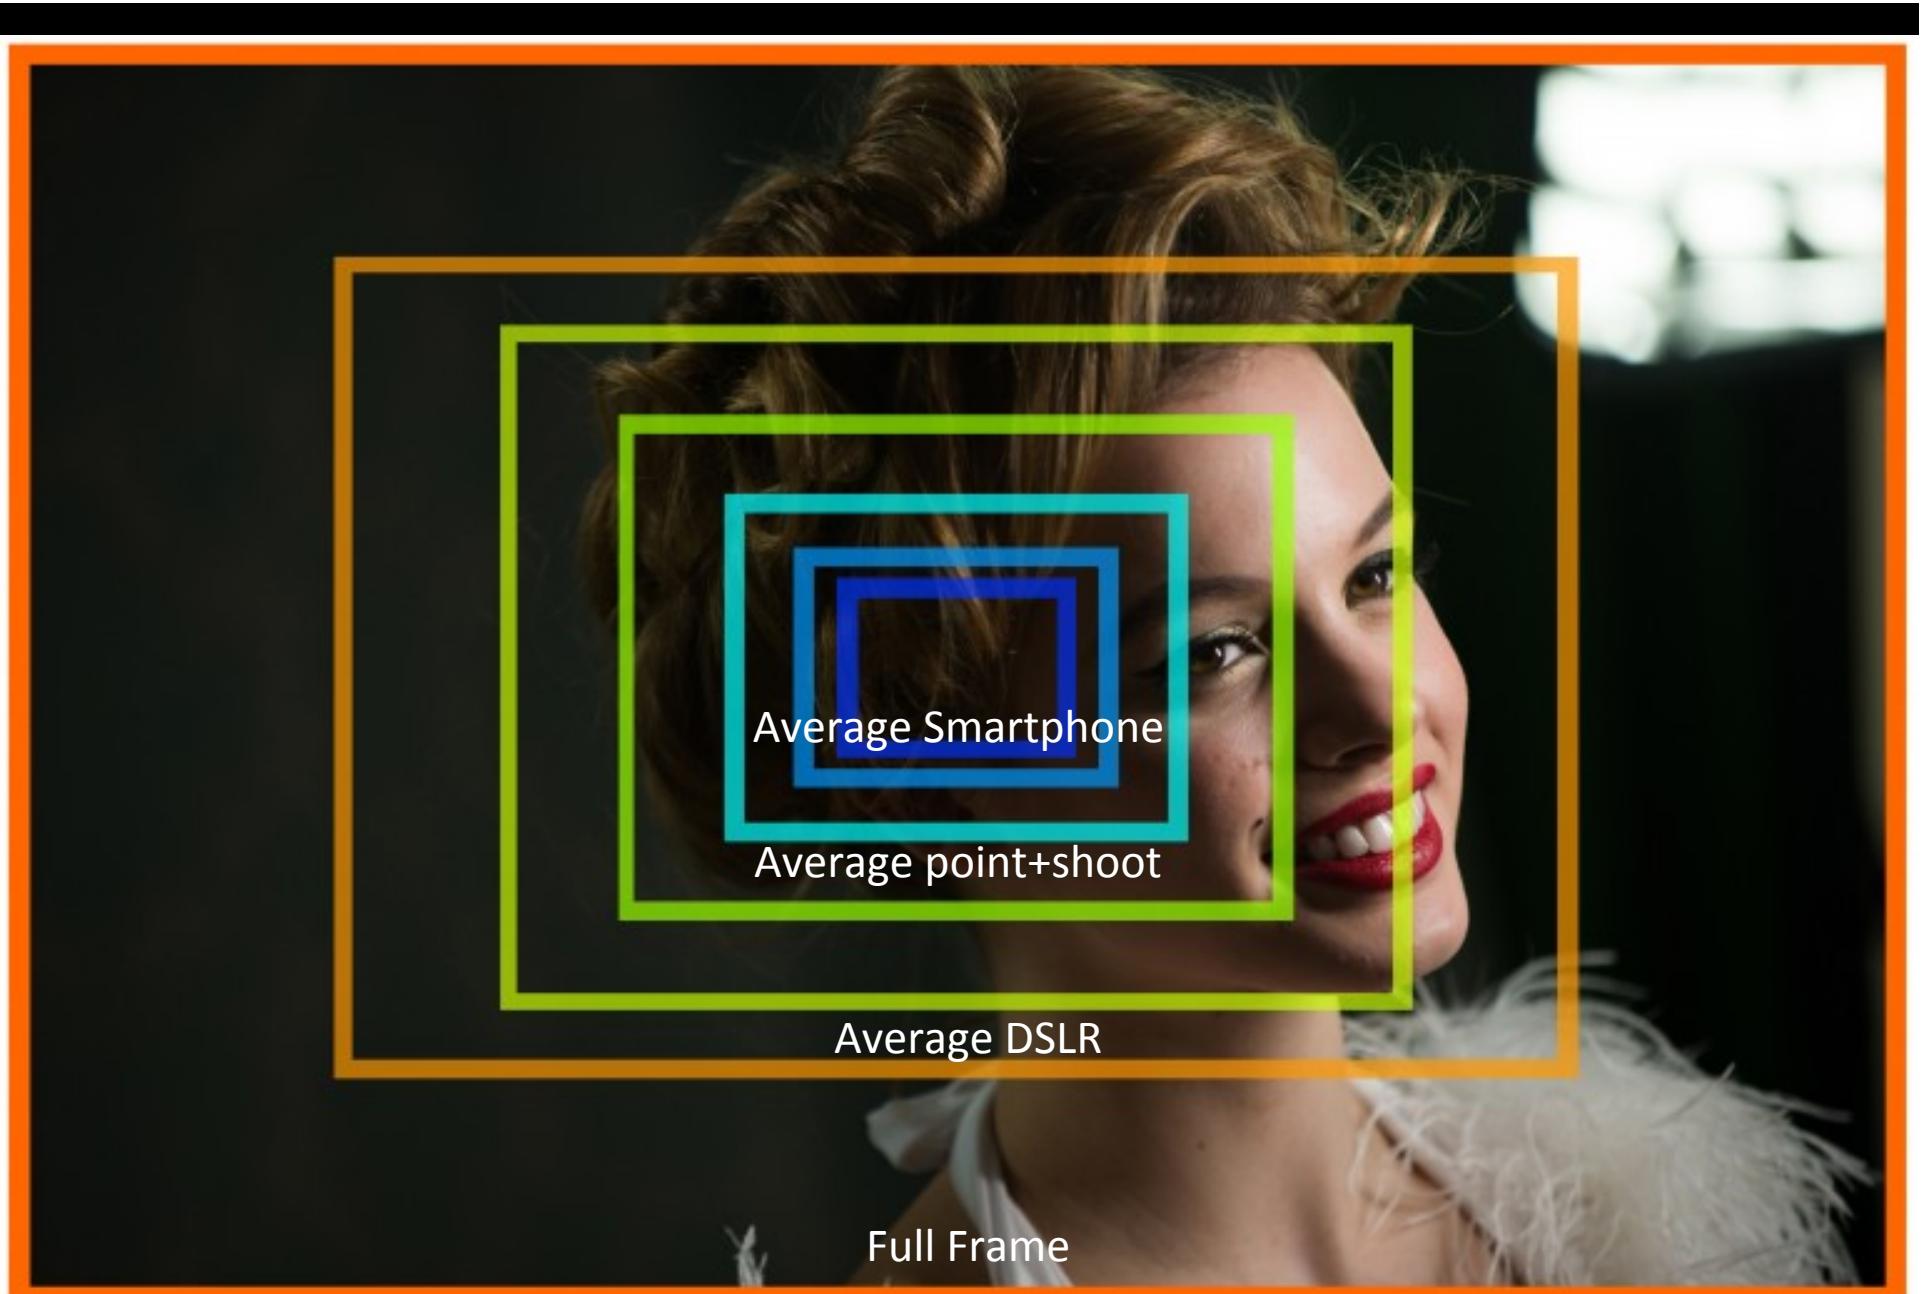


\$5



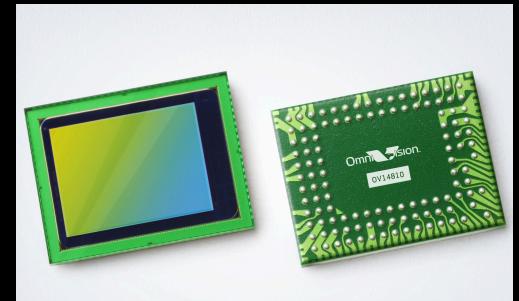
Mobile phone camera

Can these two ever be equally good at taking pictures?



# Where are the opportunities?

- Optics won't get a lot better.
- Pixels can't get much smaller.
  - Limited by optical wavelengths
- Sensor won't get much larger (in mobile)
  - Limited by cost and form factor
- Need:
  - Better Capture Protocols
  - Clever Algorithms



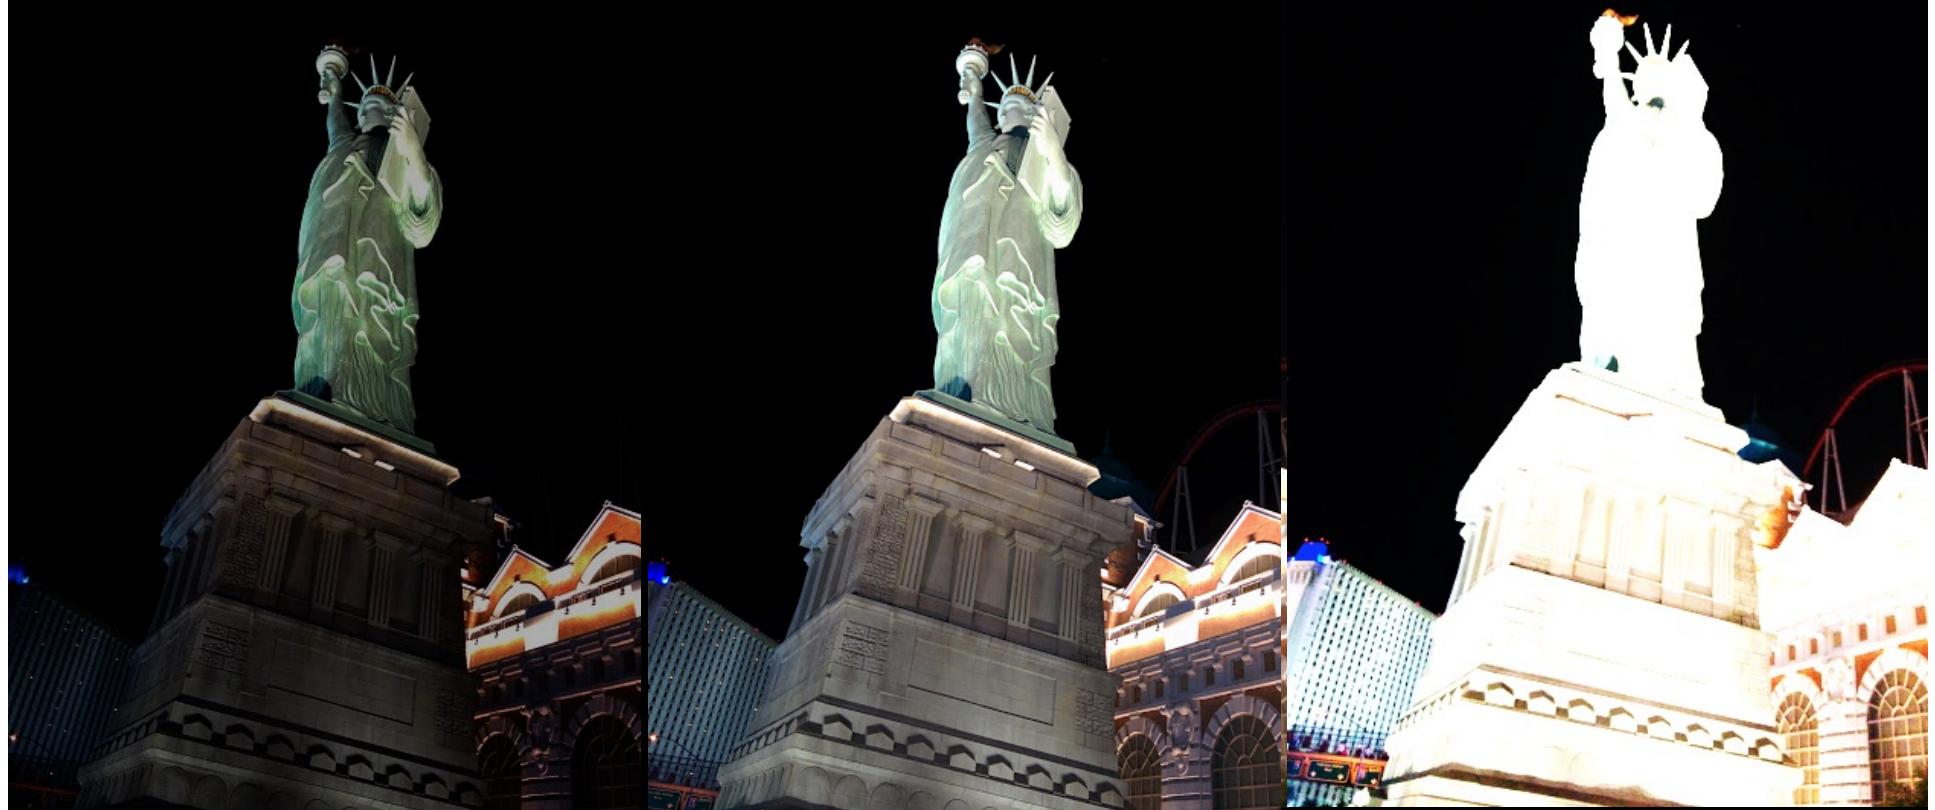
# Burst Photography

Locked Exposure  
& Focus



# HDR Burst Photography

Locked Focus  
Variable Exposure



Google Glass



# Both are from Nexus 5

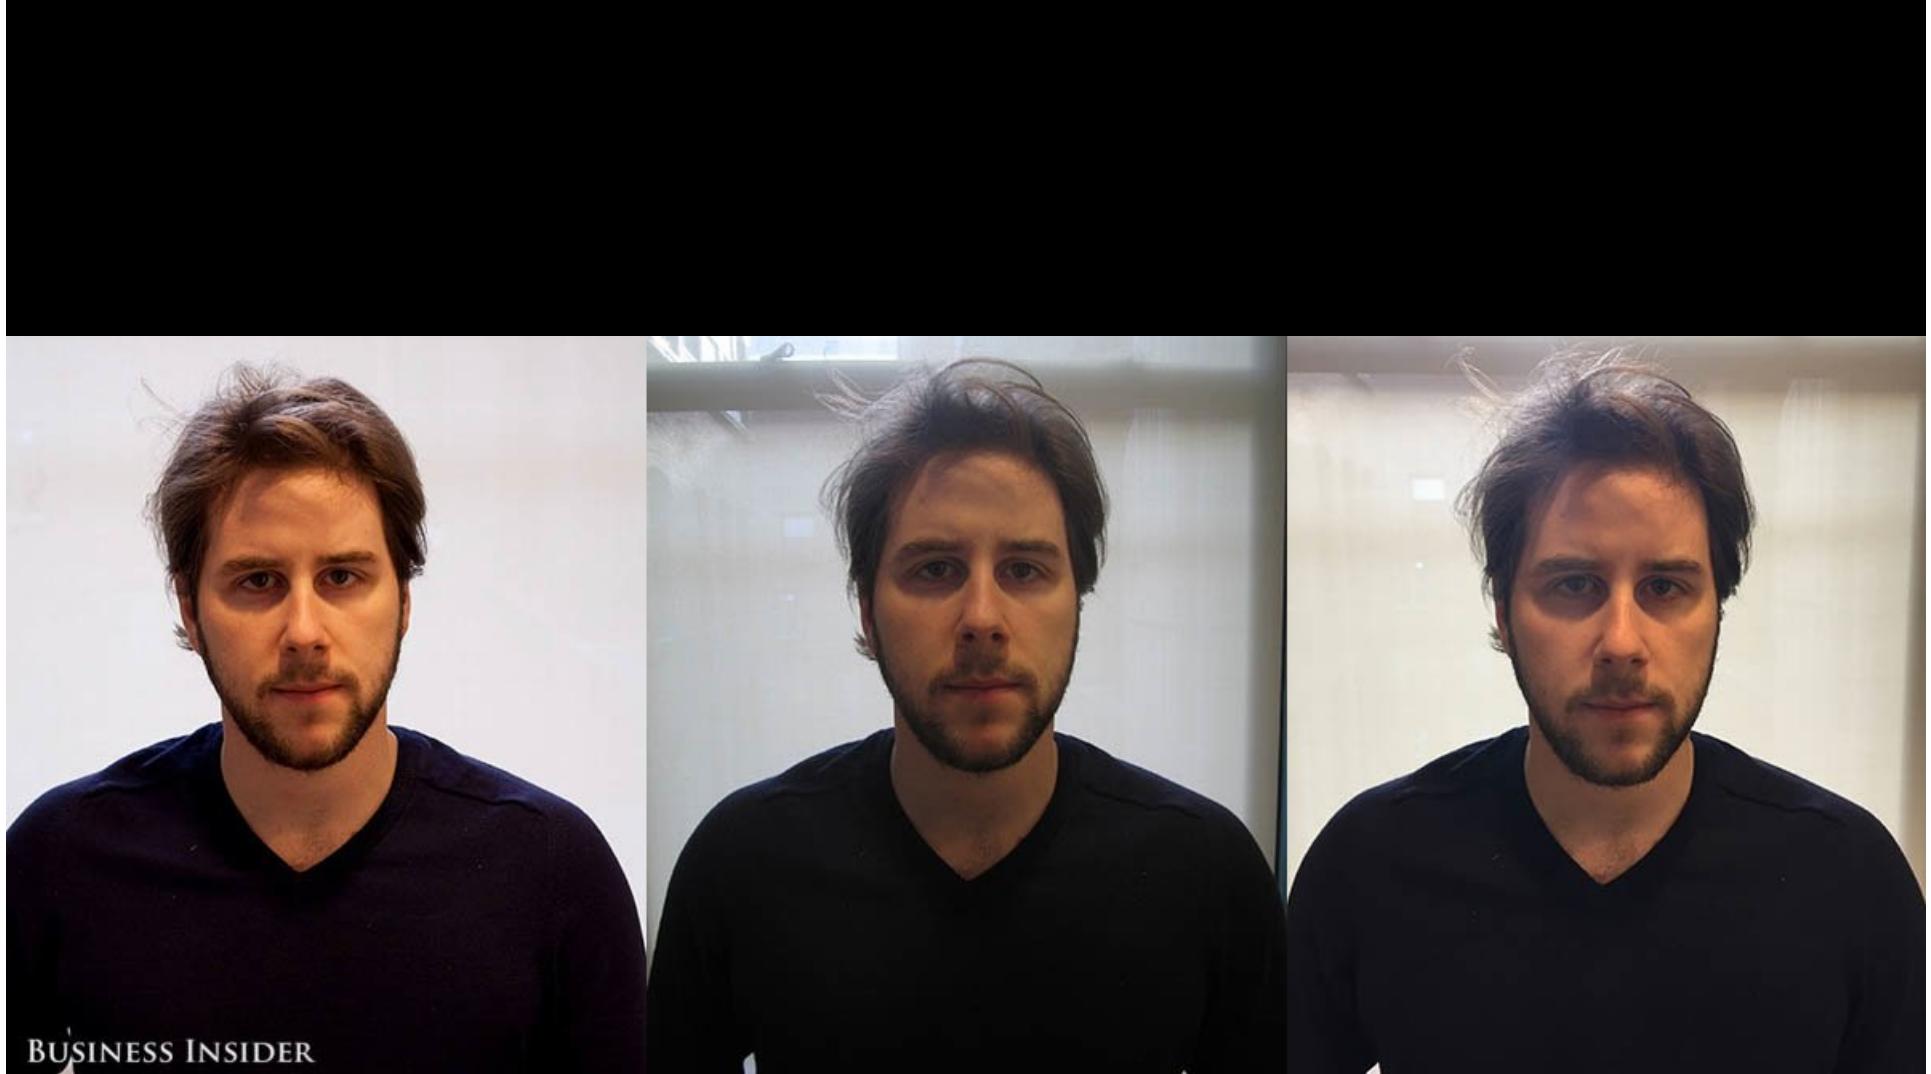


Standard shot



With Lens Blur

Credit: Sascha Haeberling



BUSINESS INSIDER

# Dual Aims of Computational Imaging

## Capture what I see (Digital Photography)

- Take a nice picture. (noise, dynamic range, etc.)
- Make me a better photographer.
- Do it with my simple camera.

## Capture what I can't see (Science, Military)

- See in the dark
- View fast, slow, small, or faint phenomena
- Detect, magnify subtle changes

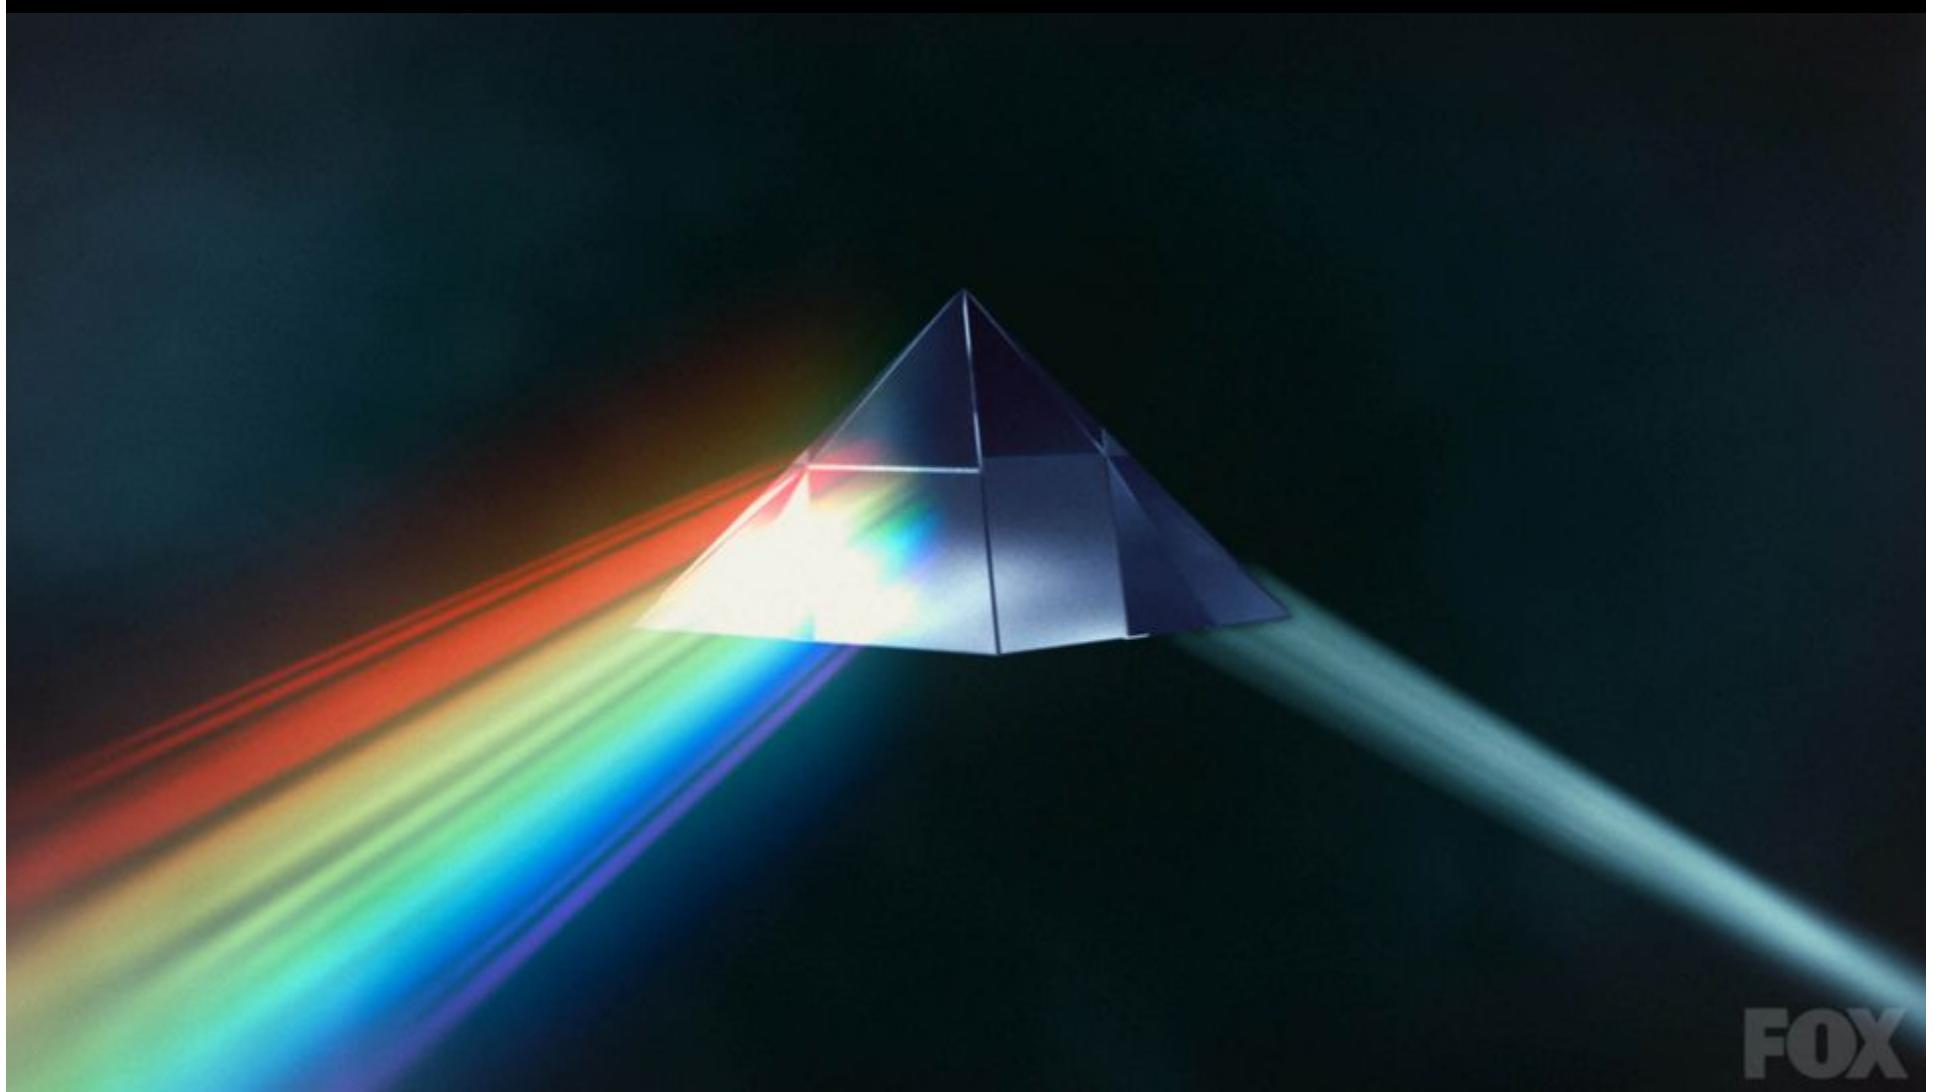
# Yosemite Valley, California

At night



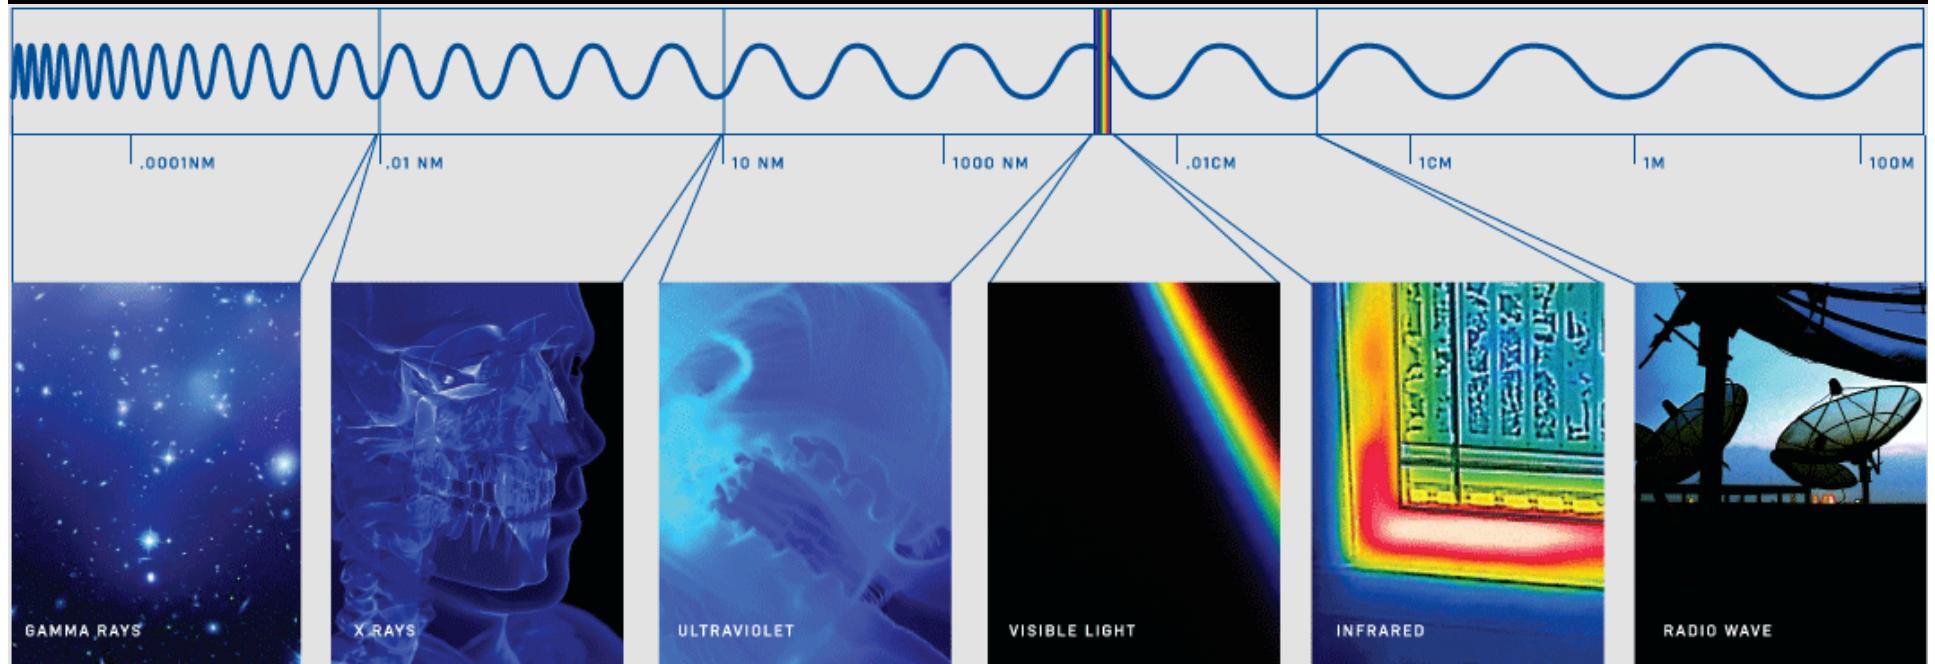
(Jesse Levinson Canon 10D, 28mm f/4, 3 min, ISO 100, 4 image pano)

Can't "see" most of nature that surrounds us

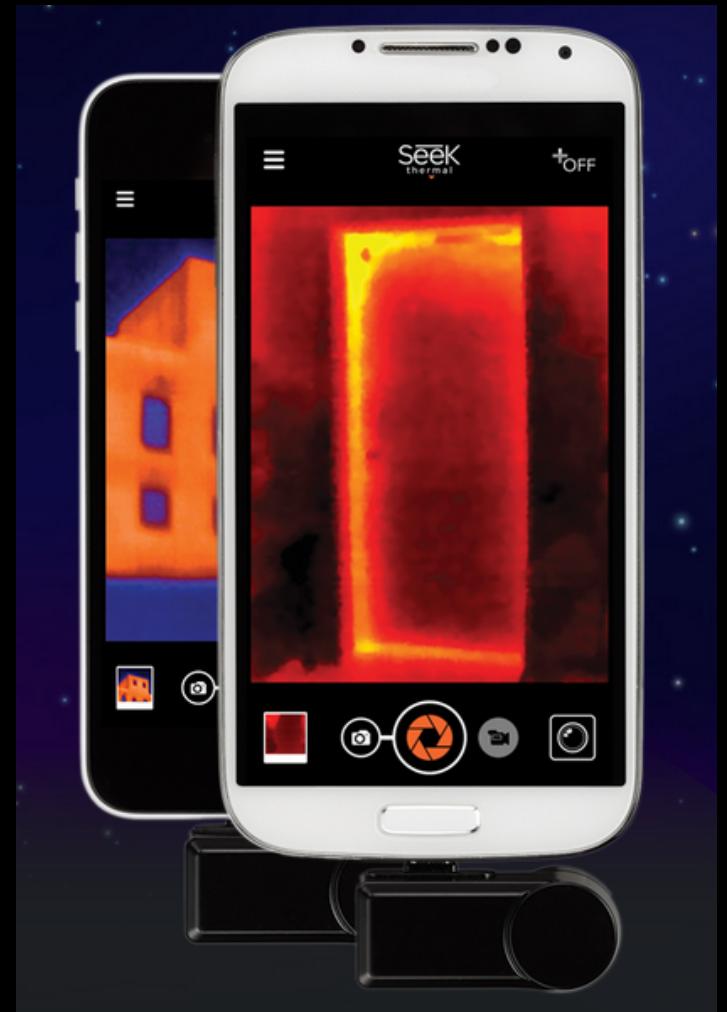


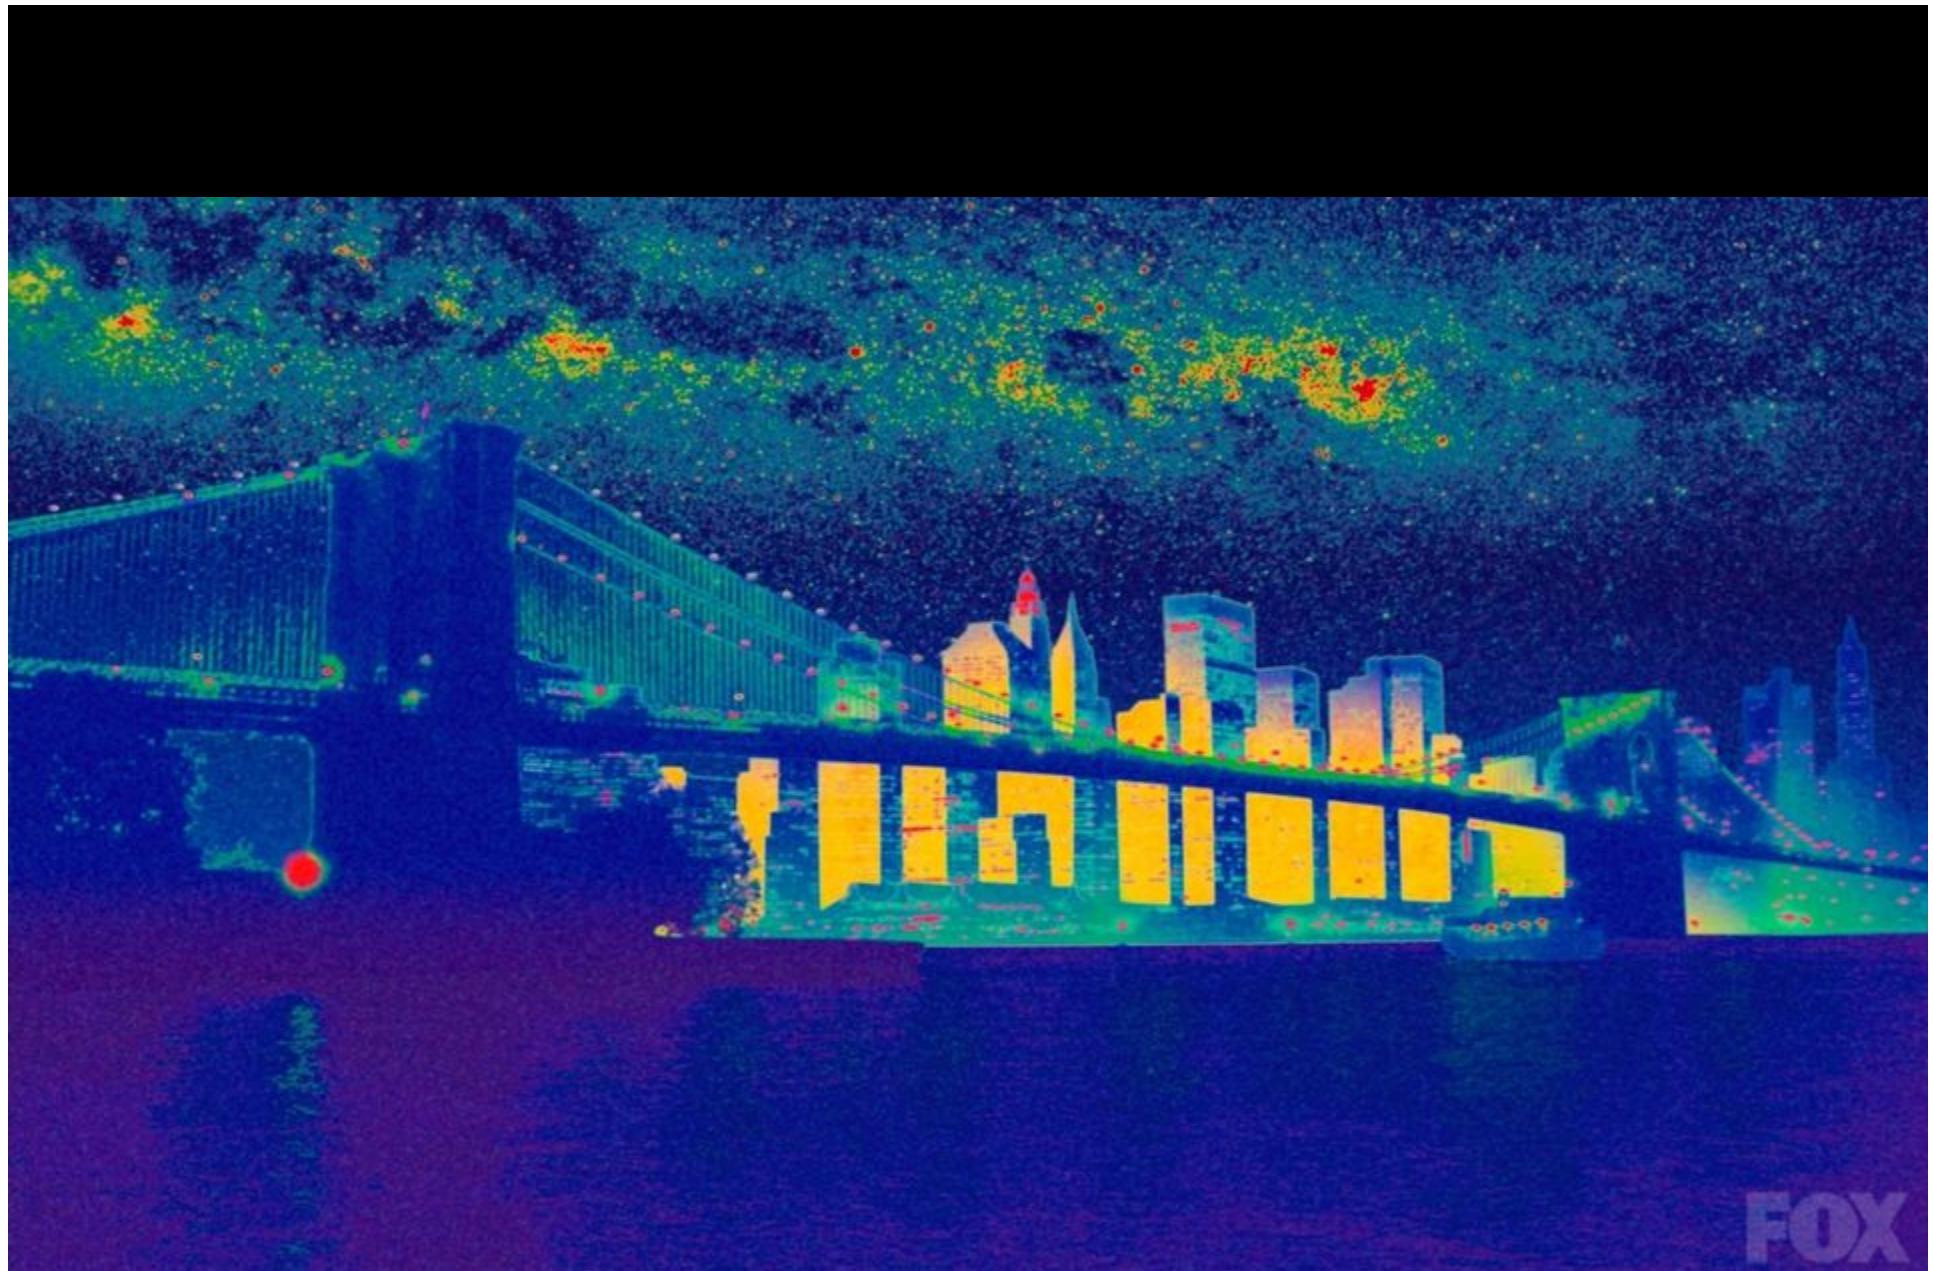
FOX

# Can't "see" most of nature that surrounds us



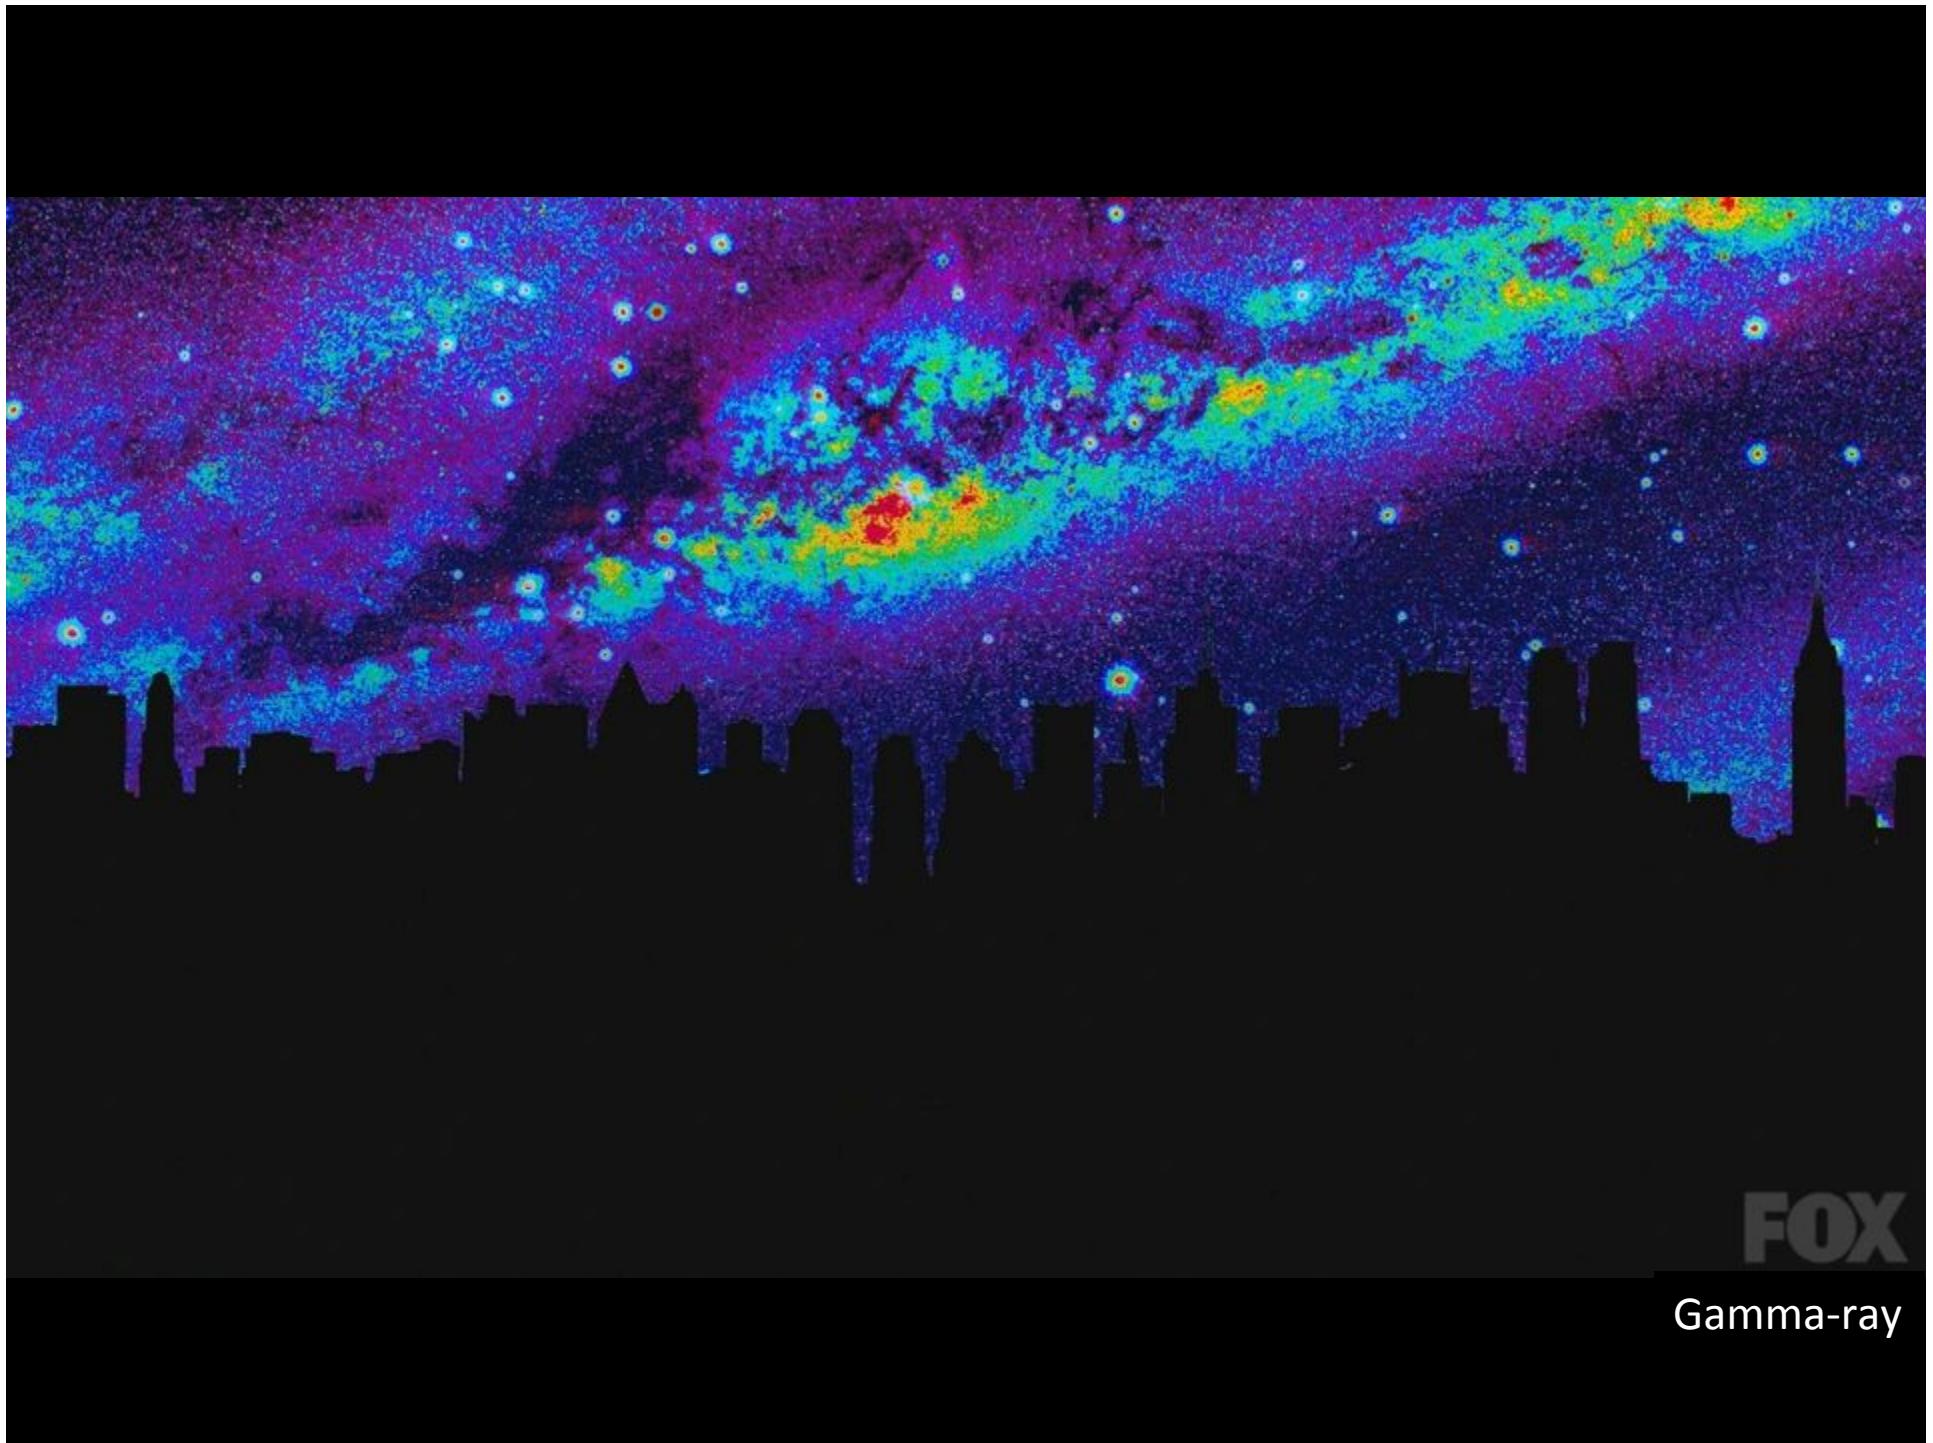
# Infrared is already here





FOX

Infrared



**FOX**

Gamma-ray



**FOX**

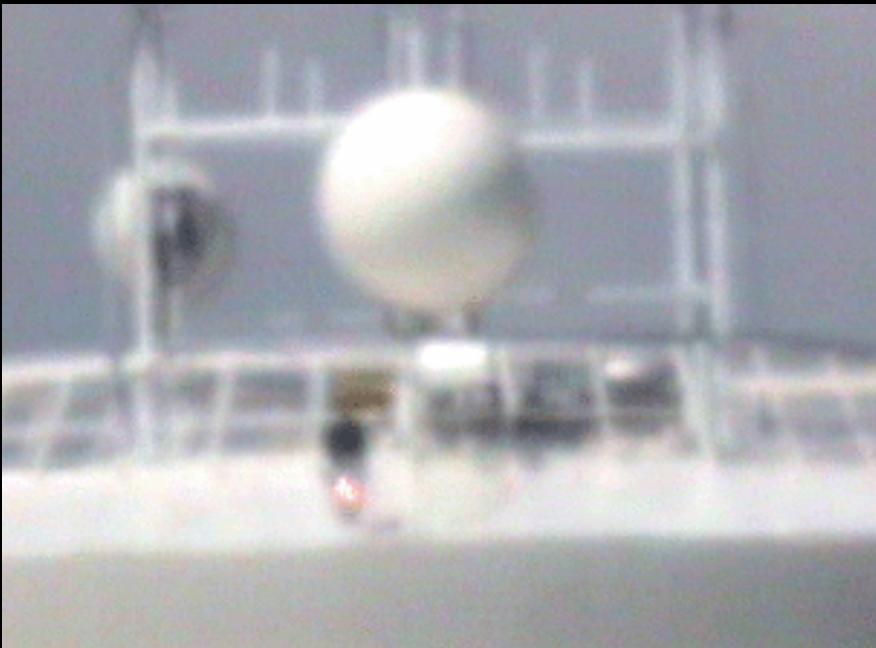
X-ray

# Seeing Far Away



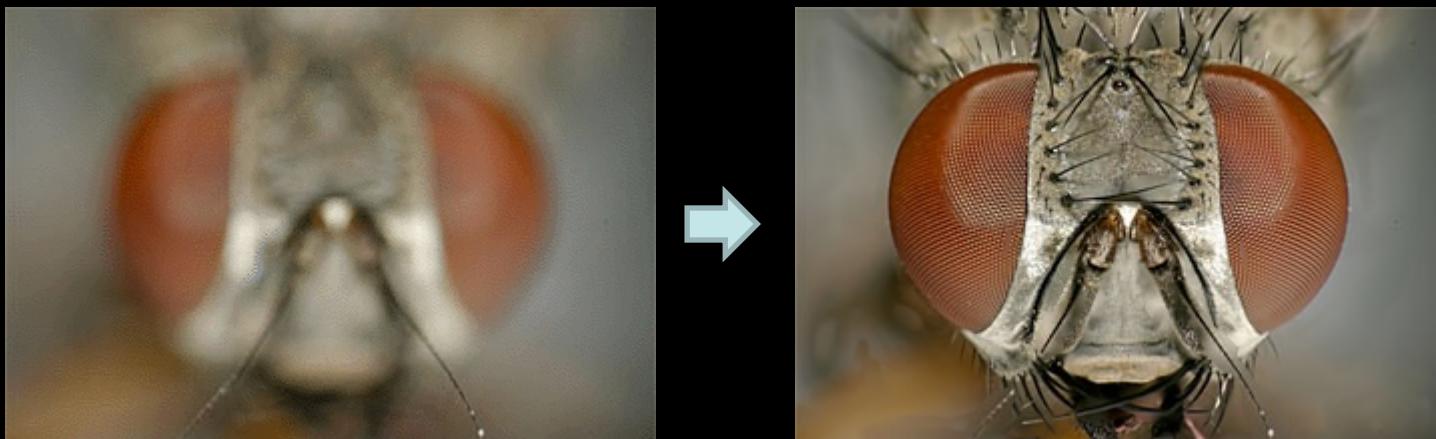
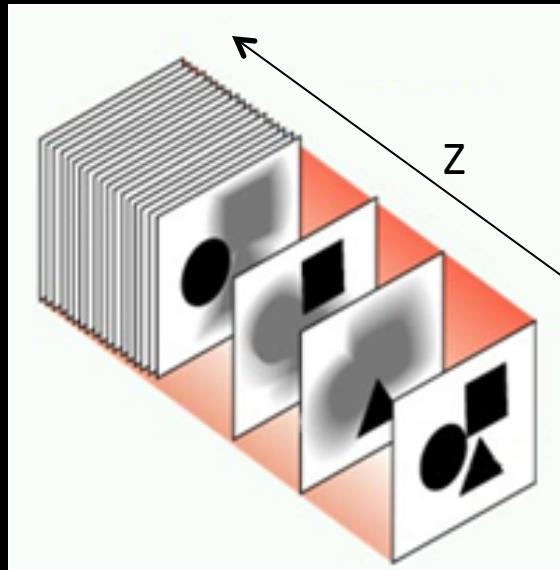
Credit : Florian Kainz

# Seeing Far Away



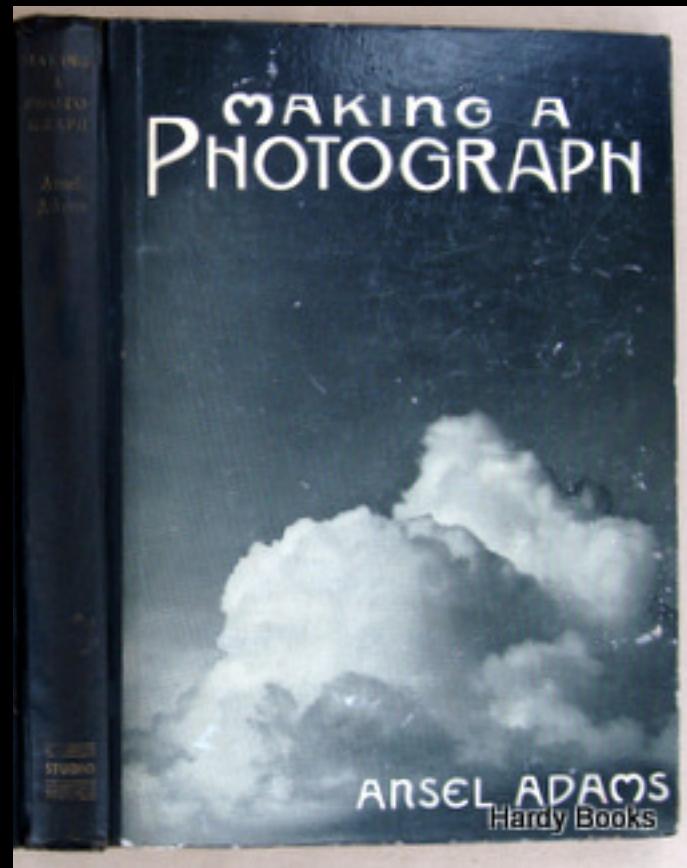
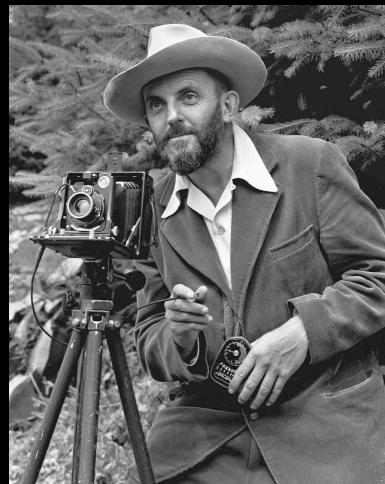
Top part of a water tower imaged at a (horizontal) distance of 2.4 km

# Seeing Small Things



X. Zhu, P. Milanfar, "Removing Atmospheric Turbulence", PAMI, 2012

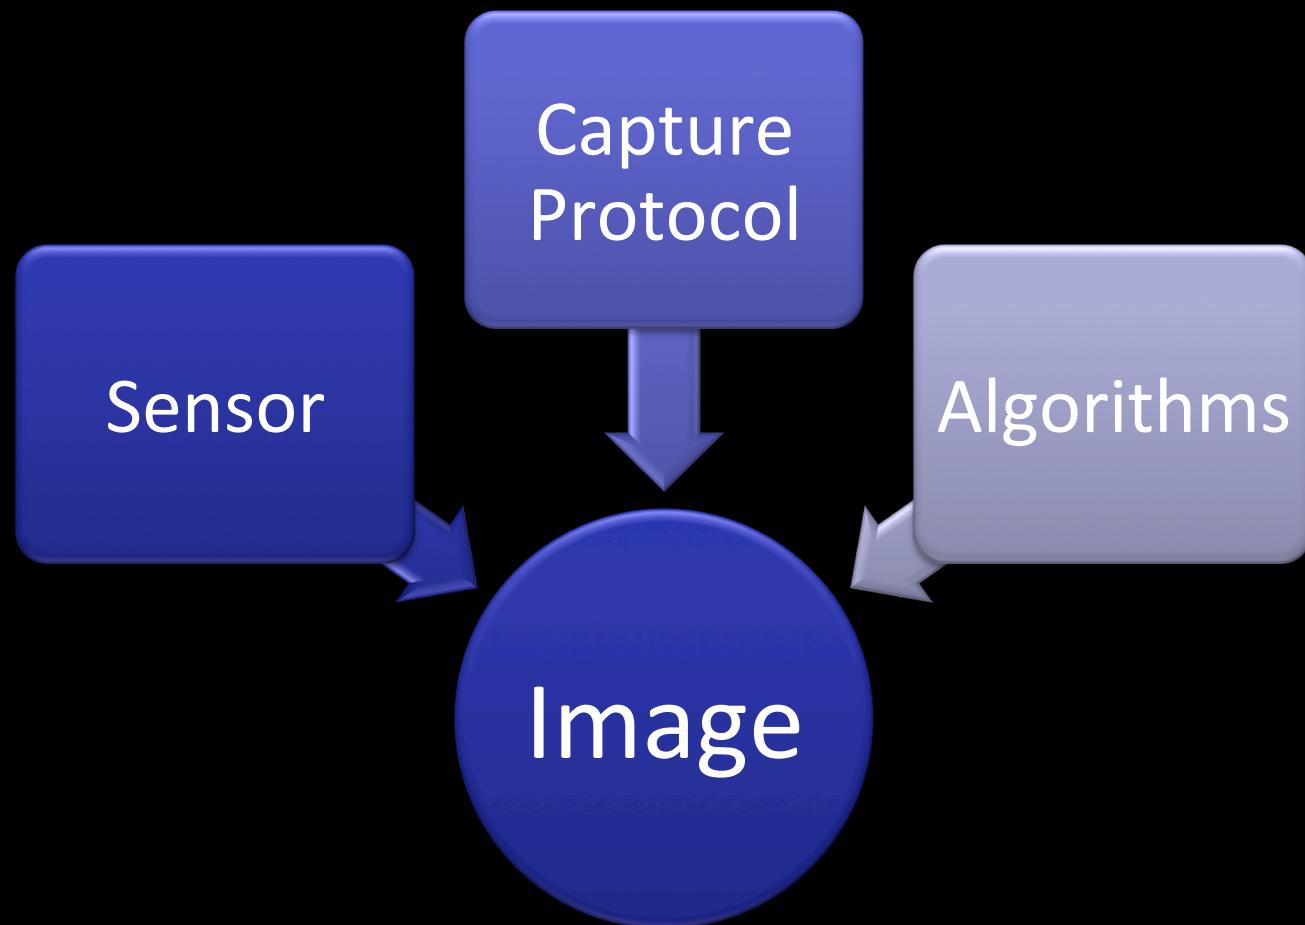
# How Do We Make a High Quality Image?



Ansel Adams' first book, published 1935

# How Do We Make a High Quality Image?

- Integration of Sensor and Computation plays a key role in the image formation process.



# Computational Photography on Android

**devcamera.org**

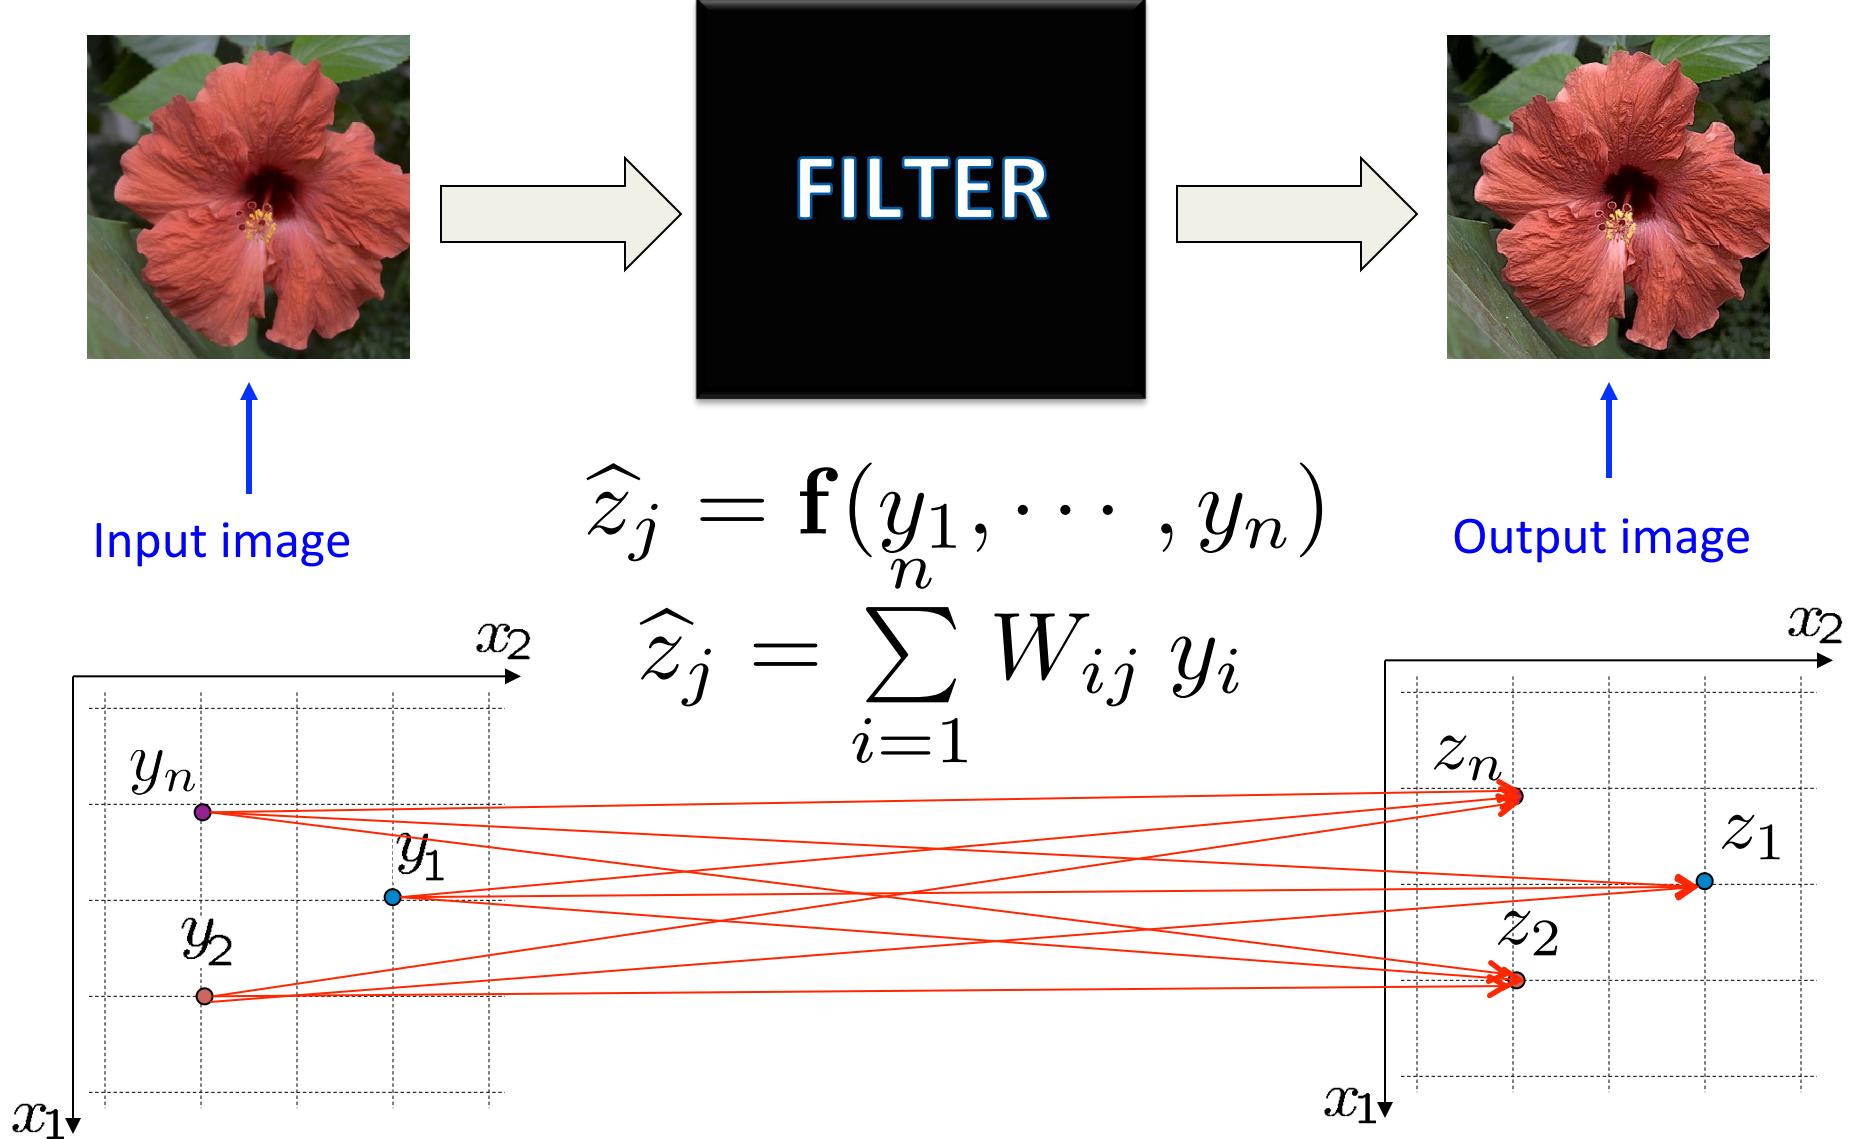
APK for parameterized image sequence  
capture using Android 5.0 + (Lollipop)

<https://youtu.be/92fgcUNCHic?t=29m56s>

# Part 2:

## A general filtering framework

# The Black Box

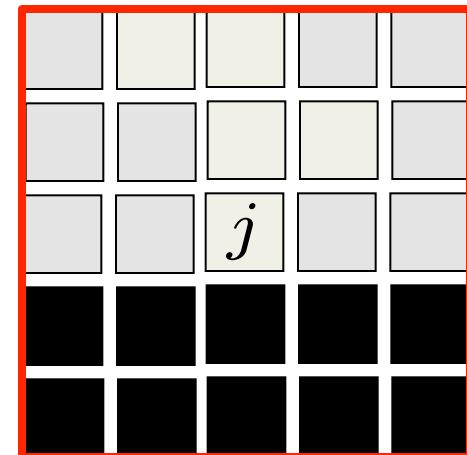


# How to design the black box

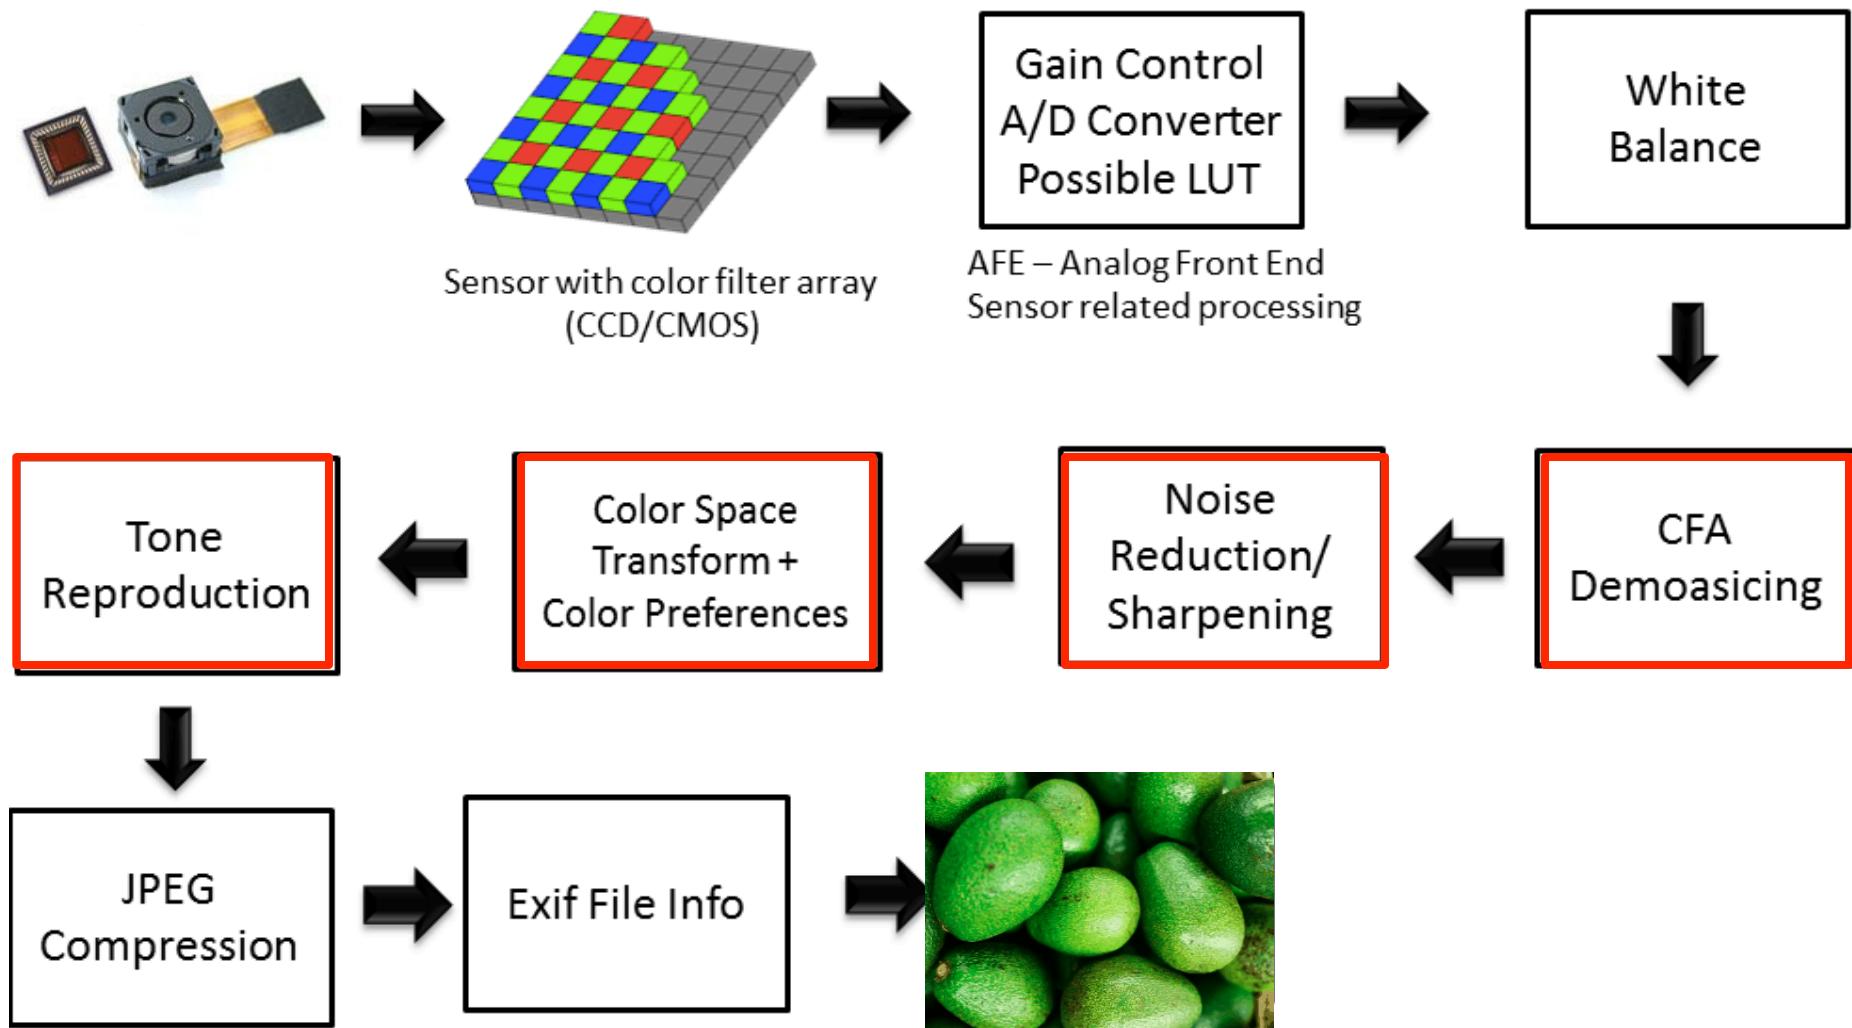
$$\hat{z}_j = \sum_{i=1}^n W_{ij} y_i$$



$y_i \in$

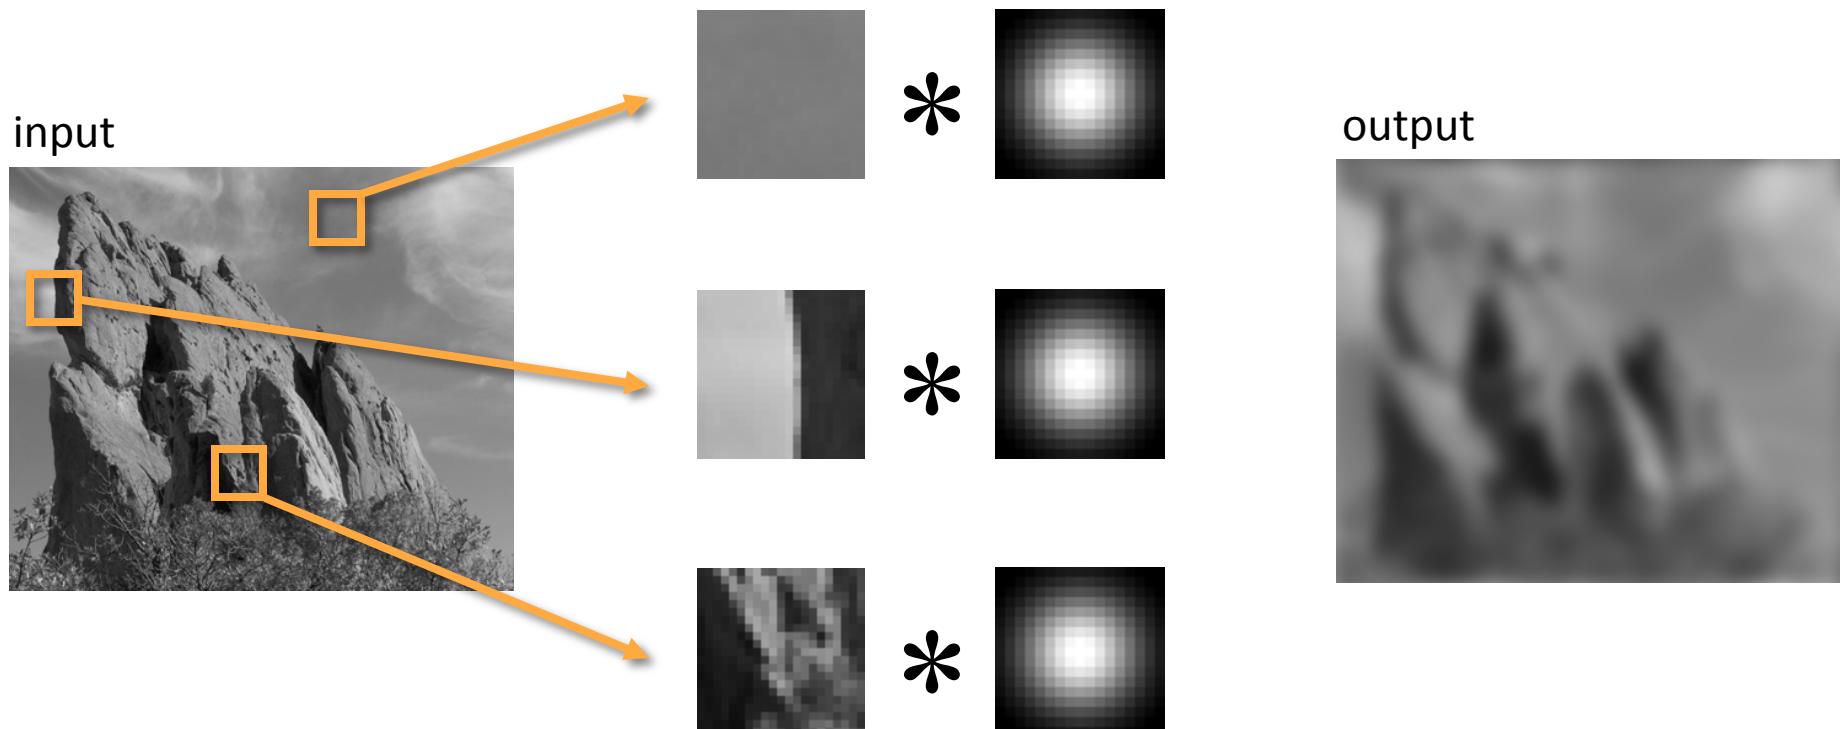


# Model for Many Parts of Pipeline



# Same weights everywhere

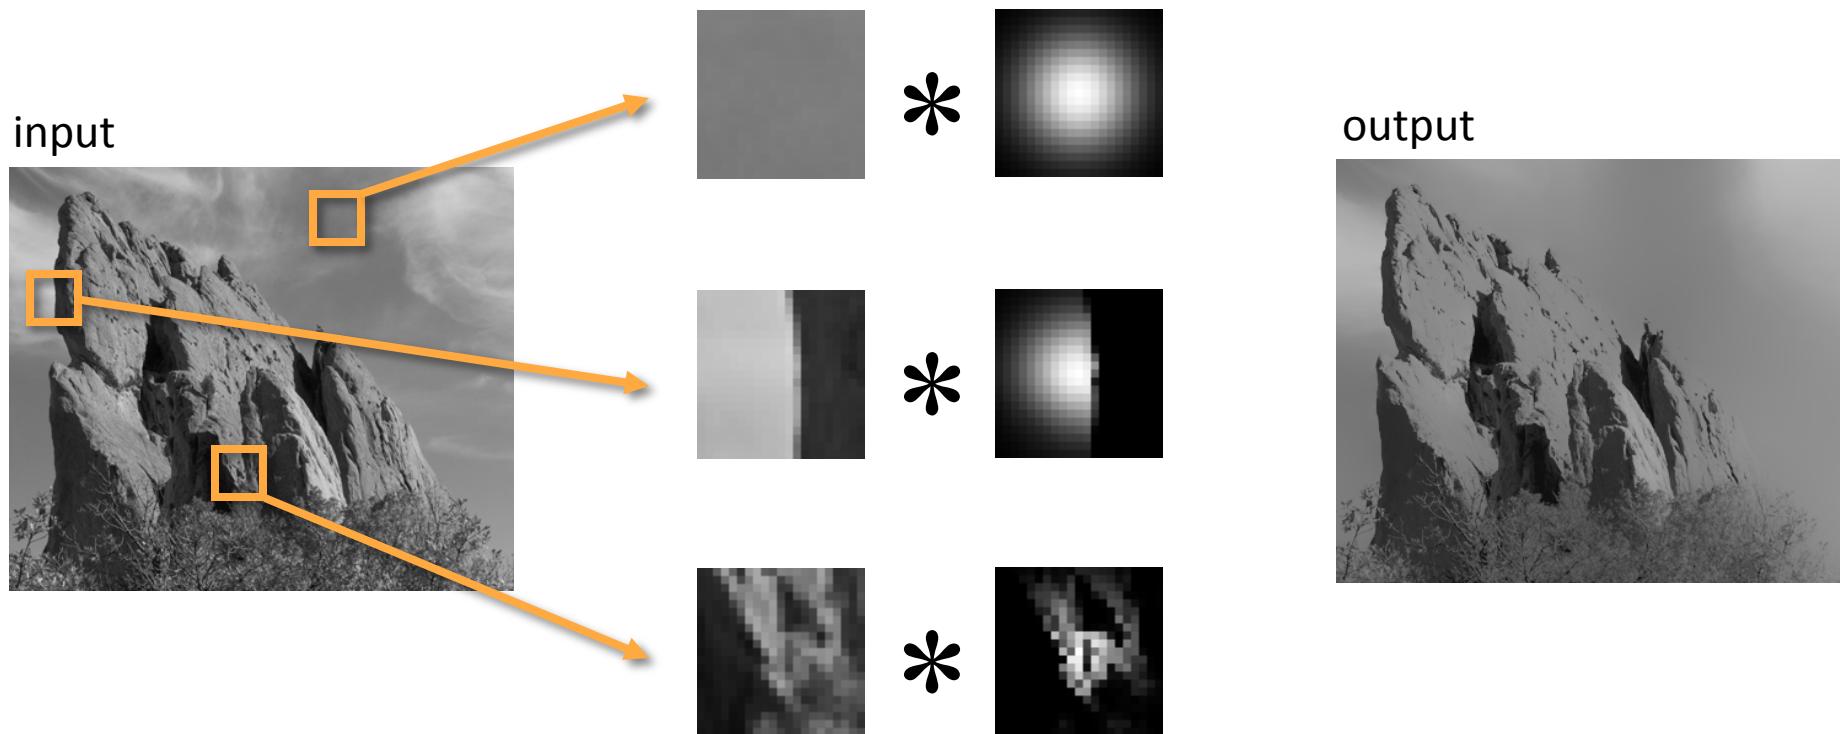
$$\hat{z}_j = \sum_{i=1}^n W_i y_i$$



Same Gaussian kernel everywhere.

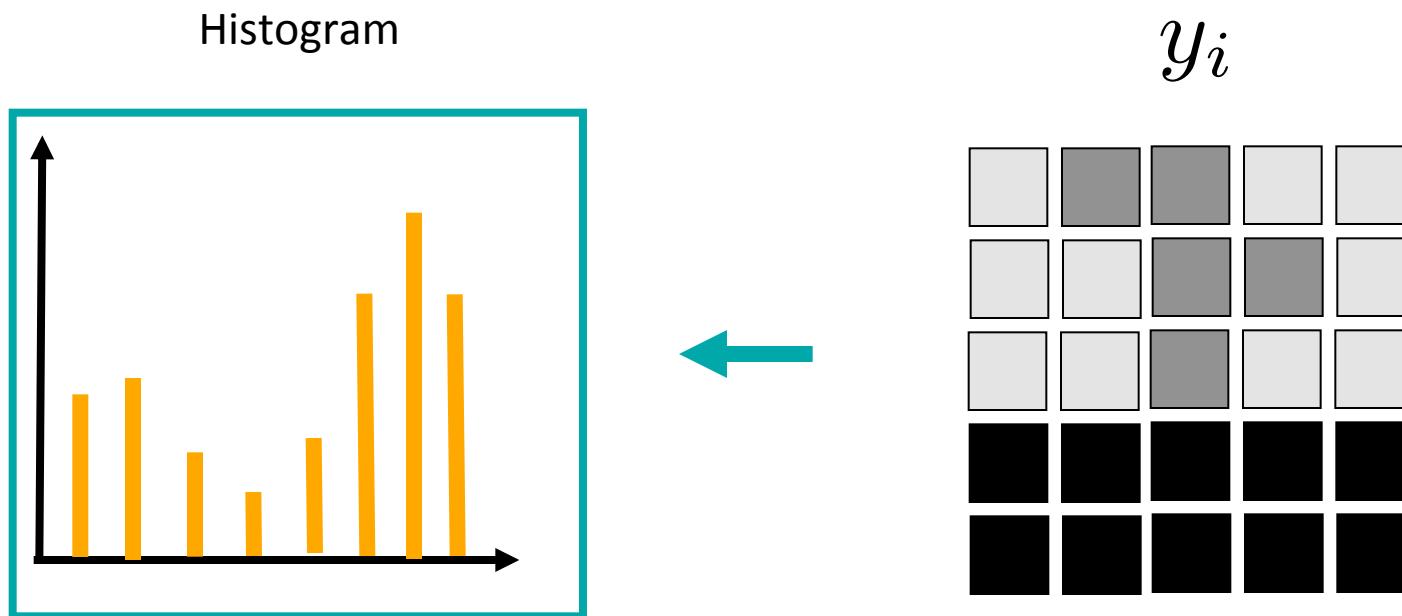
# Data-adaptive weights

$$\hat{z}_j = \sum_{i=1}^n W_{ij} y_i$$



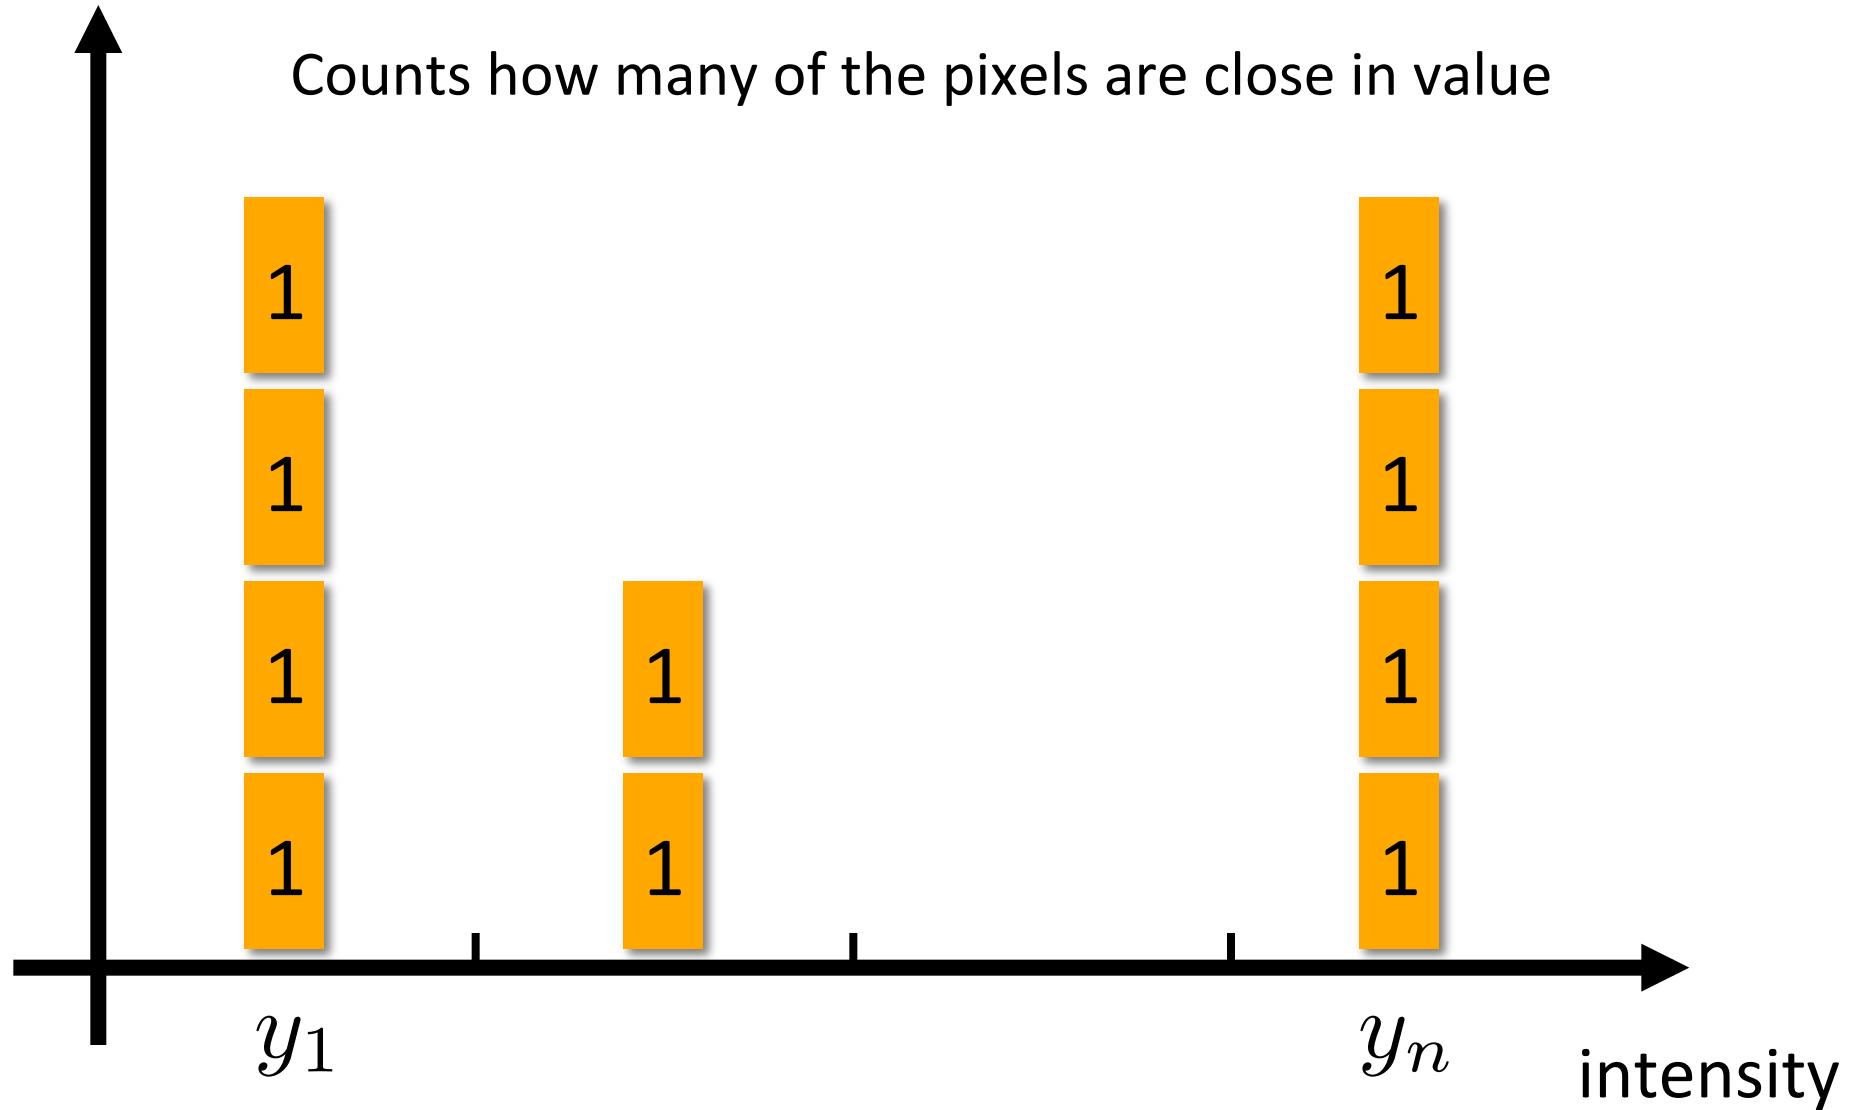
The kernel shape depends on the image content.

# Method 1: Consider the histogram



Just for illustration – we don't really make the histogram.

# pixels



# pixels

Counts how many of the pixels are close in value to  $y_j$

$$H(y_j) = \sum_{i=1}^n K(y_i - y_j)$$

$K(\cdot)$

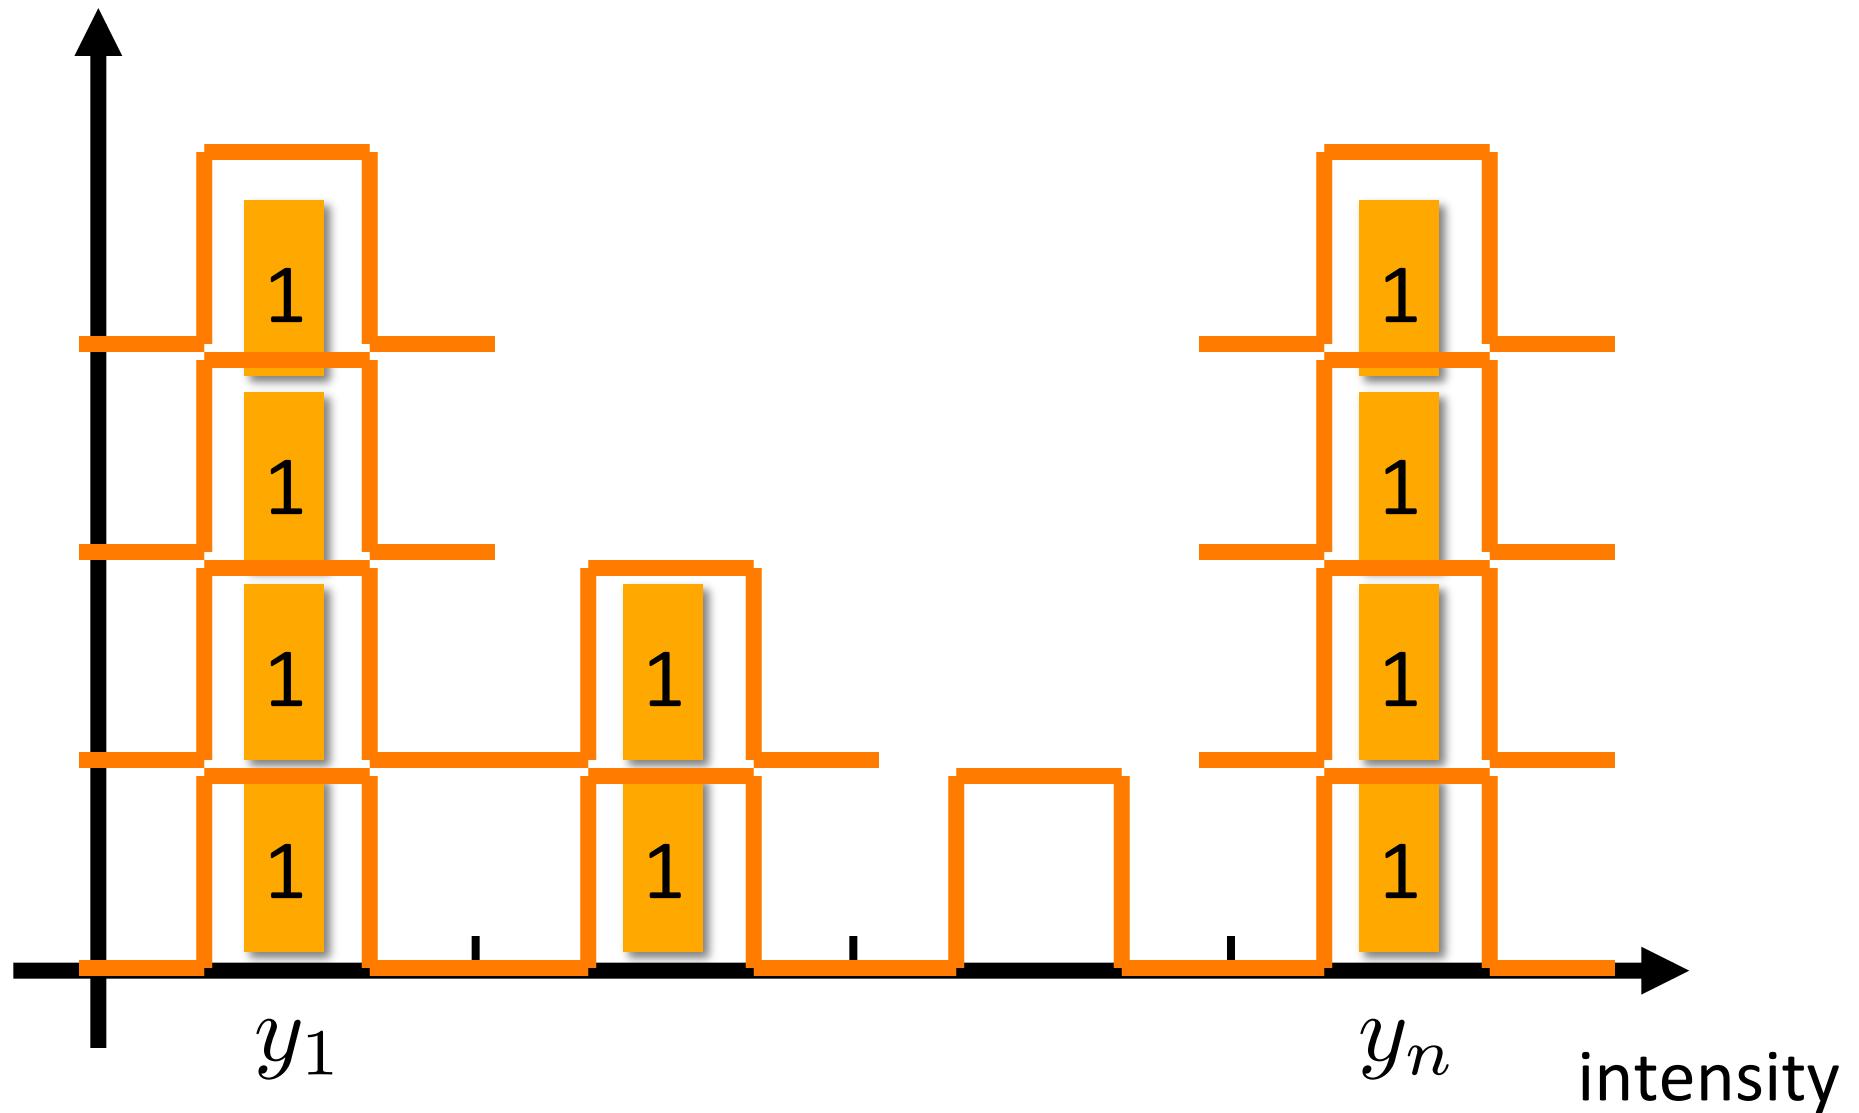
1



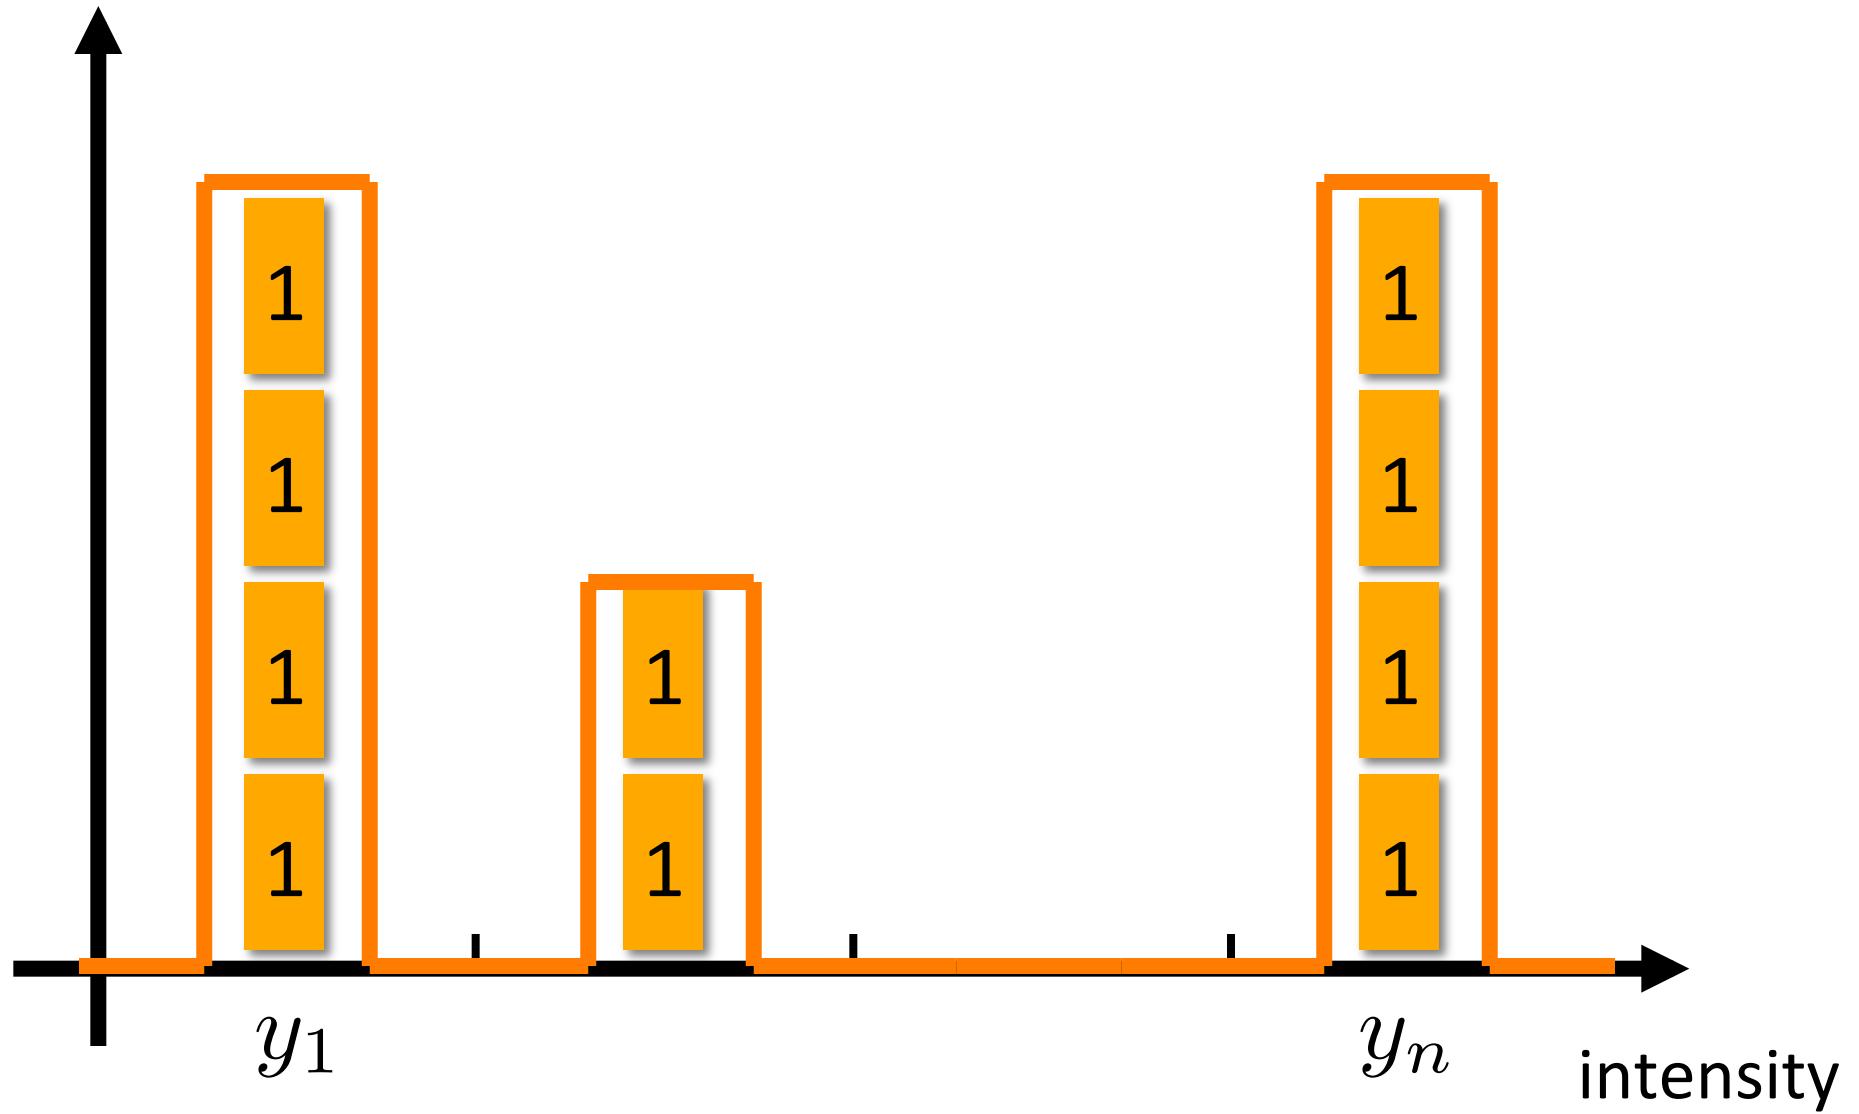
$y_j$

intensity

# pixels



# pixels



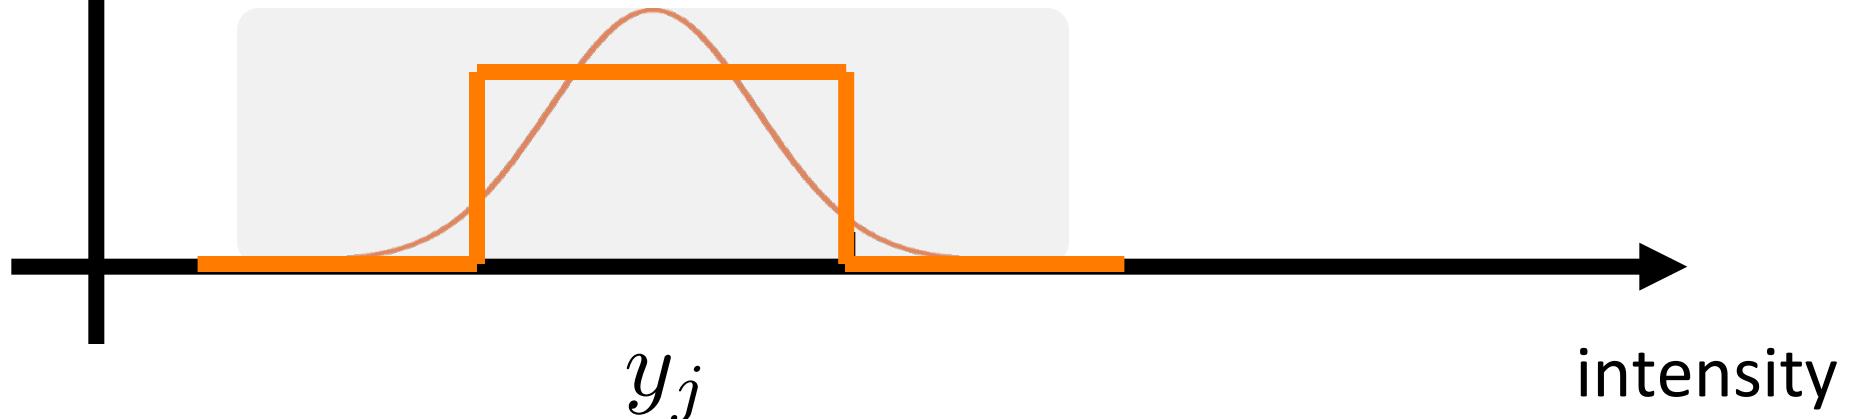
# pixels

# Why only the box function?

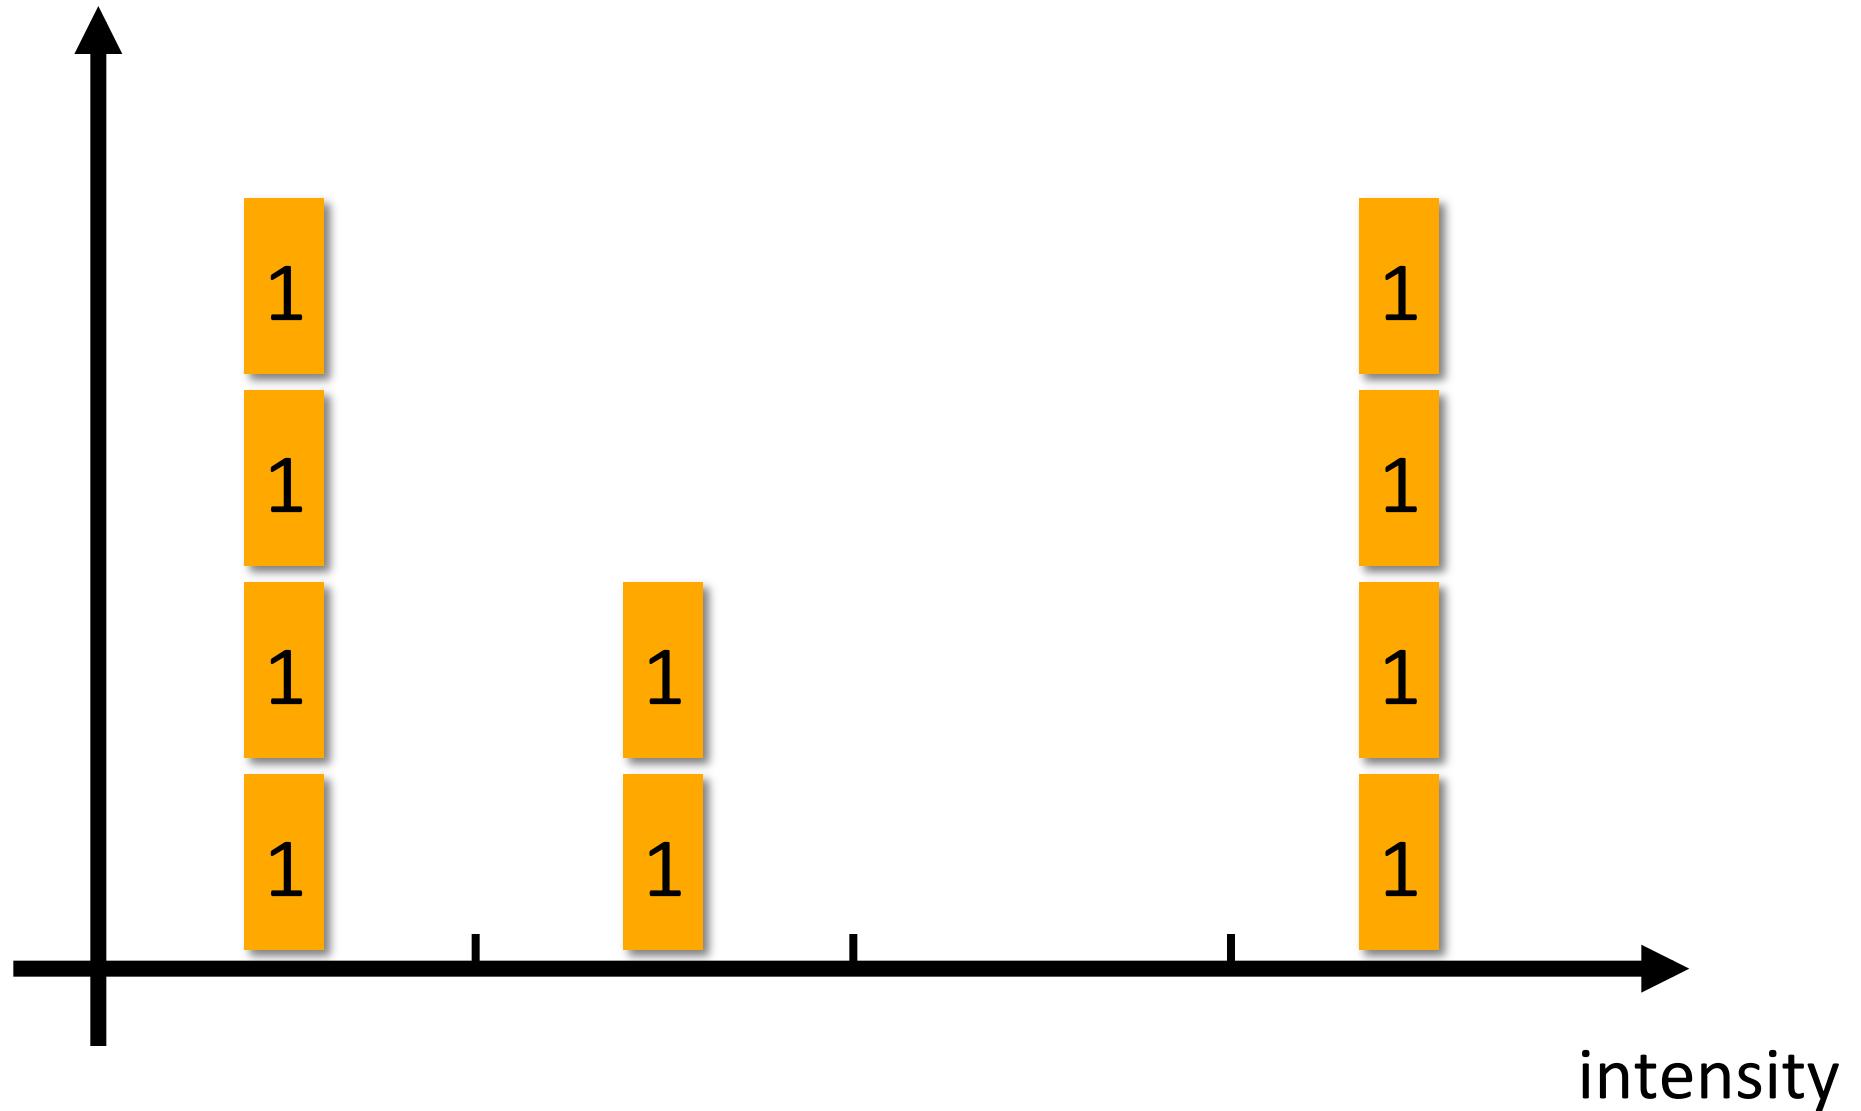
**Soft** count of how many of the pixels are close in value to  $y_j$

$$H(y_j) = \sum_{i=1}^n K(y_i - y_j)$$

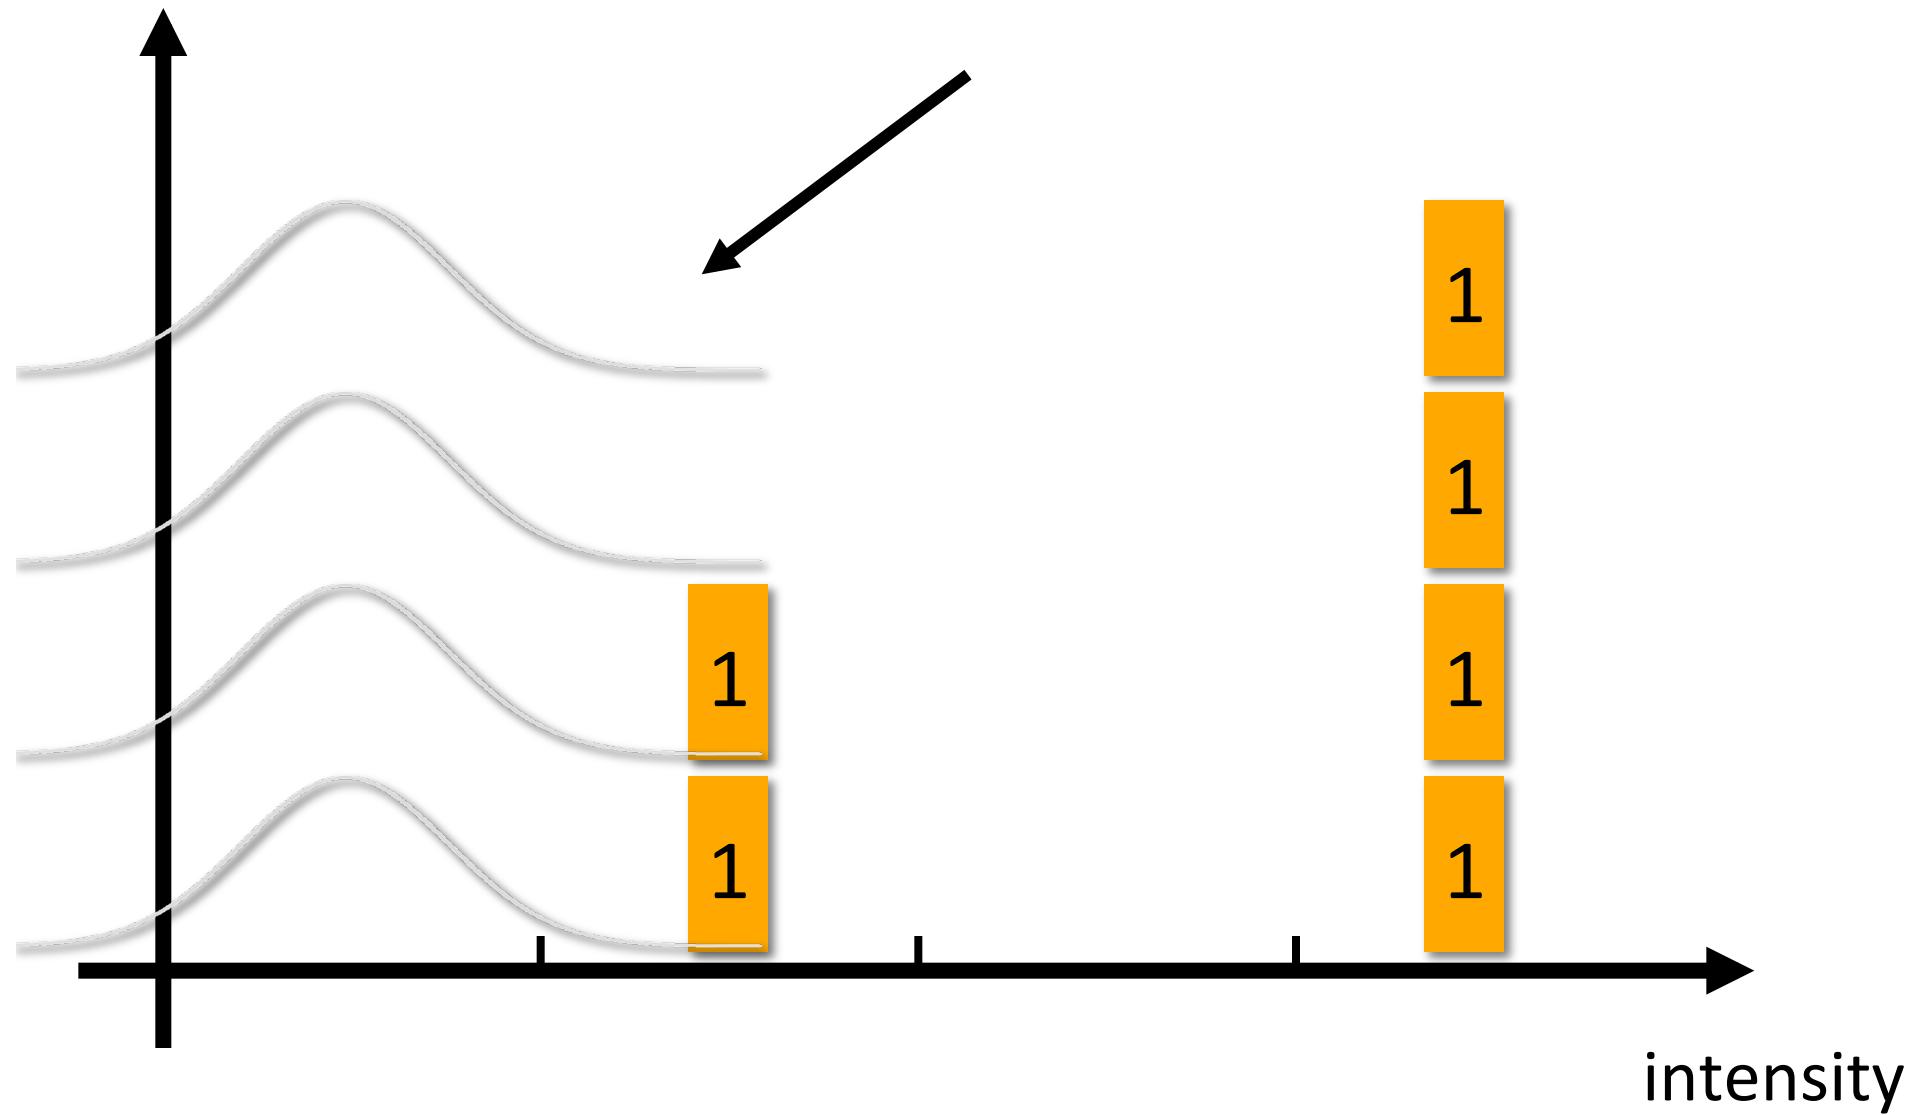
$$K(y_i - y_j) = \exp(-|y_i - y_j|^2/\sigma^2)$$

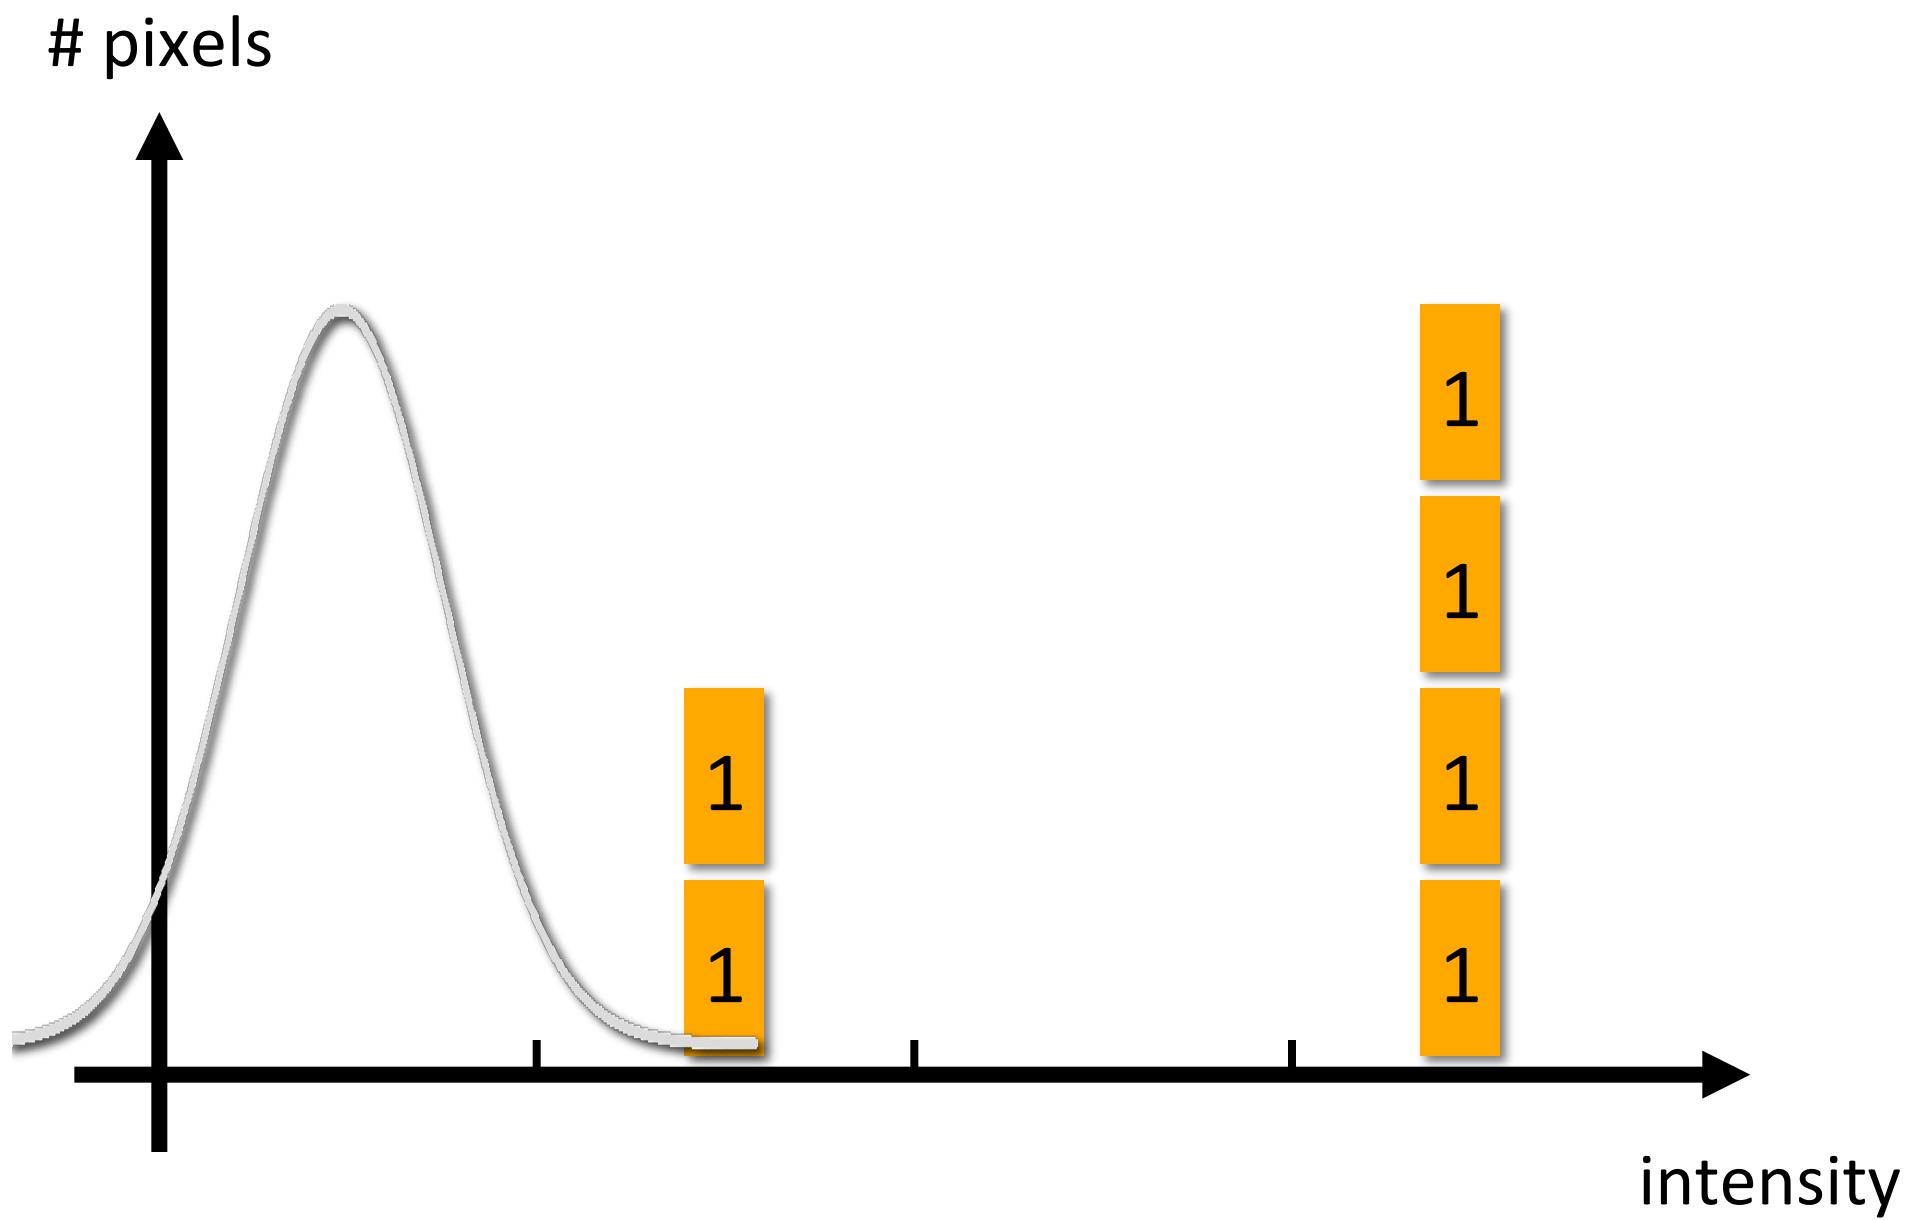


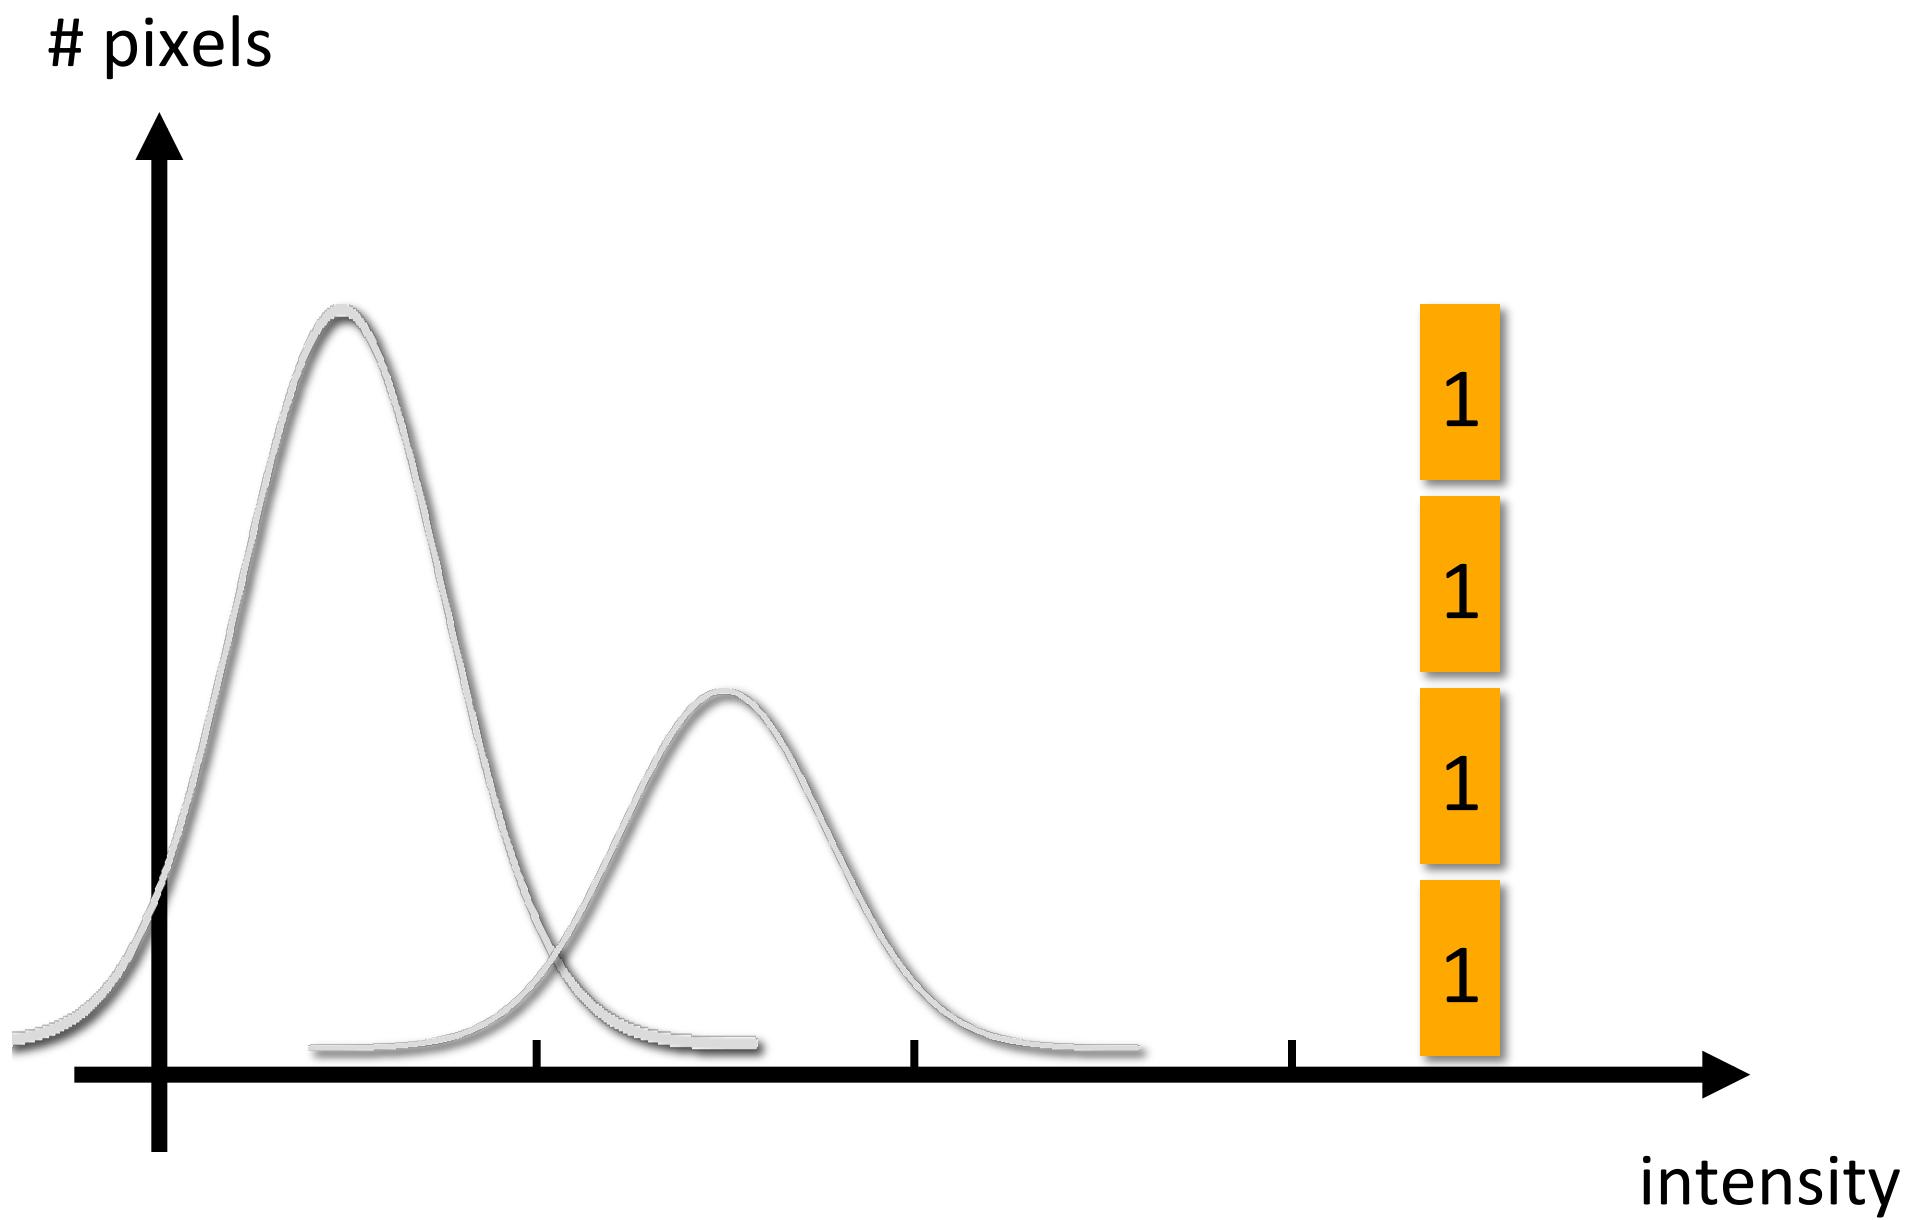
# pixels

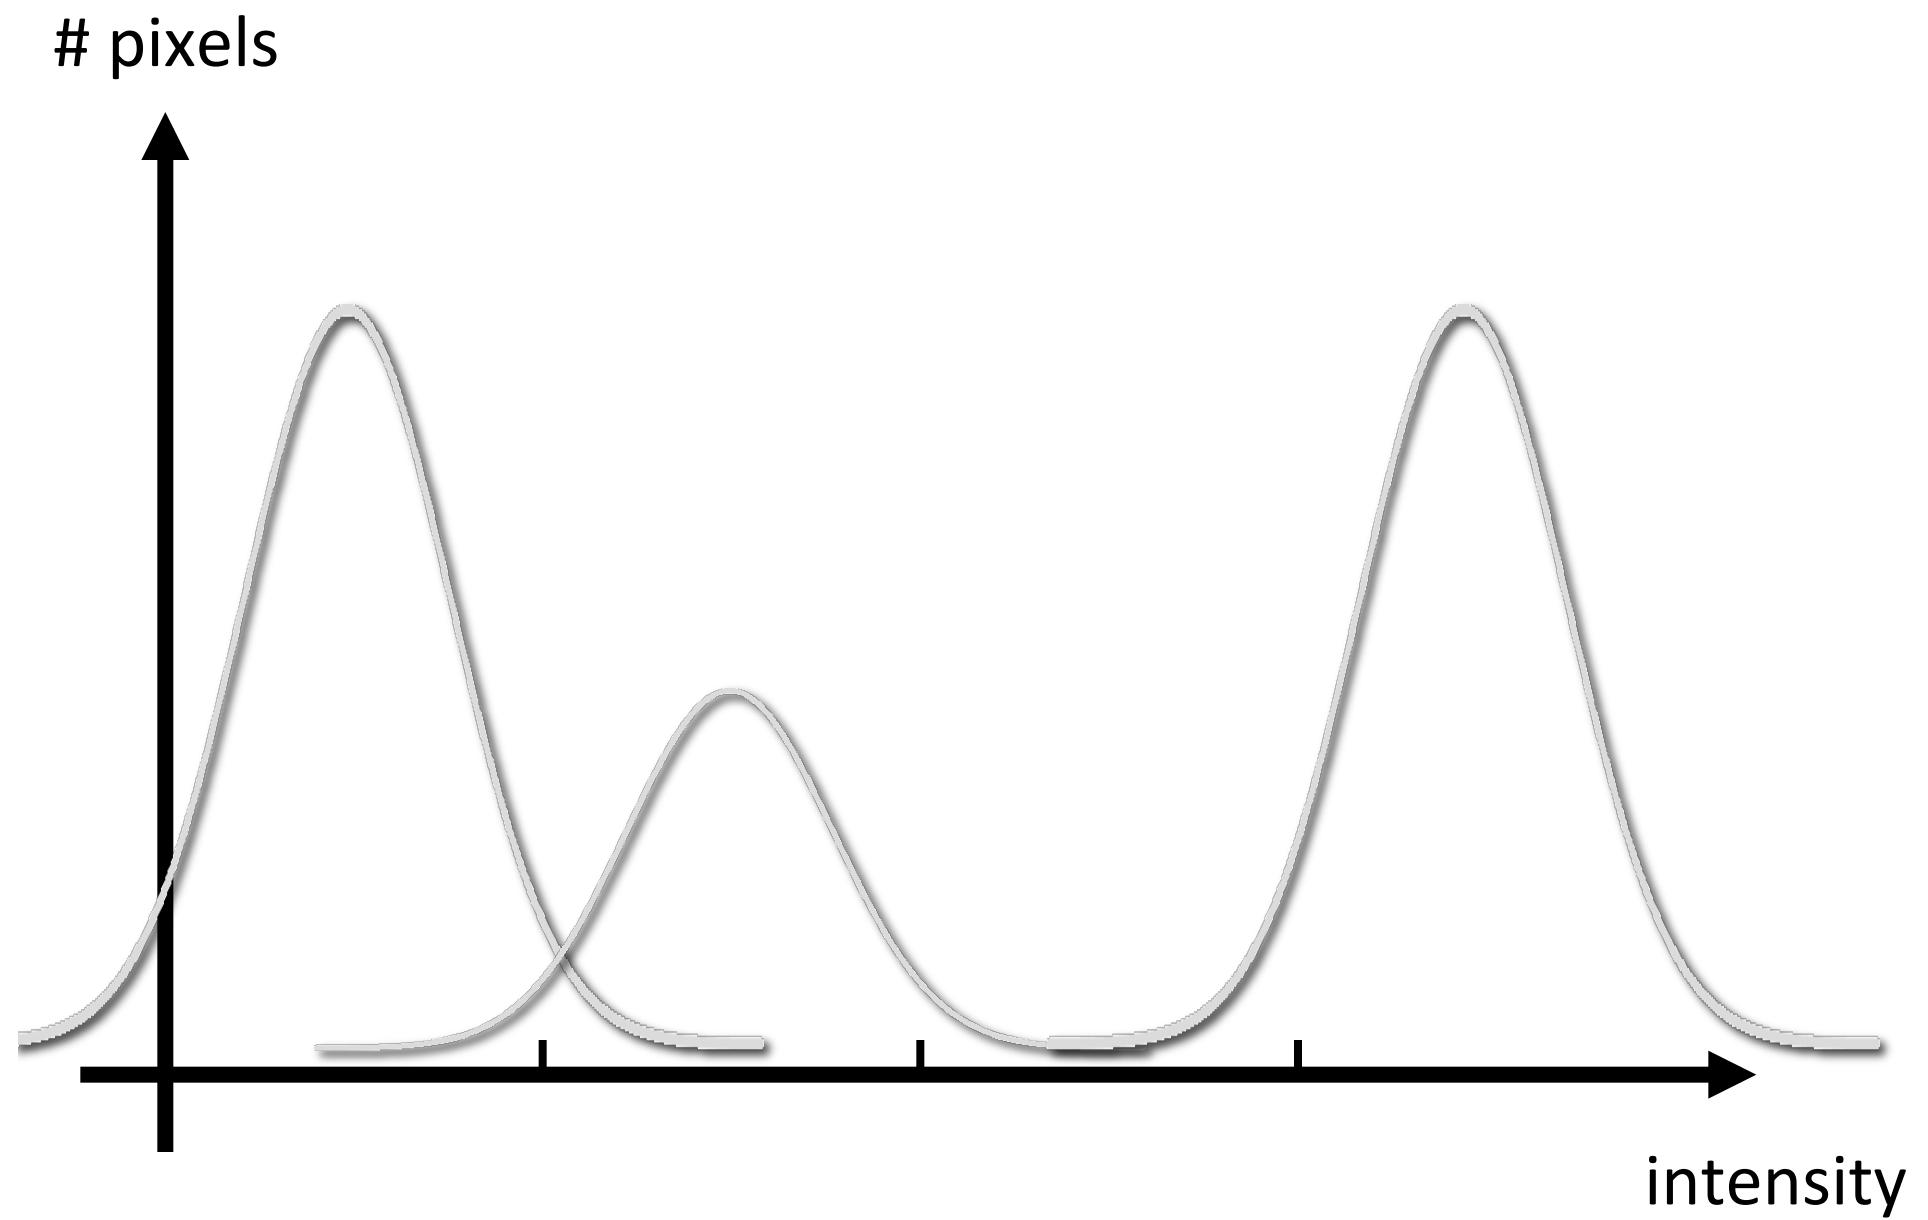


# pixels

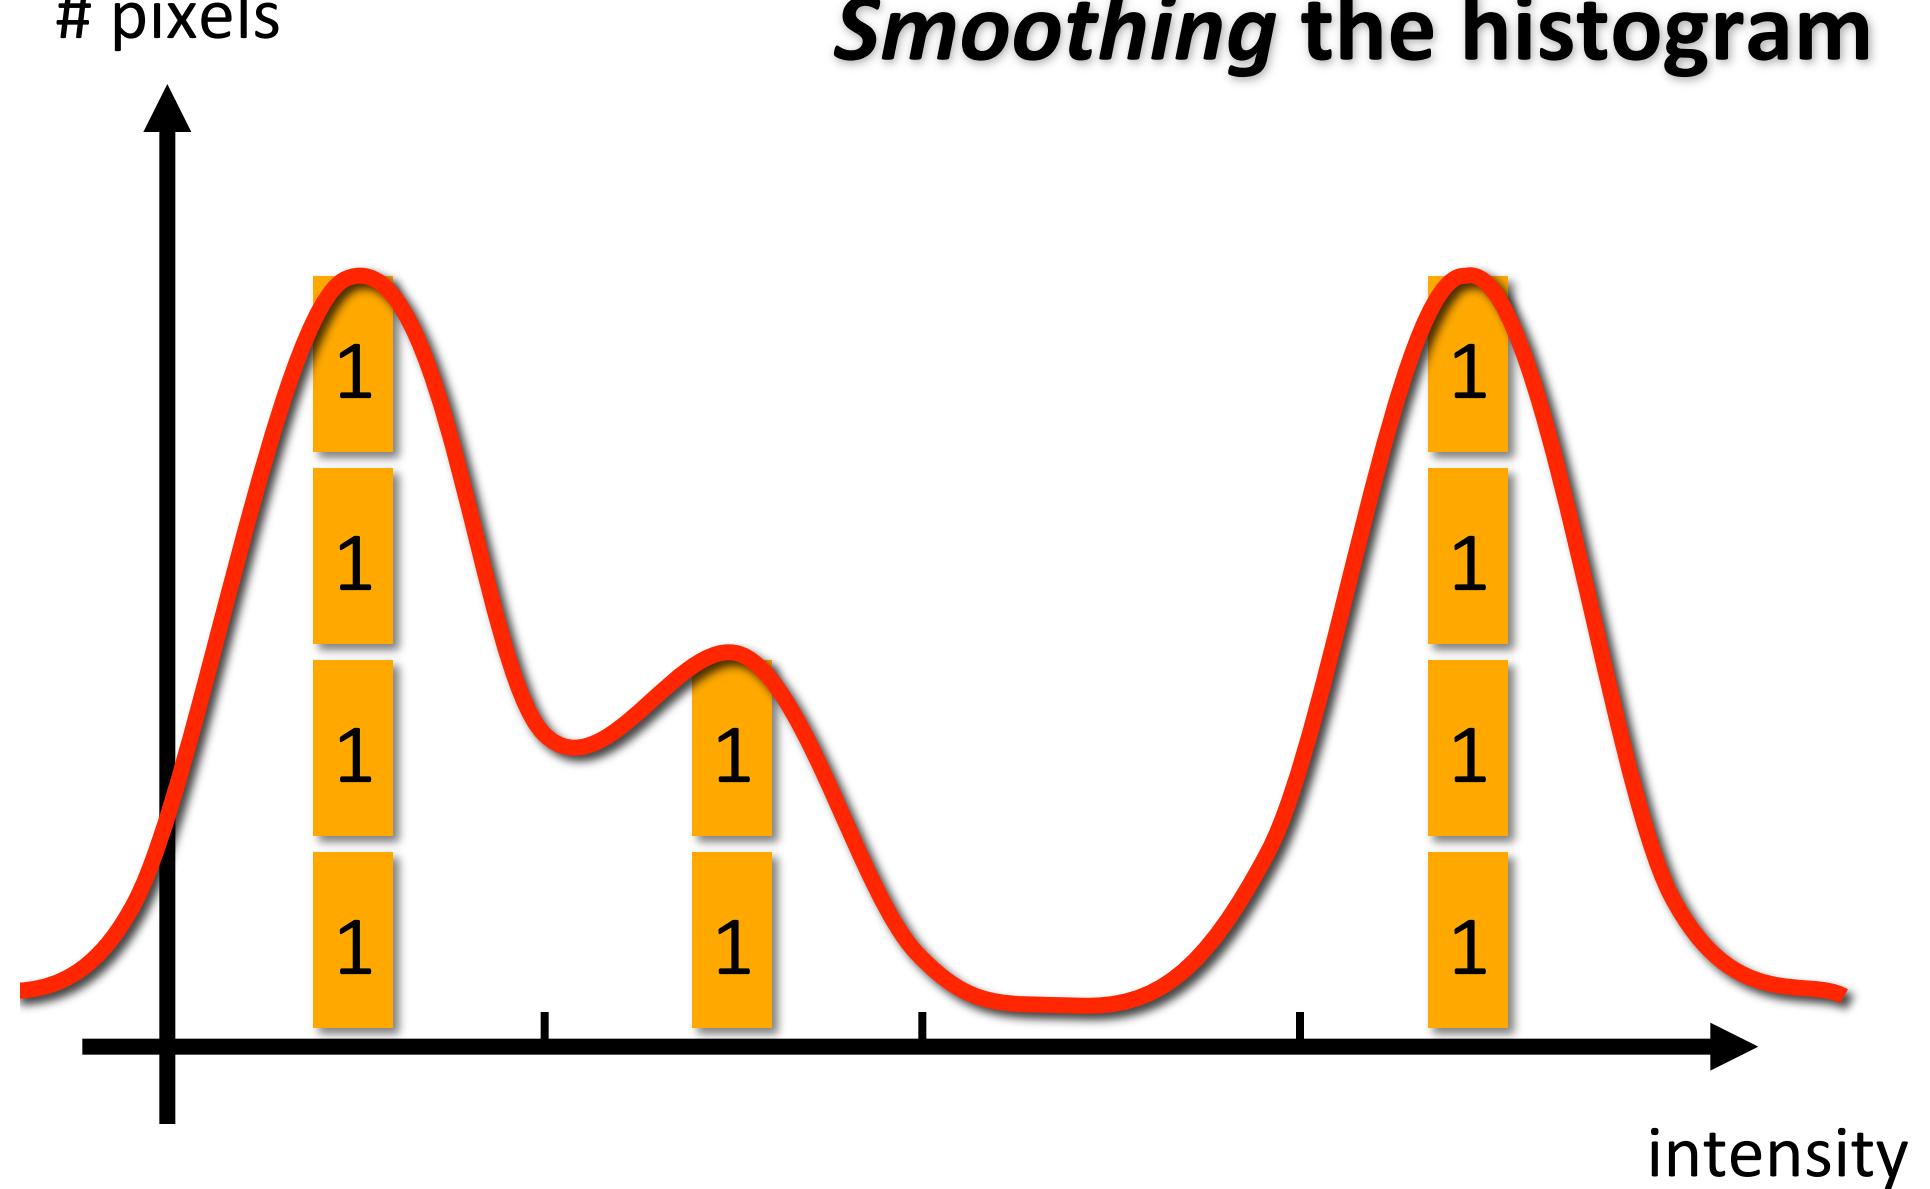






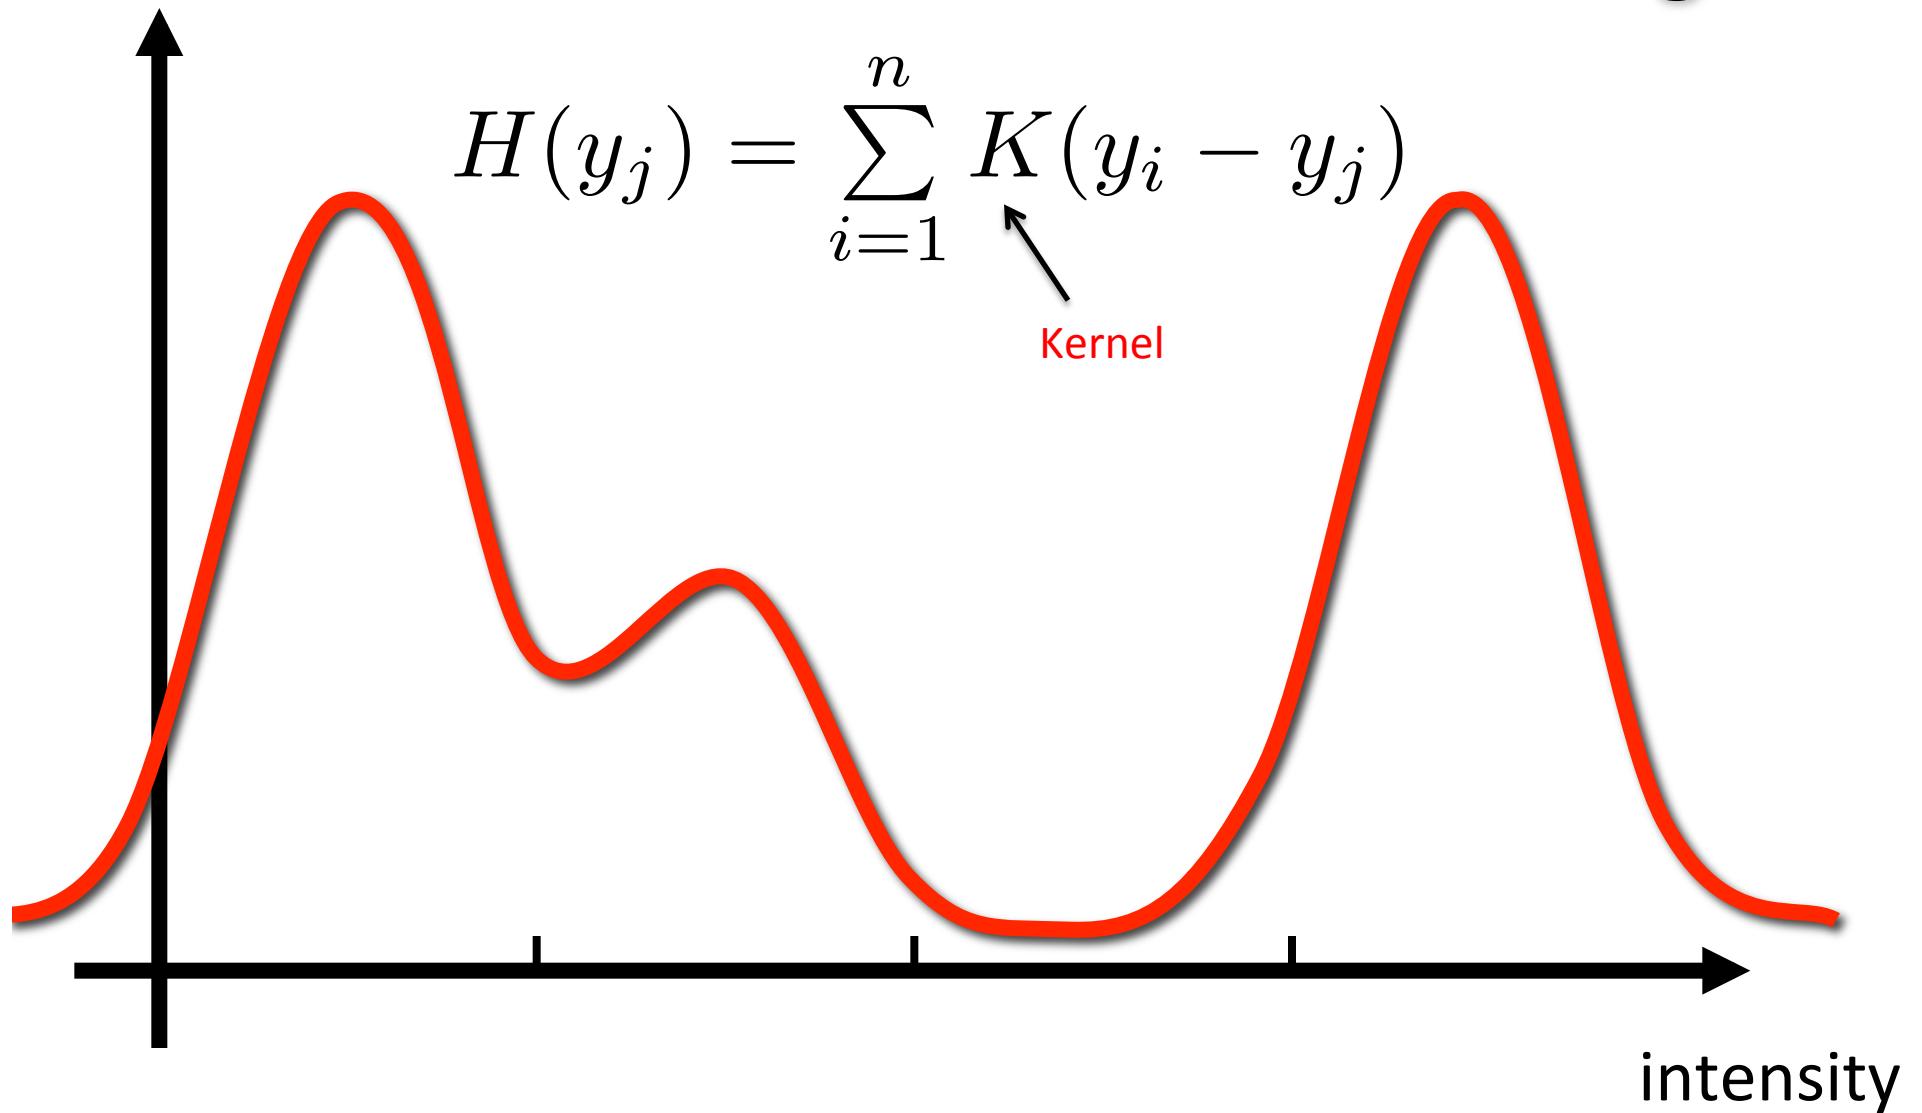


## *Smoothing the histogram*

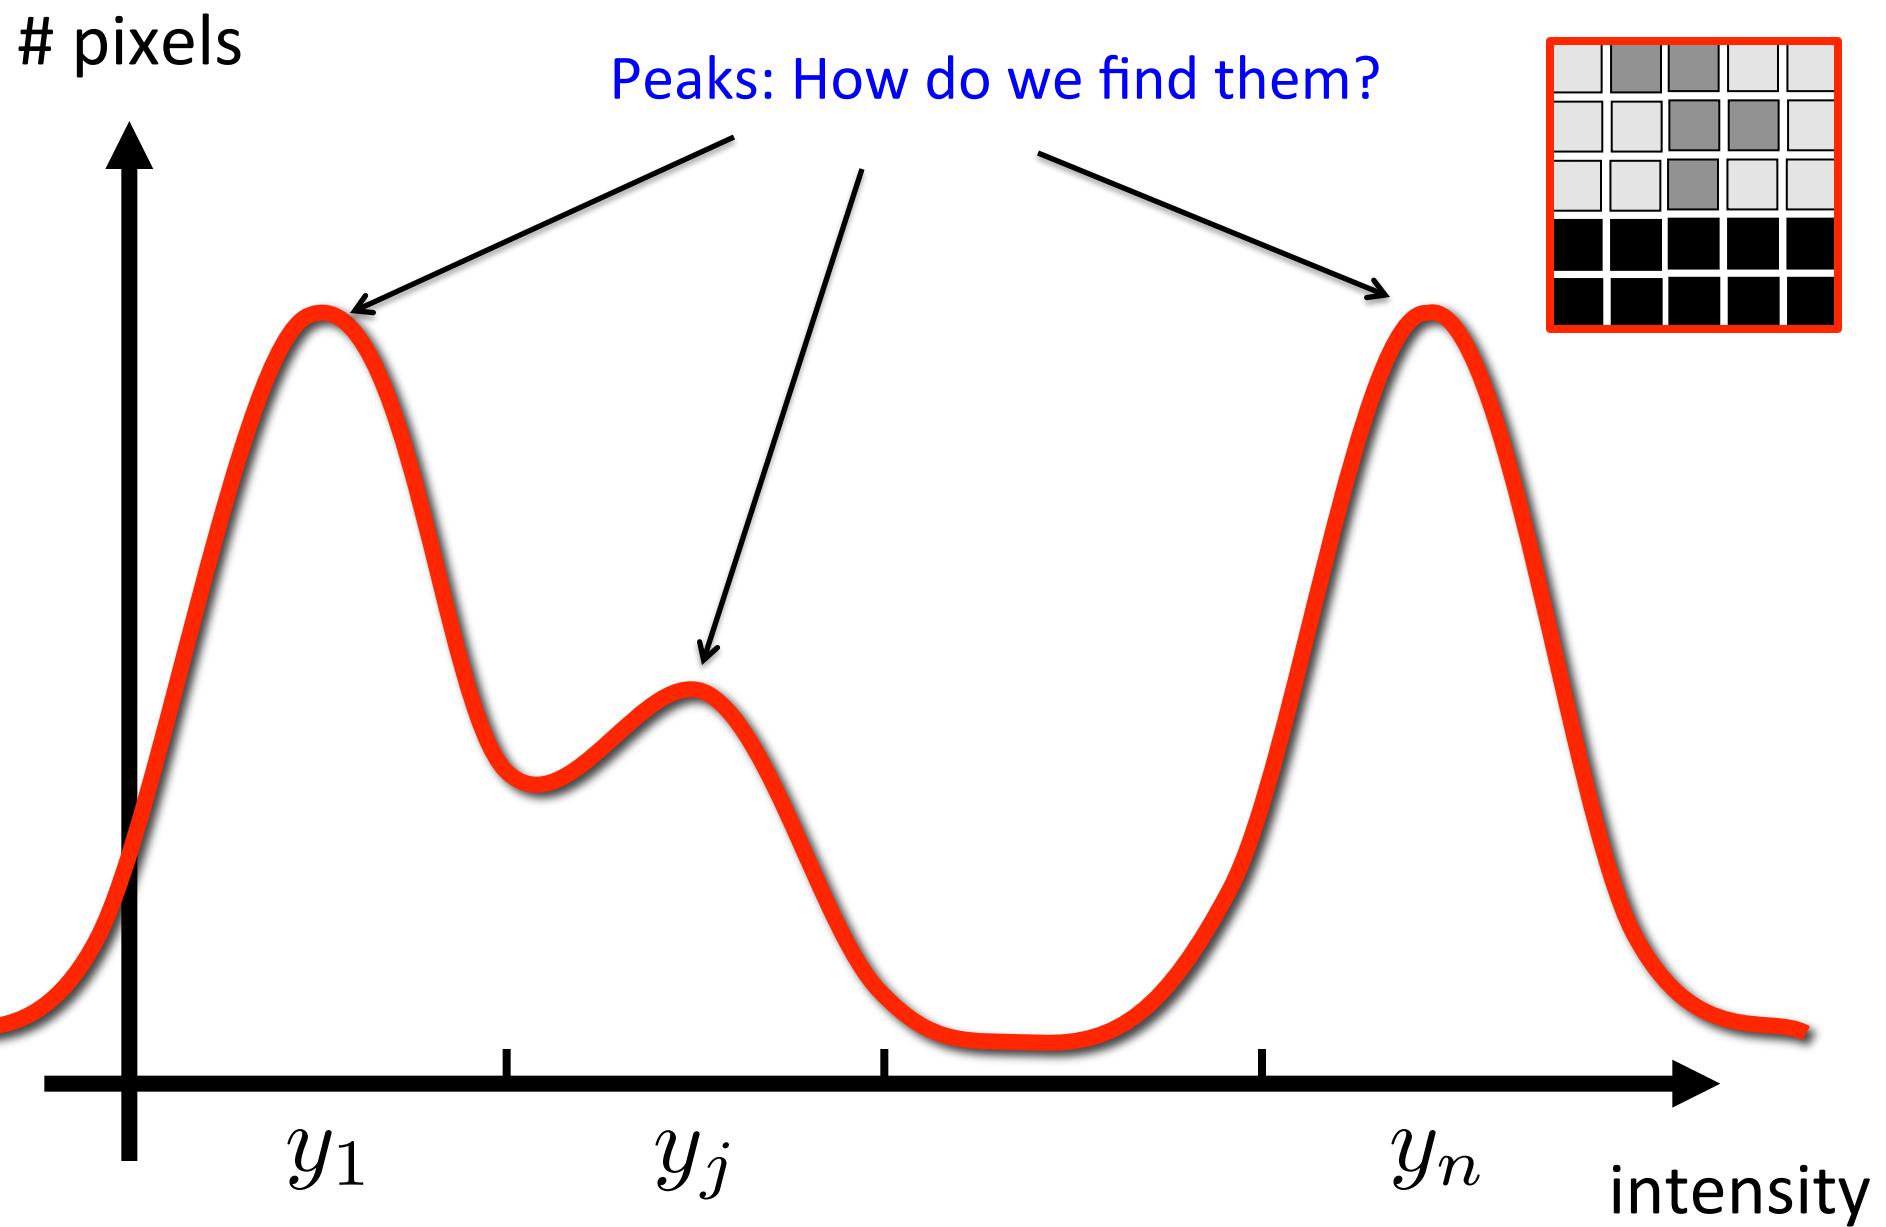


# pixels

## *Smoothed histogram*



# Modes (peaks) of the Histogram



# Mode Finding on the Histogram

Start with  $H(y_j) = \sum_{i=1}^n K(y_i - y_j)$

Fix  $y_j = t$

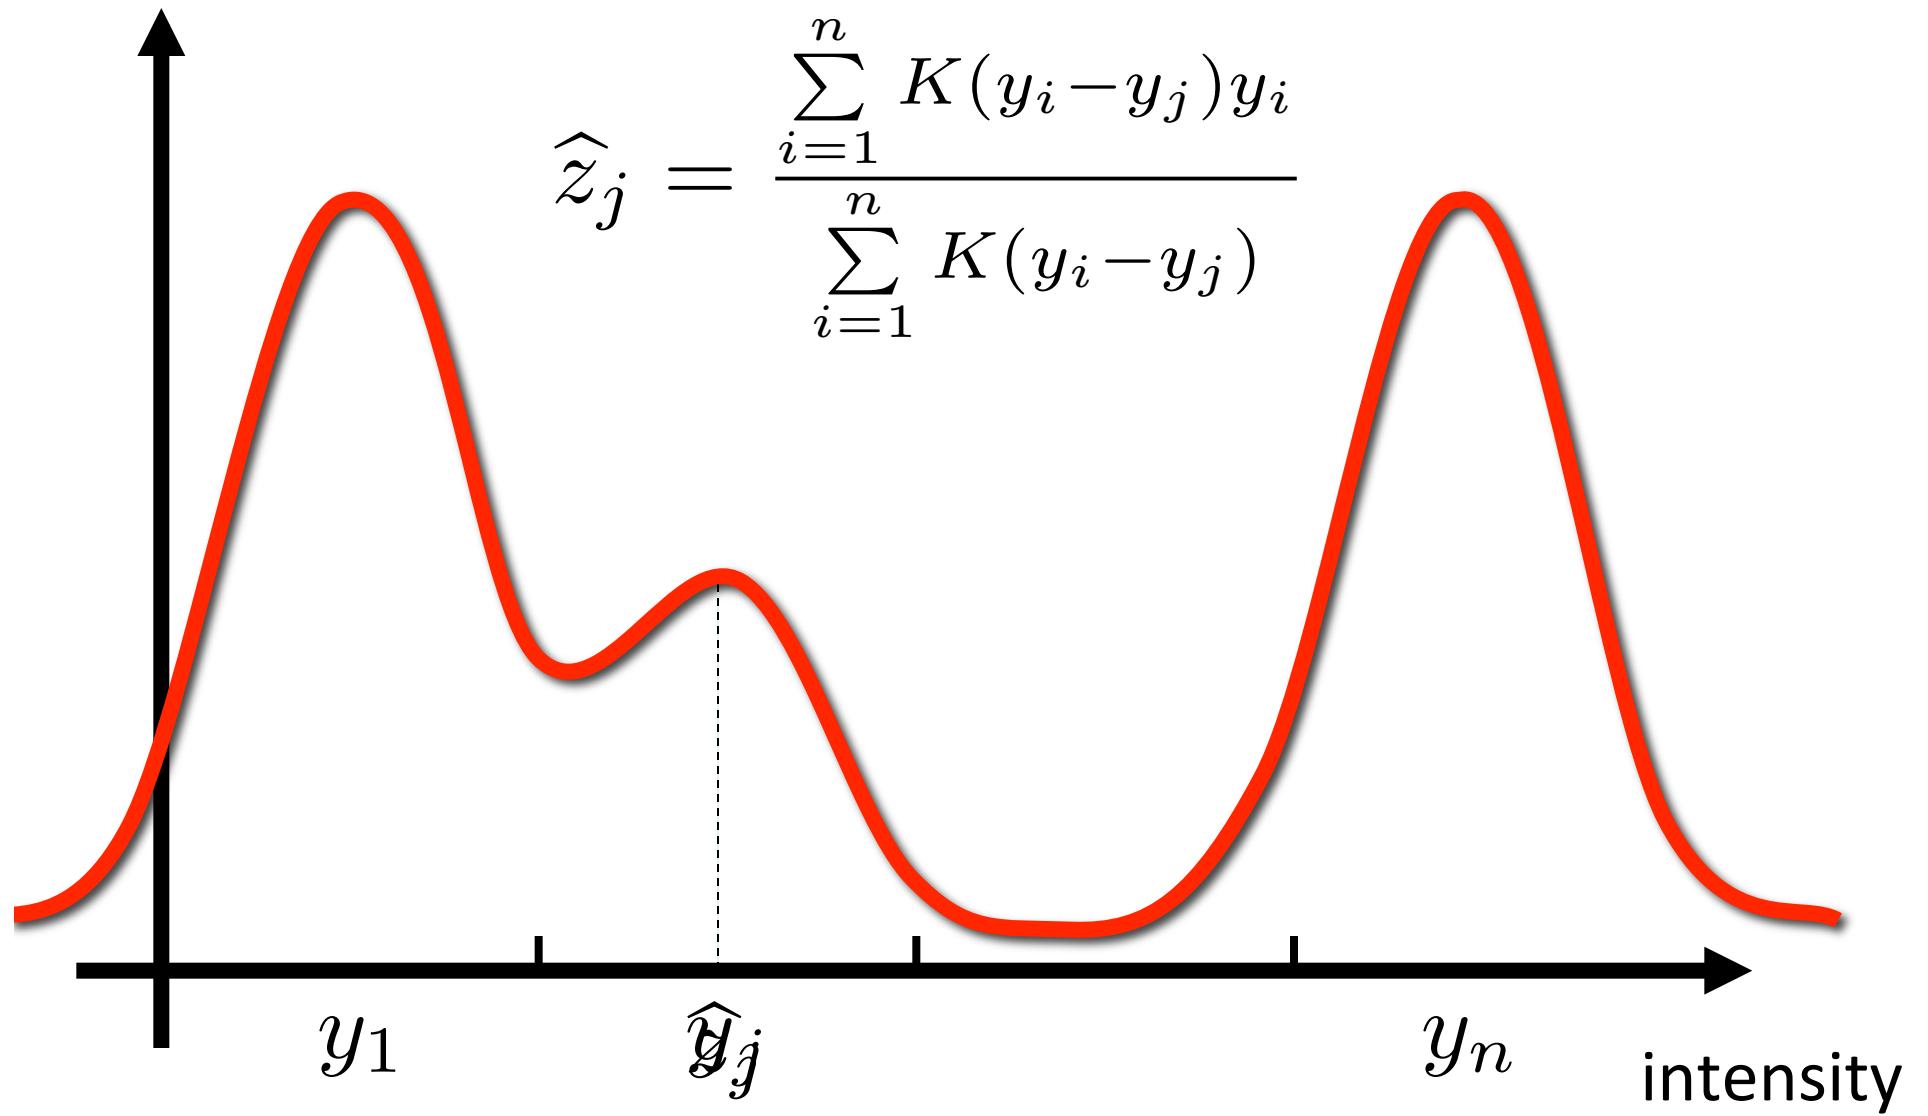
Differentiate and solve:  $\frac{\partial H(t)}{\partial t} = 0$

$$\hat{z}_j = \frac{\sum_{i=1}^n K(y_i - y_j) y_i}{\sum_{i=1}^n K(y_i - y_j)}$$

(Assuming Gaussian Kernel, for now.)

# Peak of the Histogram Nearest (in value) to $y_j$

# pixels



# Let's take a closer look

$$W_{ij} = \frac{K_{ij}}{\sum_i K_{ij}}$$

*For each pixel, must normalize so the weights sum to 1.*

$$\hat{z}_j = \sum_{i=1}^n W_{ij} y_i$$

*Just a weighted average*

$$\sum_{i=1}^n W_{ij} = 1$$

# Bilateral or non-local means kernels

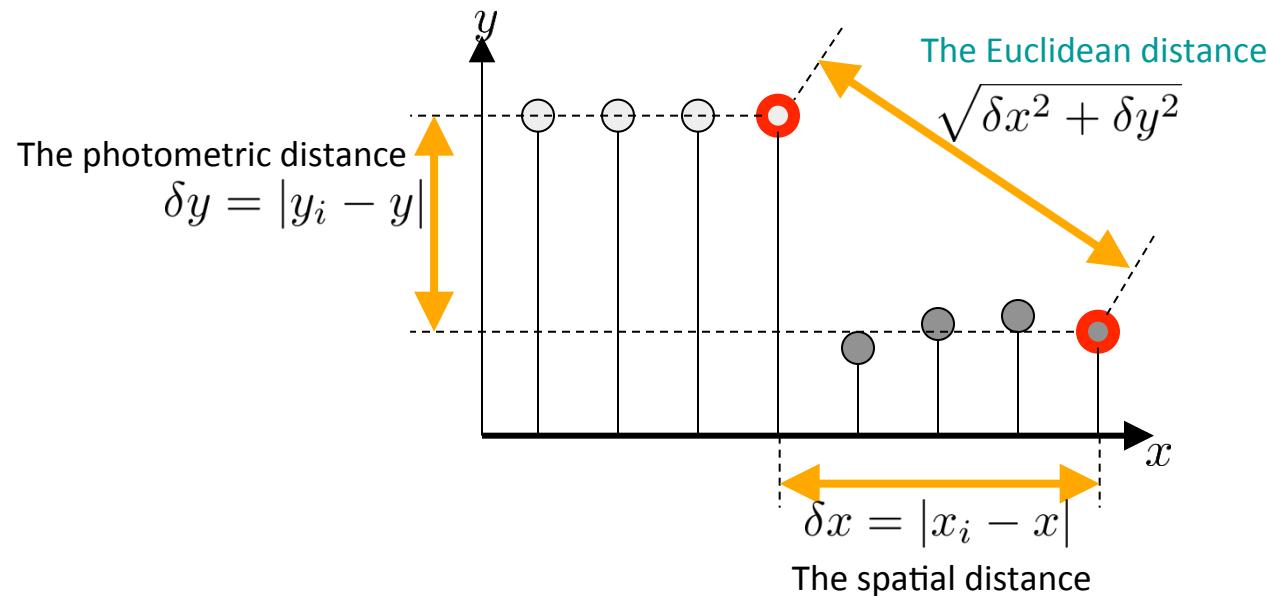
Pixel **values**:

$$K(y_i - y_j) = \exp(-|y_i - y_j|^2/h_y^2)$$

Pixel **positions**:

$$K(x_i - x_j) = \exp(-|x_i - x_j|^2/h_x^2)$$

$$K_{ij} = K(y_i - y_j) K(x_i - x_j)$$



# Faster Normalized Filtering

$$\hat{z}_j = \frac{\sum_{i=1}^n K_{ij} y_i}{\sum_{i=1}^n K_{ij}}$$

$$\hat{z}_j = \frac{\sum_{i=1}^n K_{ij} y_i}{\sum_{i=1}^n K_{ij} - 1}$$

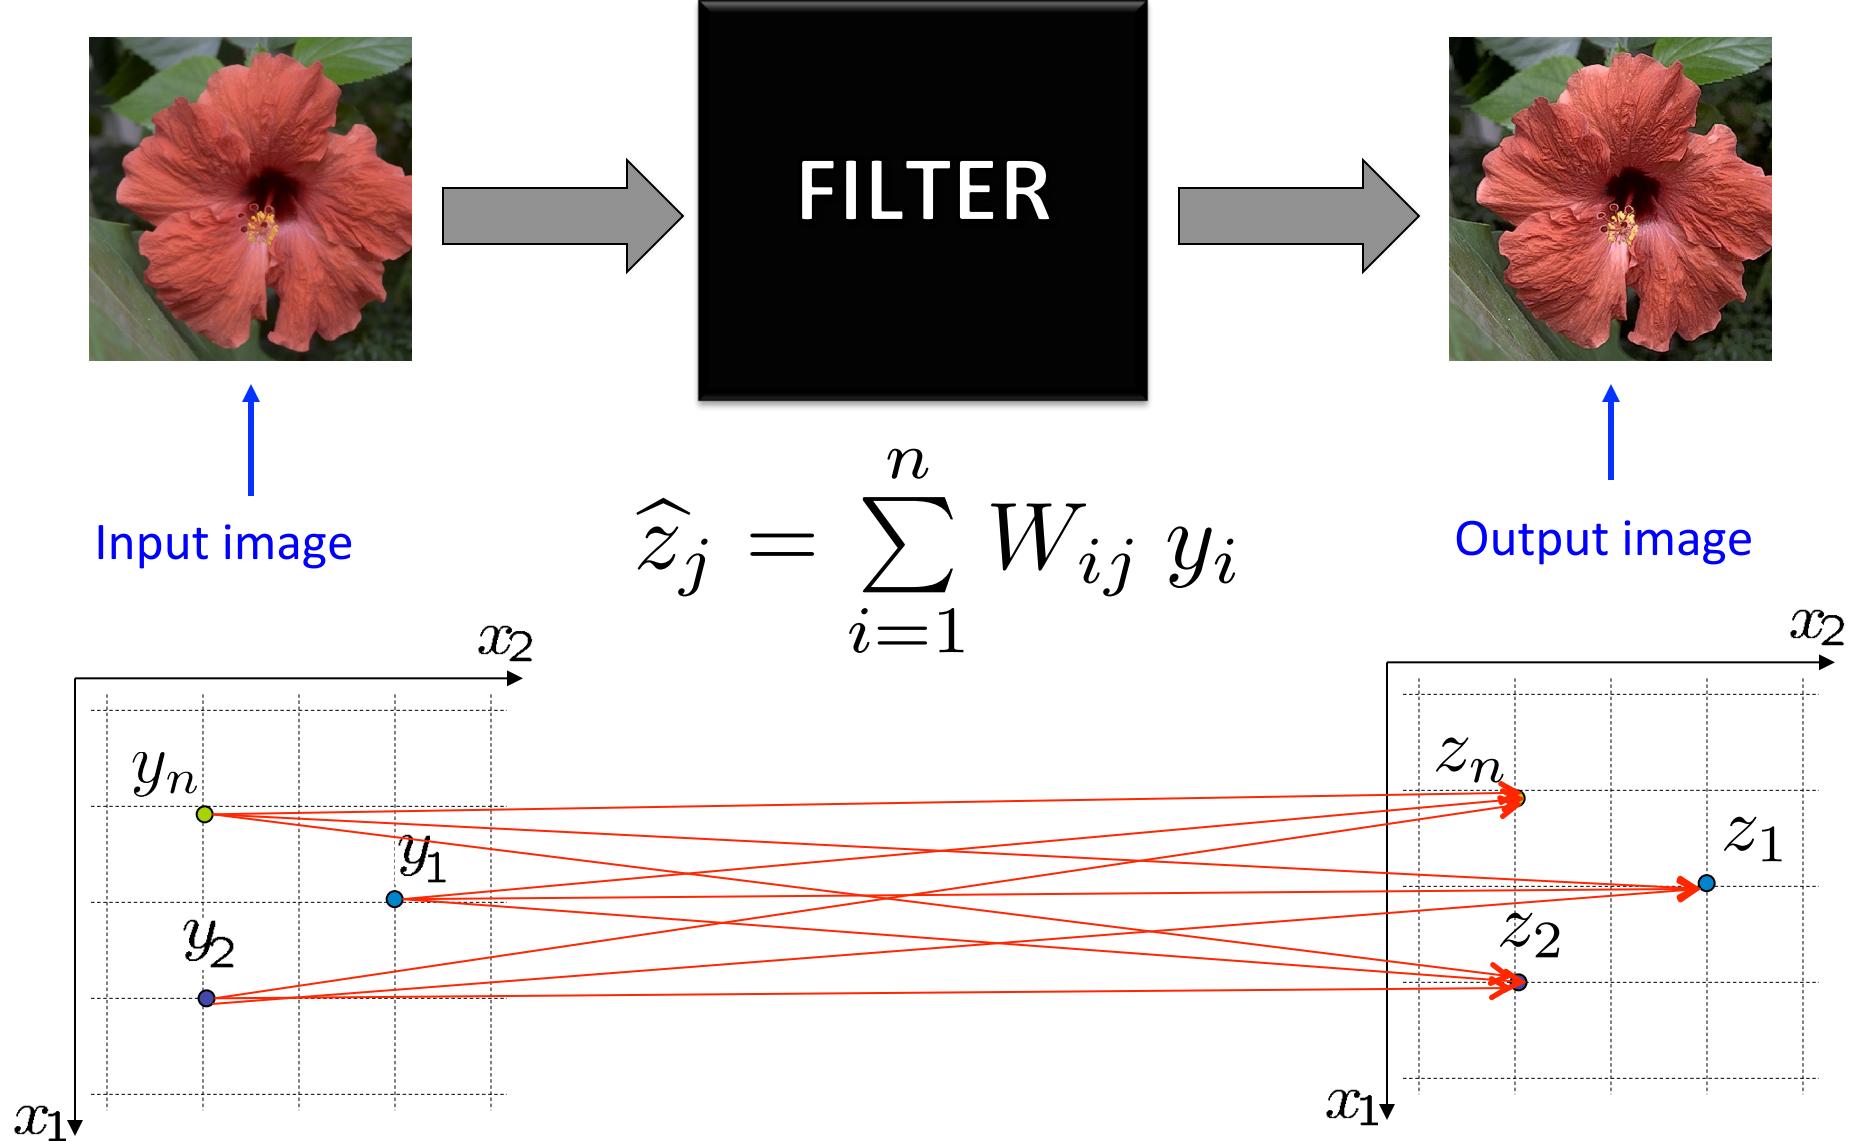
Filter once

Filter twice

$$\hat{z}_j = \frac{\sum_{i=1}^n K_{ij} y_i}{d_j}$$

Divide Pixel-wise

# The Black Box



# Global (or Local) Filters

Each output pixel  $\longrightarrow \hat{z}_j = \sum_{i=1}^n W_{ij} y_i$  All input pixels

$$\hat{z}_j = \sum_{i=1}^n W_{ij} y_i$$

$$\mathbf{w}_j^T = [W_{1j}, W_{2j}, \dots, W_{nj}]$$
$$\mathbf{W} = \begin{bmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{bmatrix}$$

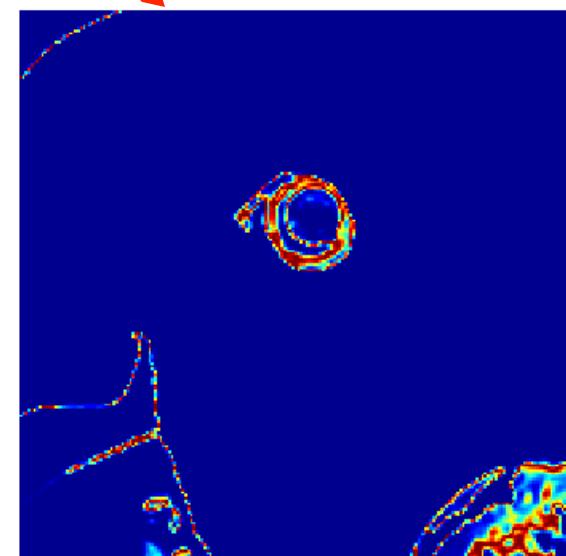
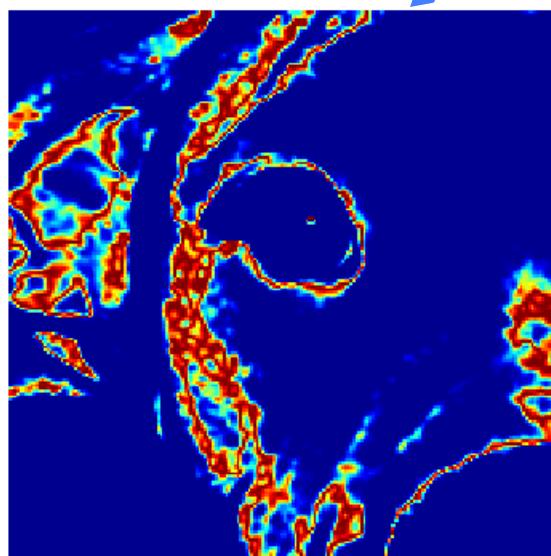
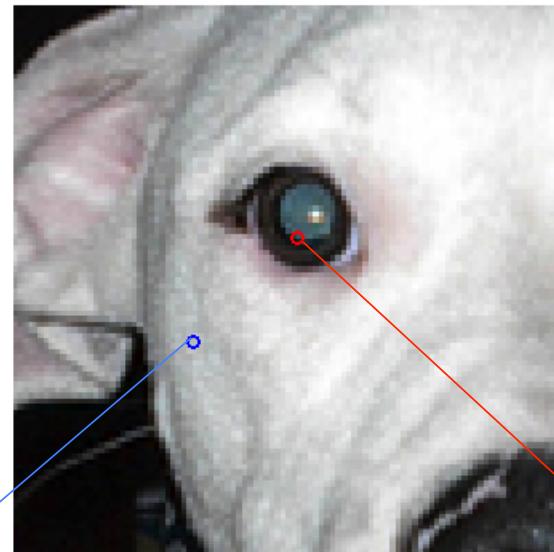
rows

$$\hat{\mathbf{z}} = \mathbf{W} \mathbf{y}$$

Data-dependent matrix

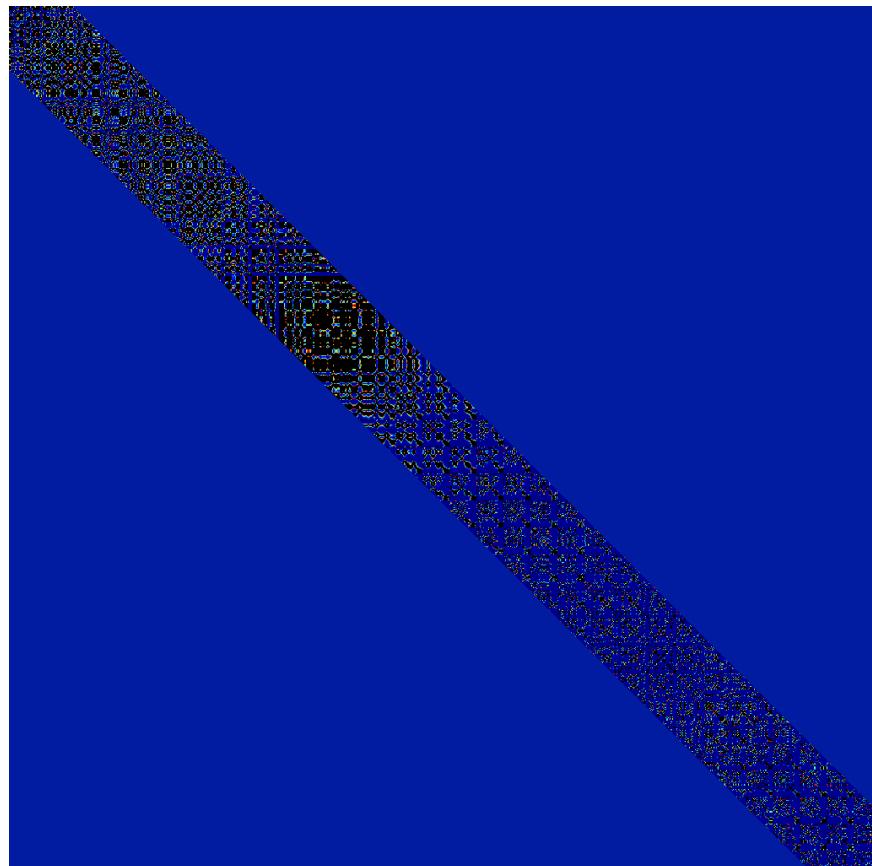
Image scanned into a vector

# Global Affinities



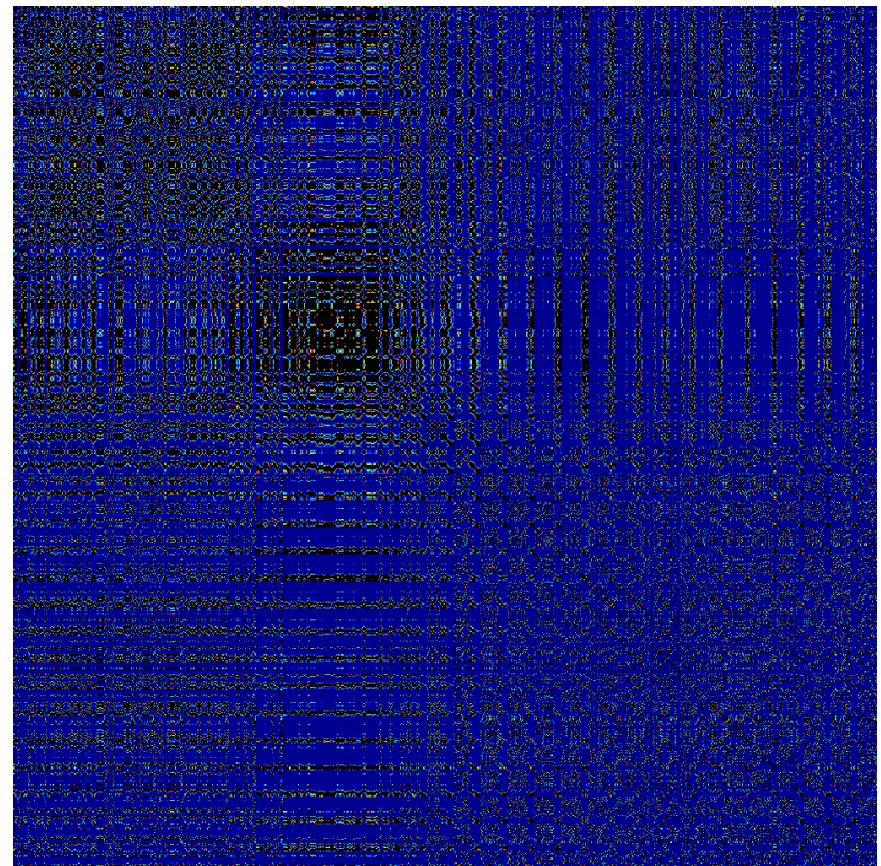
# Local vs. Global Filter Matrix $W$

Local Filters



Sparse, but high-rank

Global Filters



Dense but low-rank

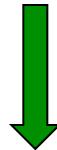
# A Linear Embedding Approach

$$\hat{\mathbf{z}} = \mathbf{W} \mathbf{y}$$

$$\mathbf{W} = \mathbf{V} \mathbf{S} \mathbf{V}^T$$

Orthogonal Eigenvectors  
 $\mathbf{V} = [\mathbf{v}_1, \dots, \mathbf{v}_n]$

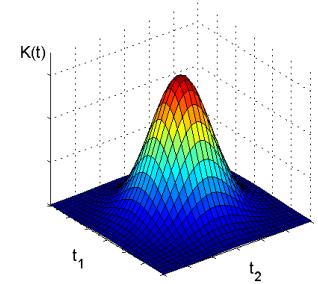
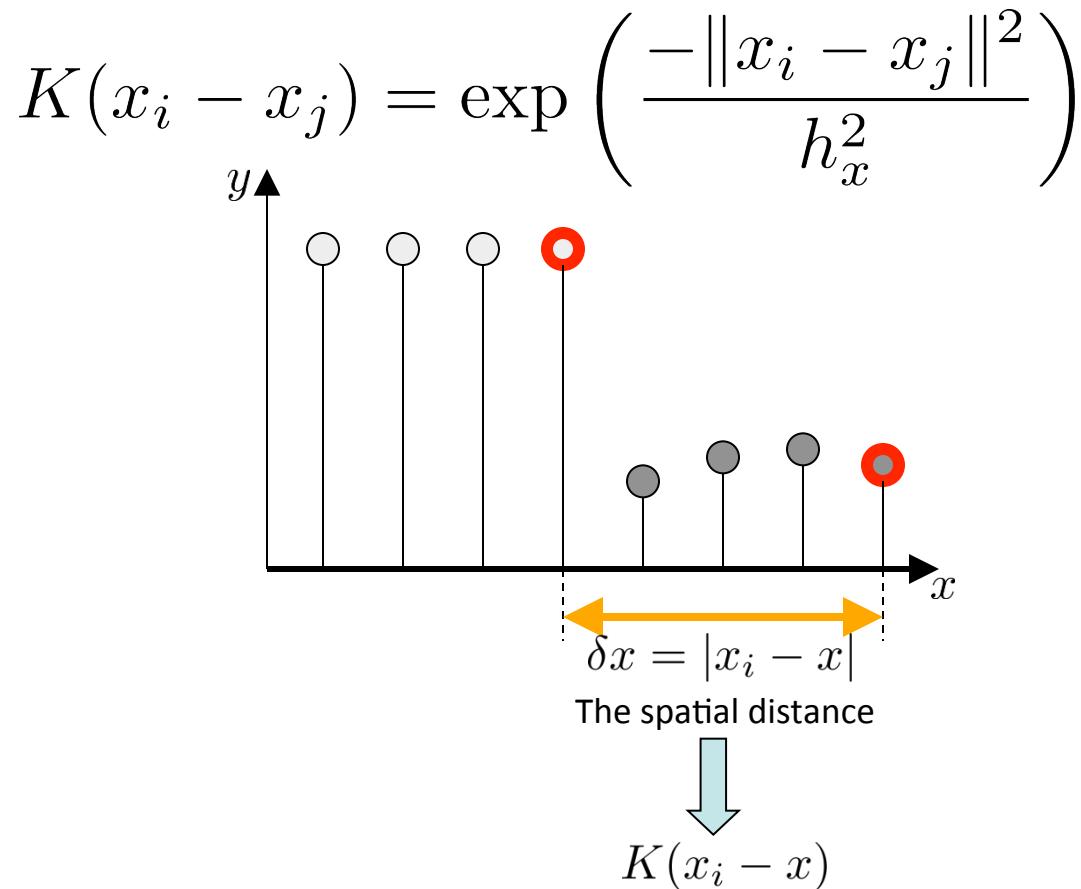
Eigenvalues  
 $\mathbf{S} = \text{diag}[\lambda_1, \dots, \lambda_n]$   
 $0 \leq \lambda_n \leq \dots < \lambda_1 = 1$



$$\hat{\mathbf{z}} \approx \sum_i \alpha_i \times \text{top eigenvectors of } \mathbf{W}$$

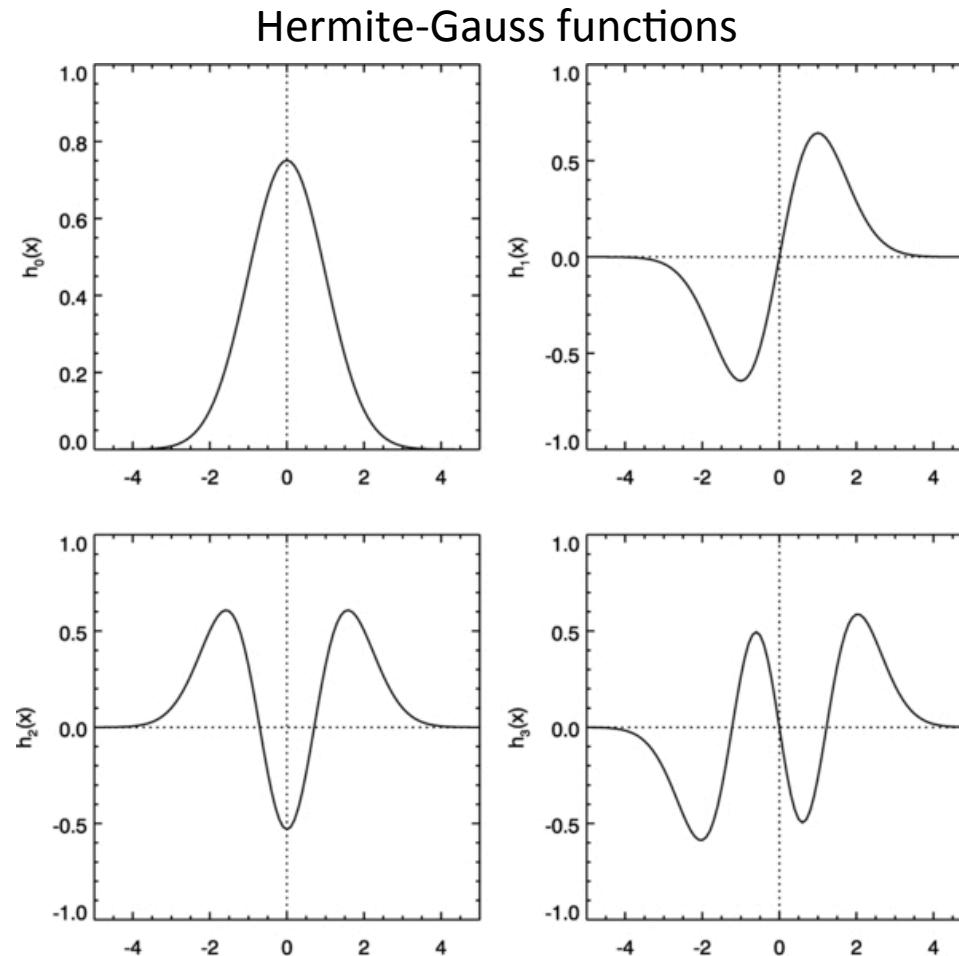
*Nonlinear Filtering is now “Linear”*

# Choice of Kernels



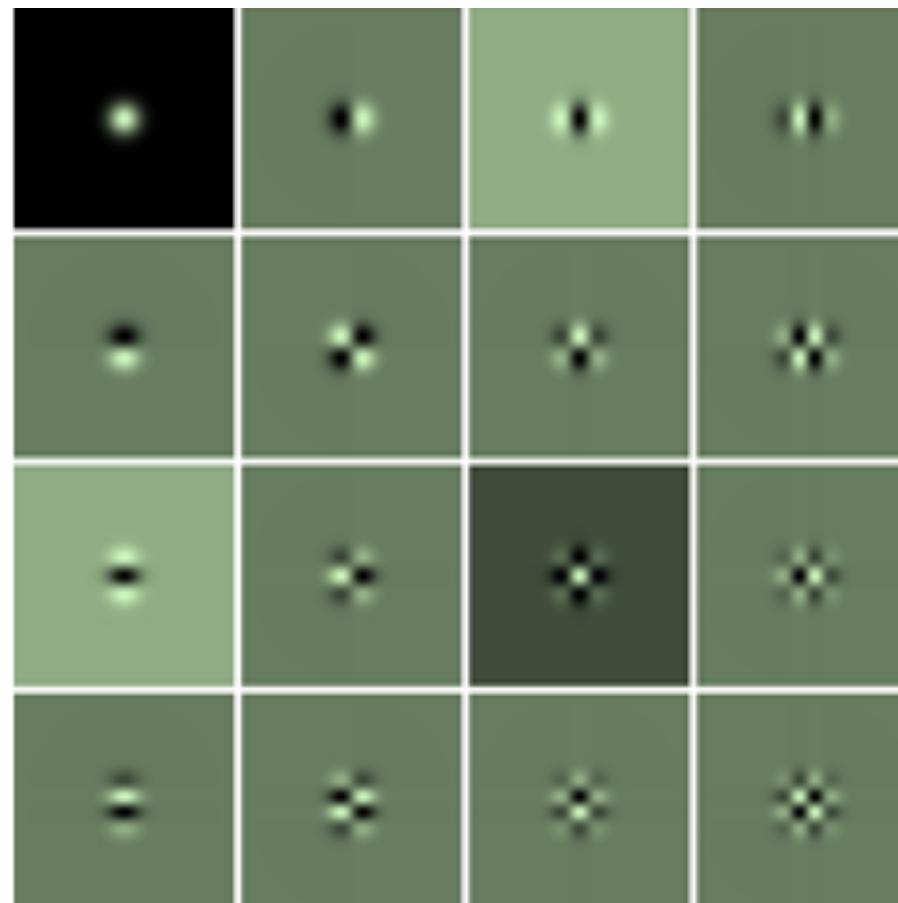
- Classical Gaussian Linear Filters

# Eigenvectors of Linear Gaussian Kernel



# Eigenvectors of Linear Gaussian Kernel

Hermite-Gauss functions

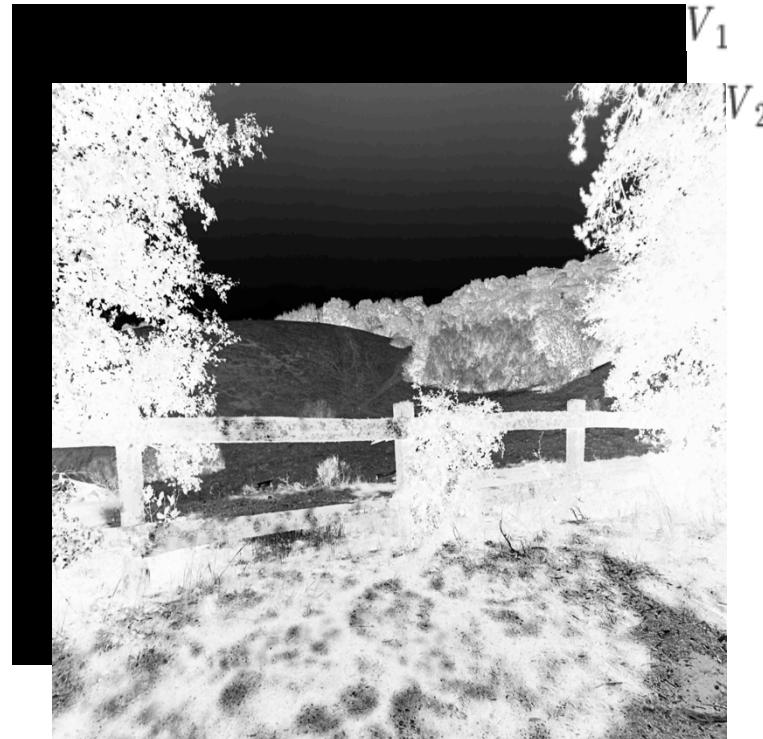


# Filter Eigenvectors (For NLM)

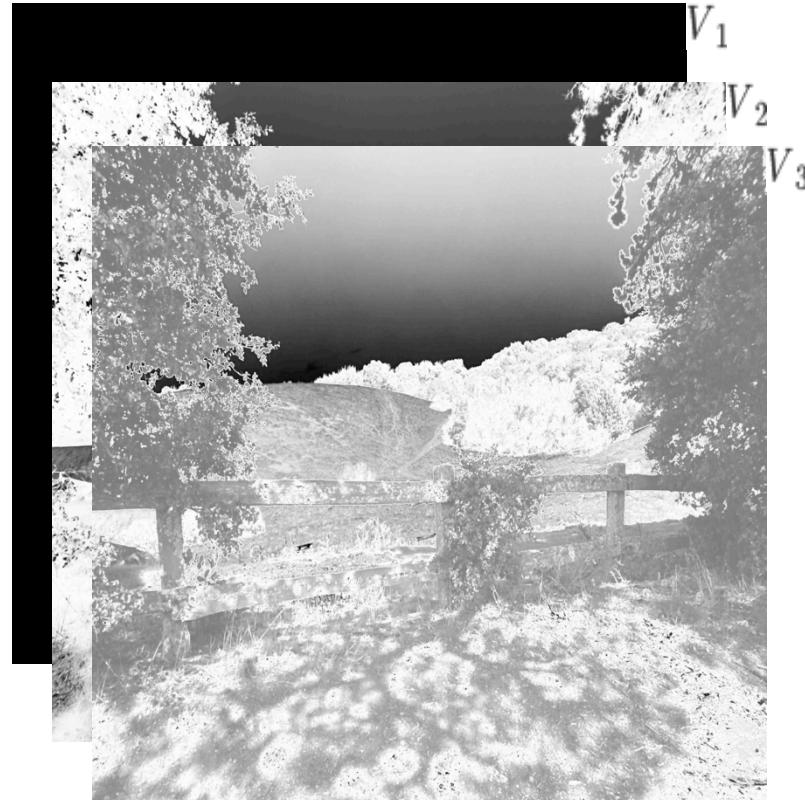
(constant image)



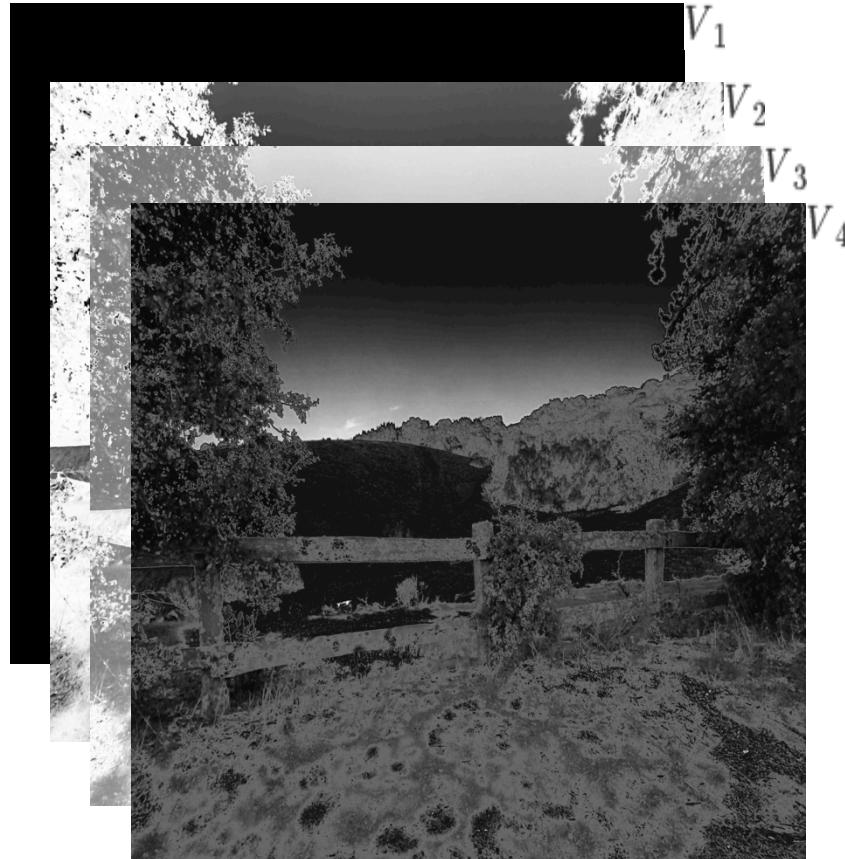
# Filter Eigenvectors



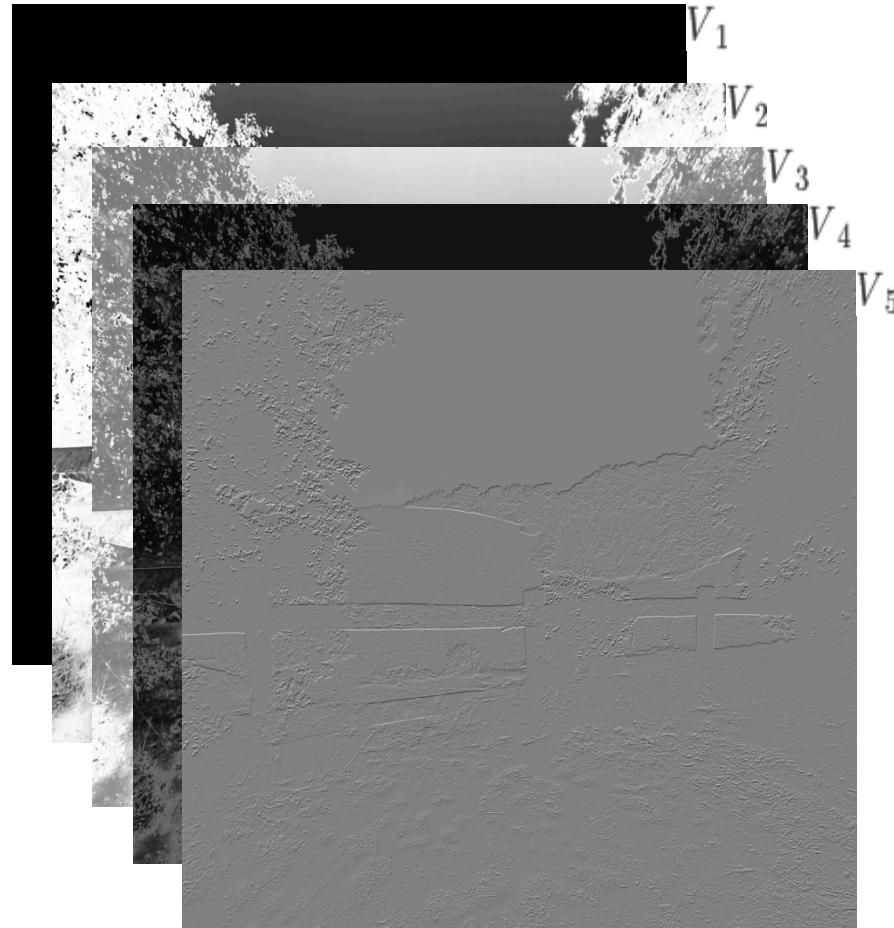
# Filter Eigenvectors



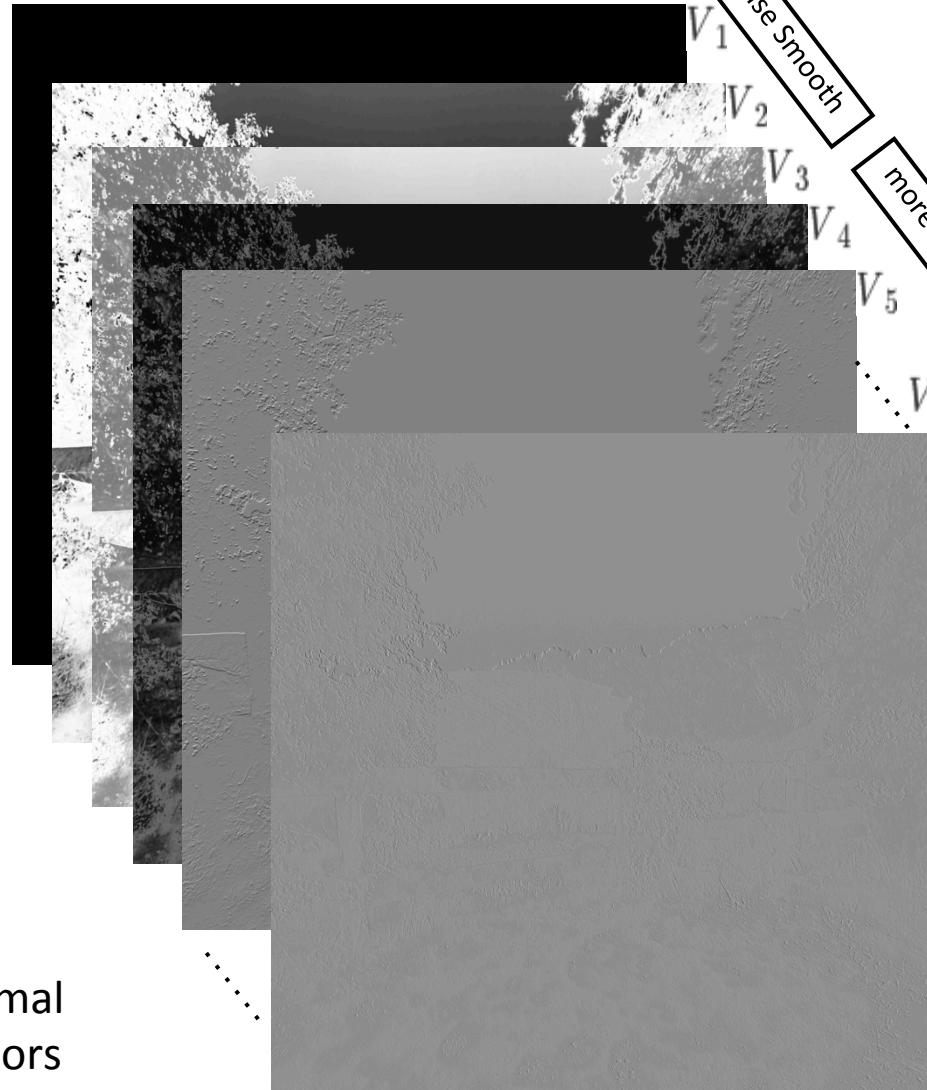
# Filter Eigenvectors



# Filter Eigenvectors



# Filter Eigenvectors



# Nystrom Approximation



Spatially uniform sampling

**A:** Sampled pixels ( $m$ )

**B:** Remaining pixels ( $n-m$ )

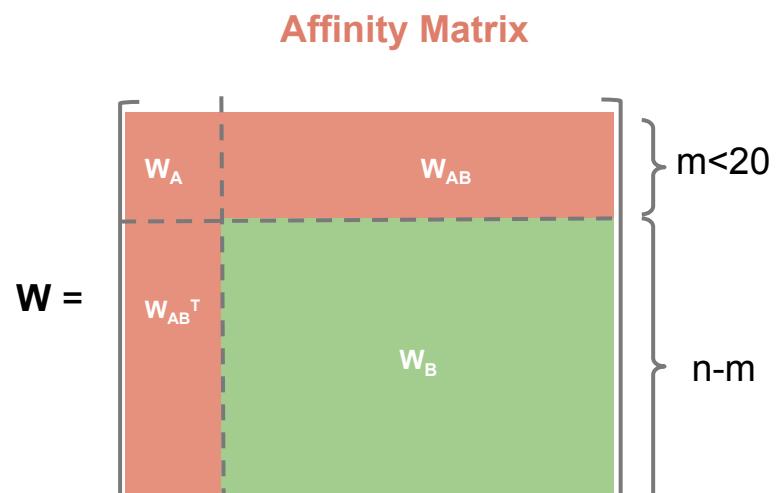
# Nystrom Approximation



Spatially uniform sampling

A: Sampled pixels (m)

B: Remaining pixels (n-m)



Explicitly computed weights

Approximated weights

# Other filters from a base filter

$$\text{Output Image} \quad \downarrow \quad \text{Input Image} \quad \downarrow$$
$$\widehat{\mathbf{z}} = \mathbf{V} f(\mathbf{S}) \mathbf{V}^T \mathbf{y}$$

  
**Filter controls**

 **Shrinkage/Boosting function**  
 $f(\mathbf{S}) = \text{diag}[f(\lambda_1), \dots, f(\lambda_n)]$



$$\widehat{\mathbf{z}} \approx \sum_i f(\alpha_i) \times \text{top eigenvectors of } \mathbf{W}$$

# Nonlinear Filtering is now “Linear”

$$\text{Output Image} \quad \downarrow \quad \text{Input Image} \quad \downarrow$$
$$\widehat{\mathbf{z}} = \mathbf{V} f(\mathbf{S}) \mathbf{V}^T \mathbf{y}$$

  
**Filter controls**

↗ Shrinkage/Boosting function  
 $f(\mathbf{S}) = \text{diag}[f(\lambda_1), \dots, f(\lambda_n)]$

---

Diffusion:  
(Perona-Malik '90, Coifman et al. '06, ....)

$$f(t) = t^p$$

Twicing, Reaction-Diffusion, Boosting, Bregman iteration:  
(Tukey '77, Nordstrom '90, Buhlmann et al. '03, Osher et al. '05)

$$f(t) = 1 - (1 - t)^p$$

$p$  can be **any** positive real number

Local Laplacian Filter



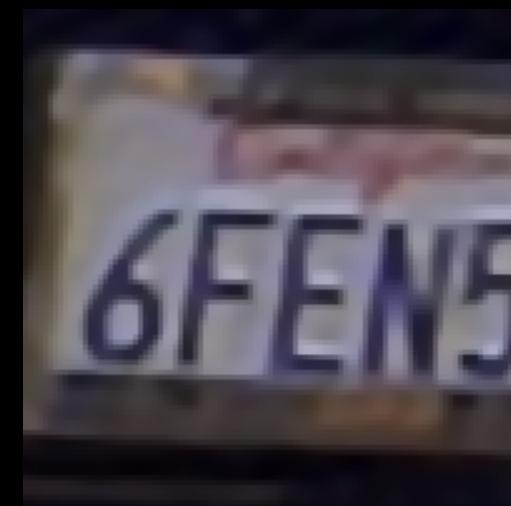
# Multi-frame Upscaling/Zoom





Fused 4 images, denoised and upscaled by 4x





# Relevant Papers

- “A Tour of Modern Image Filtering”,  
P. Milanfar, IEEE Signal Processing Magazine,  
no. 30, pp. 106–128, Jan. 2013
- “A General Framework for Regularized,  
Similarity-based Image Restoration”,  
A. Kheradmand, and P. Milanfar, IEEE Trans on  
Image Processing, vol. 23, no. 12, Dec. 2014
- “Global Image Denoising”, H. Talebi, and P.  
Milanfar, IEEE Trans on Image Processing, vol.  
23, no. 2, pp. 755-768, Feb. 2014
- “Nonlocal Image Editing”, H. Talebi, and P.  
Milanfar, IEEE Trans on Image Processing, vol.  
23, no. 10, Oct. 2014



<http://milanfar.org>

Thank you.